

System Calls, fork(), exec()

Abhijit A. M.
abhijit.comp@coep.ac.in

(C) Abhijit A.M.

Available under Creative Commons Attribution-ShareAlike License V3.0+

Credits: Slides of “OS Book” ed10.

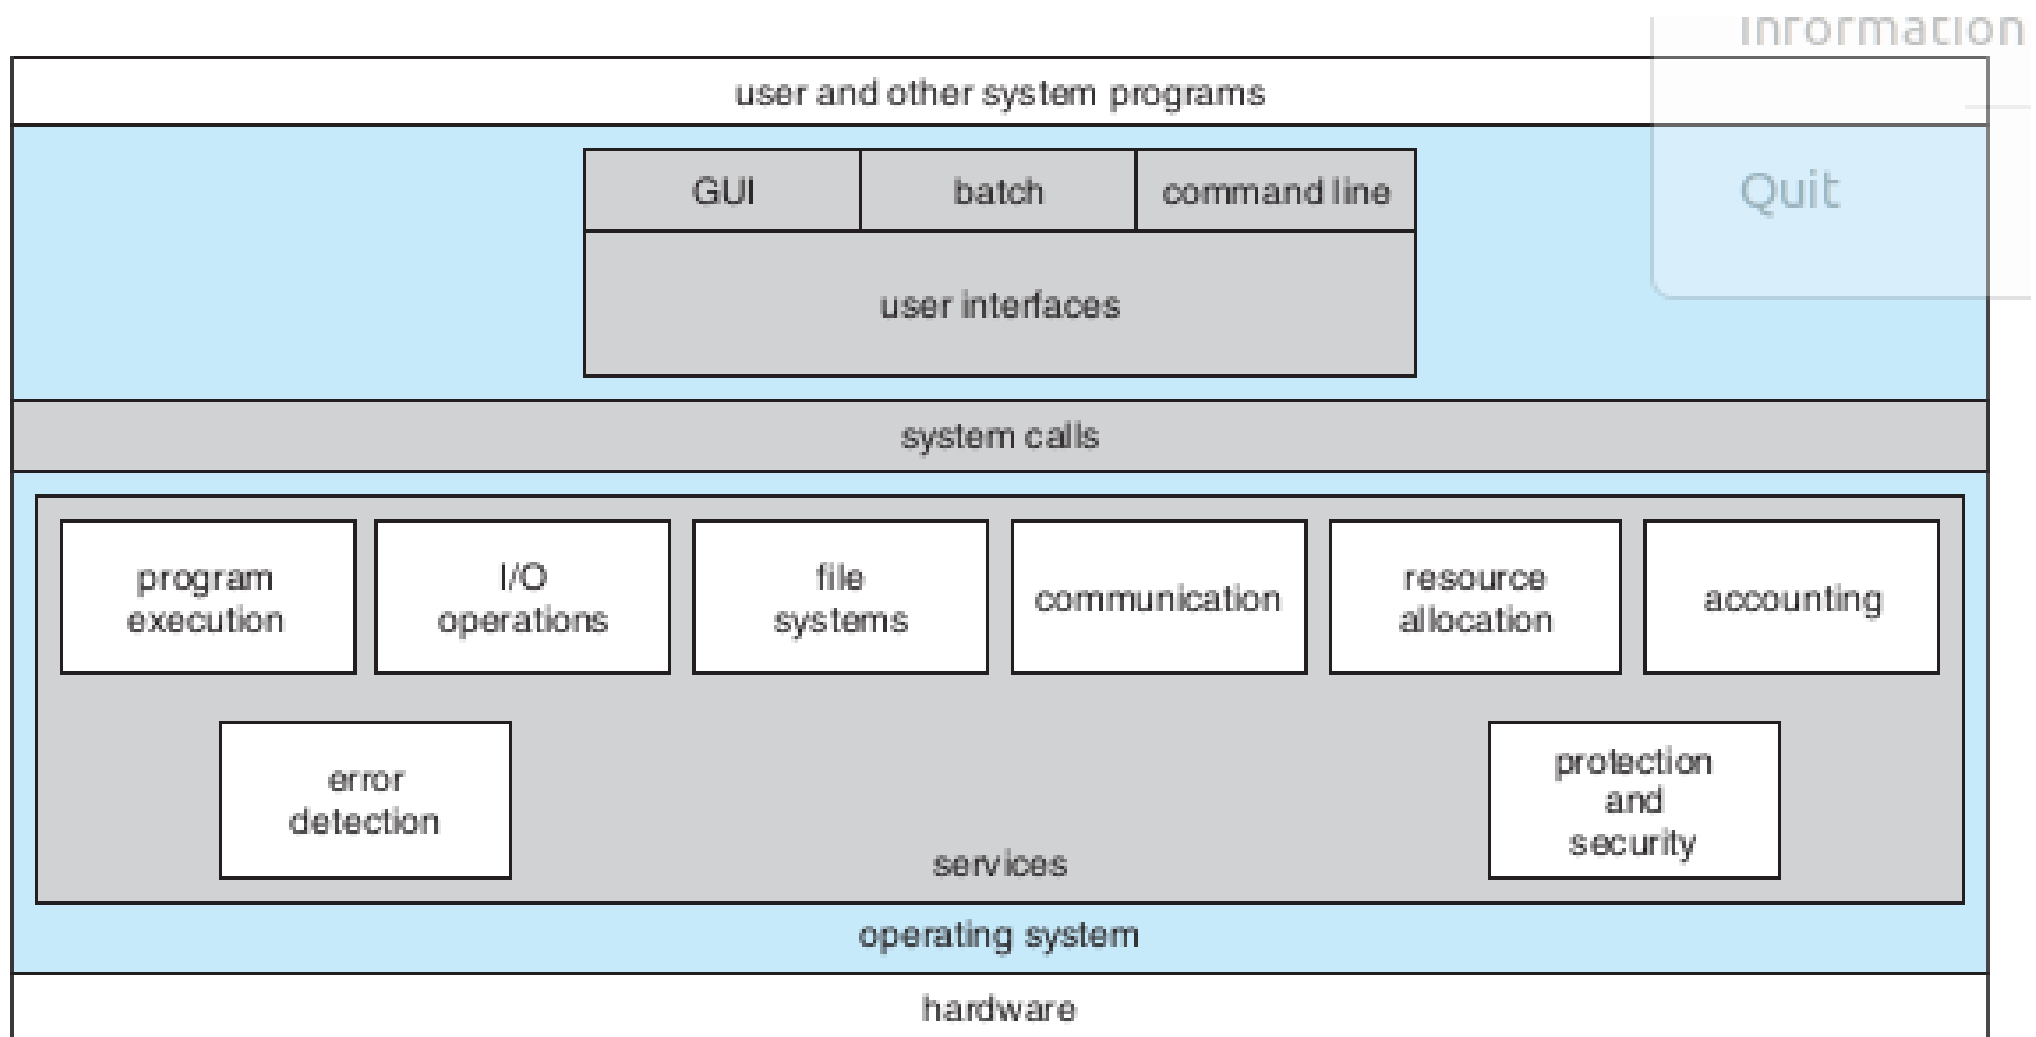


Figure 2.1 A view of operating system services.

System Calls

- **Services provided by operating system to applications**
 - Essentially available to applications by calling the particular software interrupt application
 - All system calls essentially involve the “INT 0x80” on x86 processors + Linux
 - Different arguments specified in EAX register inform the kernel about different system calls
- **The C library has wrapper functions for each of the system calls**
 - E.g. open(), read(), write(), fork(), mmap(), etc.

Types of System Calls

- **File System Related**
 - Open(), read(), write(), close(), etc.
- **Processes Related**
 - Fork(), exec(), ...
- **Memory management related**
 - Mmap(), shm_open(), ...
- **Device Management**
- **Information maintainance – time,date**
- **Communication between processes (IPC)**
- **Read man syscalls**

https://linuxhint.com/list_of_linux_syscalls/

EXAMPLES OF WINDOWS AND UNIX SYSTEM CALLS

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shm_open() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

```
int main() {
    int a = 2;
    printf("hi\n");
}
```

C Library

```
int printf("void *a, ...) {
    ...
    write(1, a, ...);
}

int write(int fd, char *, int len) {
    int ret;
    ...
    mov $5, %eax,
    mov ... %ebx,
    mov ..., %ecx
    int $0x80
    __asm__("movl %eax, -4(%ebp)");
# -4ebp is ret
    return ret;
}
```

Code schematic

-----user-kernel-mode-
boundary----

//OS code

```
int sys_write(int fd, char *, int
len) {
    figure out location on disk
    where to do the write and
    carry out the operation,
    etc.
}
```

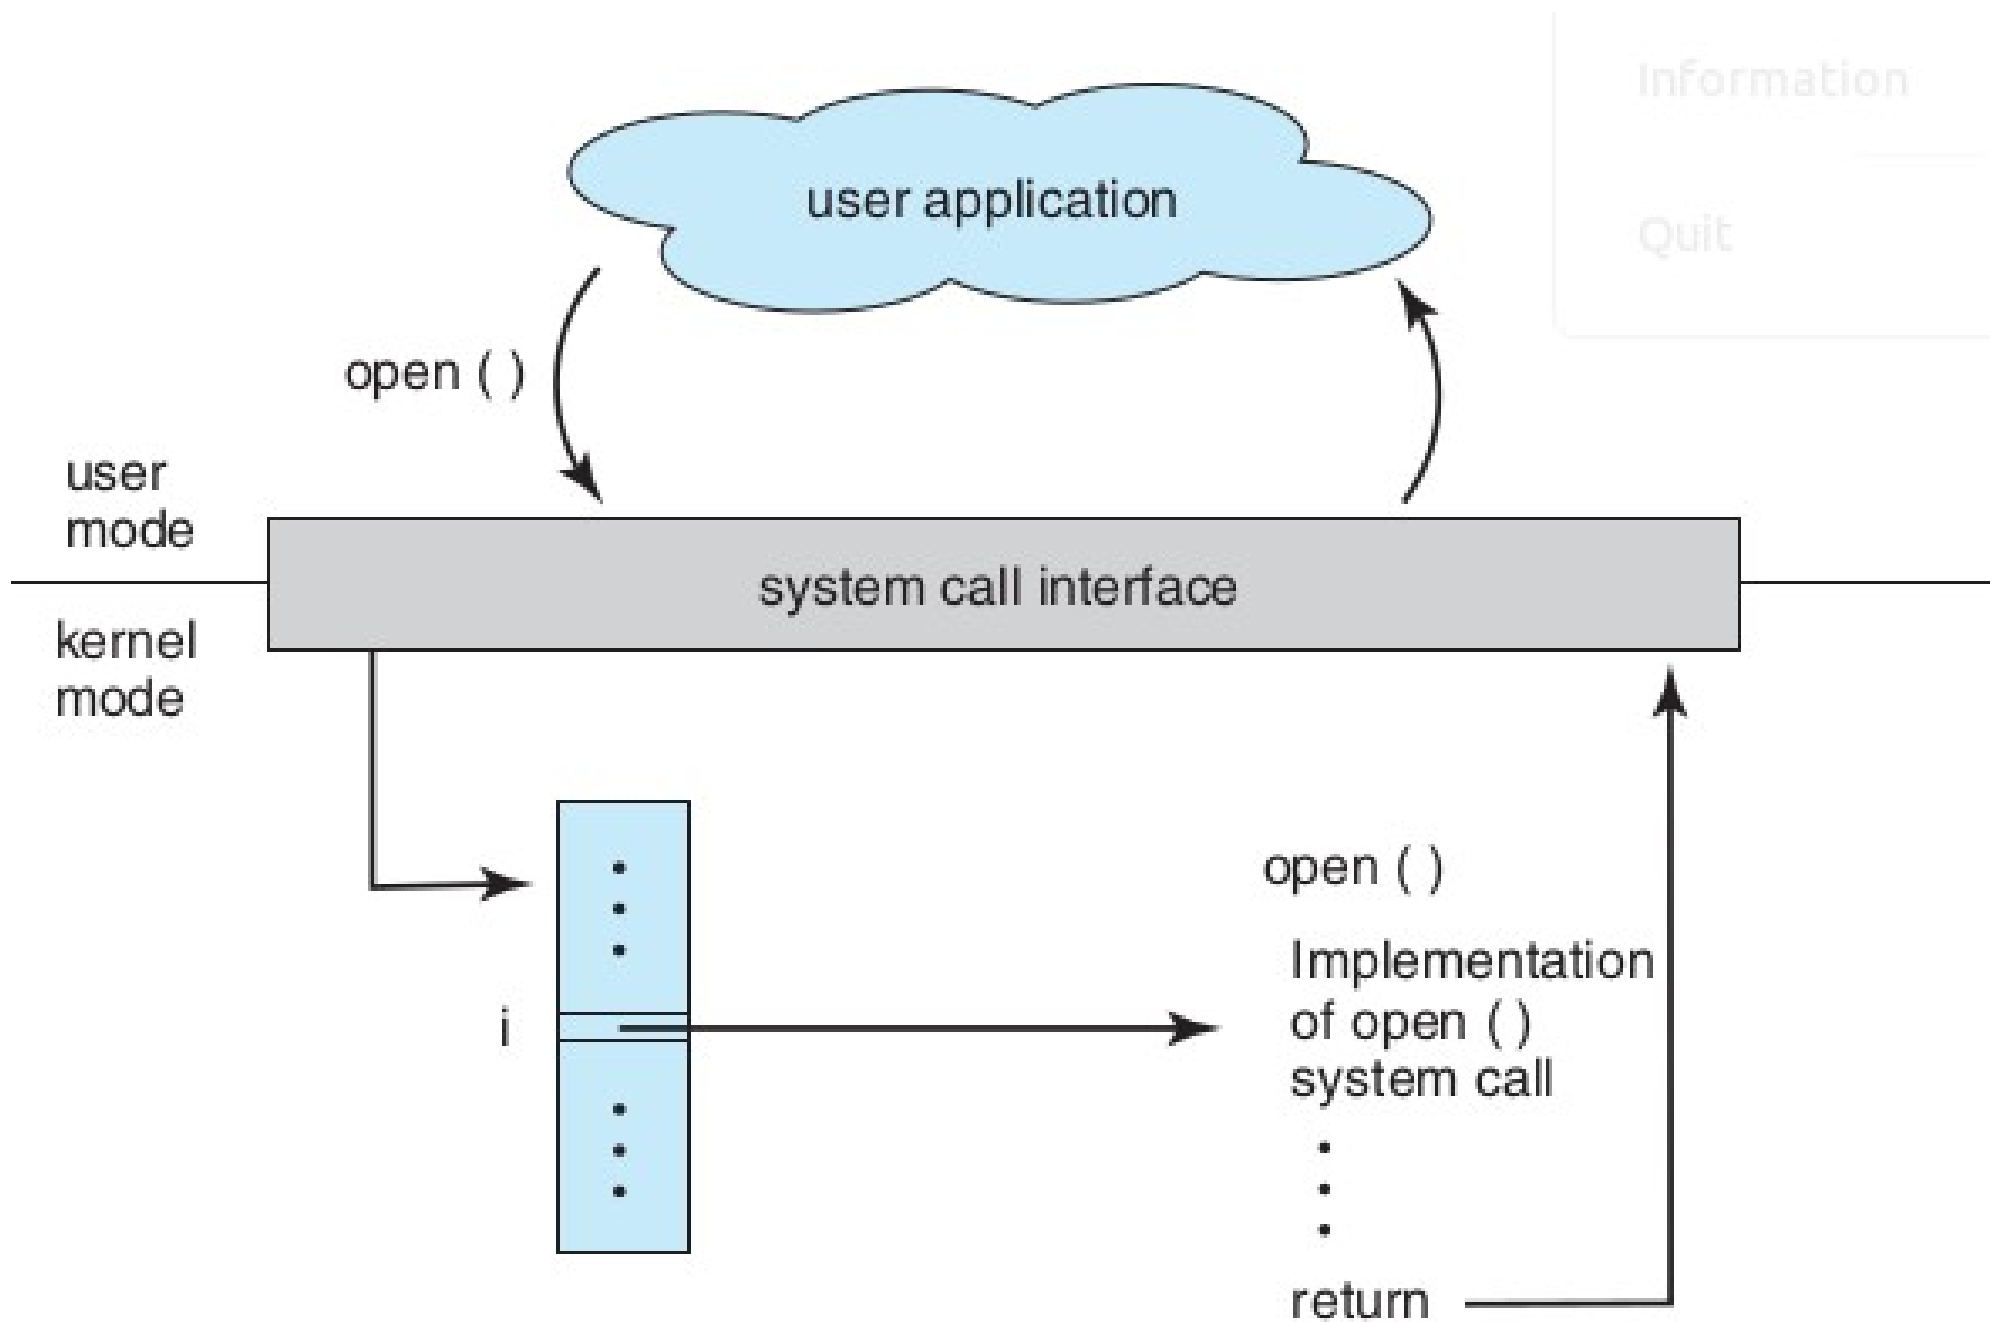


Figure 2.6 The handling of a user application invoking the `open ()` system call.

Two important system calls
Related to processes

fork() and exec()

Process

- **A program in execution**
- **Exists in RAM**
- **Scheduled by OS**
 - In a timesharing system, intermittantly scheduled by allocating a time quantum, e.g. 20 microseconds
- **The “ps” command on Linux**

Process in RAM

- **Memory is required to store the following components of a process**
 - **Code**
 - **Global variables (data)**
 - **Stack (stores local variables of functions)**
 - **Heap (stores malloced memory)**
 - **Shared libraries (e.g. code of printf, etc)**
 - **Few other things, may be**

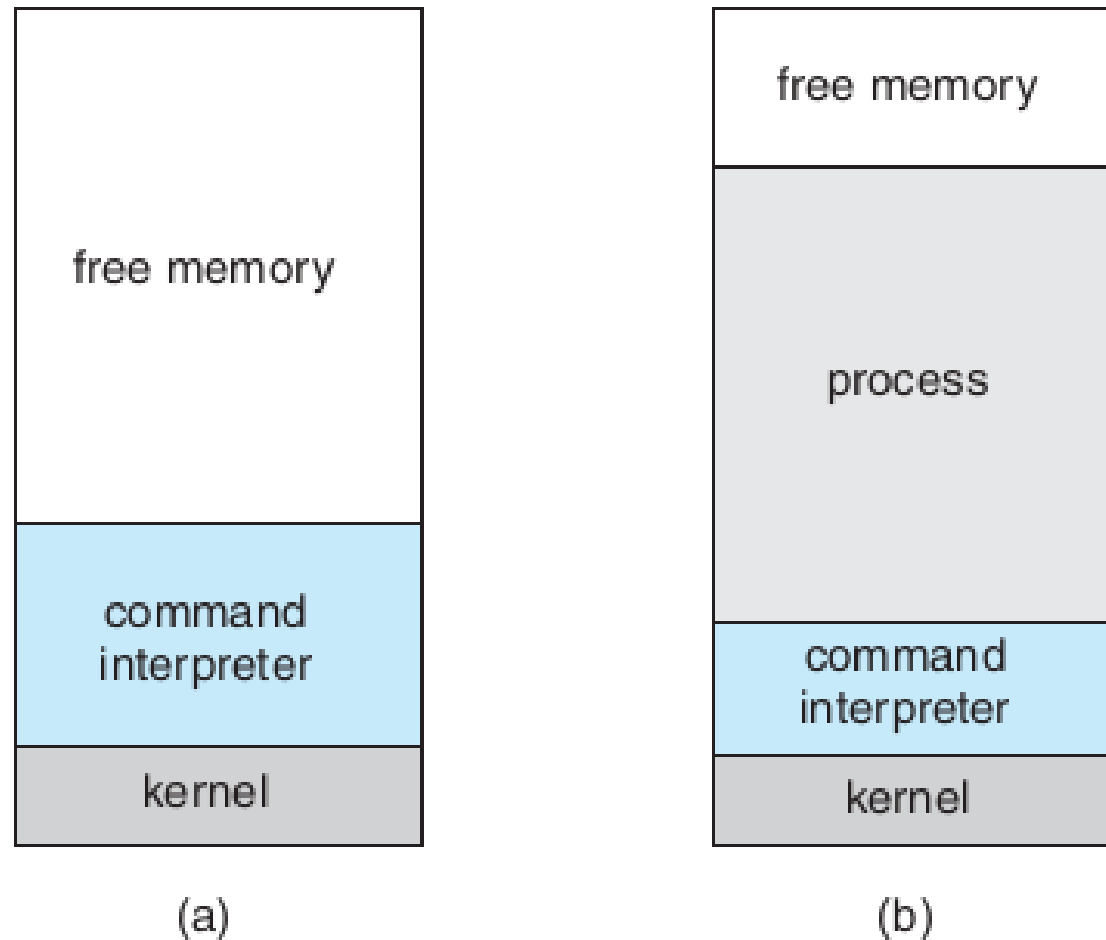
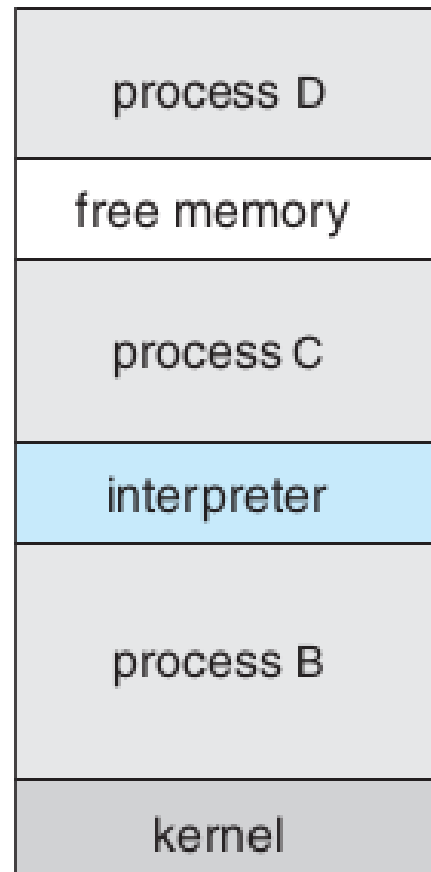


Figure 2.9 MS-DOS execution. (a) At system startup. (b) Running a program.

MS-DOS: a single tasking operating system
Only one program in RAM at a time, and only one program can run at a time



A multi tasking system
With multiple programs loaded in memory
Along with kernel
(A very simplified conceptual diagram. Things are more complex in reality)

fork()

- **A running process creates it's duplicate!**
- **After call to fork() is over**
 - **Two processes are running**
 - **Identical**
 - **The calling function returns in two places!**
 - **Caller is called parent, and the new process is called child**
 - **PID is returned to parent and 0 to child**

exec()

- **Variants: execvp(), execl(), etc.**
- **Takes the path name of an executable as an argument**
- **Overwrites the existing process using the code provided in the executable**
- **The original process is OVER ! Vanished!**
- **The new program starts running overwriting the existing process!**

Shell using fork and exec

- Demo
- The only way a process can be created on Unix/Linux is using `fork()` + `exec()`
- All processes that you see were started by some other process using `fork()` + `exec()` , except the initial *“init”* process
- *When you click on “firefox” icon, the user-interface program does a `fork()` + `exec()` to start firefox; same with a command line shell program*
- The “bash” shell you have been using is nothing but an advanced version of the shell code shown during the demo
- See the process tree starting from “init”
- Your next assignment

The boot process, once again

- BIOS
- Boot loader
- OS – kernel
- Init created by kernel by Hand(kernel mode)
- Kernel schedules init (the only process)
- Init fork-execs some programs (user mode)
 - Now these programs will be scheduled by OS
- Init -> GUI -> terminal -> shell
 - One of the typical parent-child relationships