

Notes on reading xv6 code

Abhijit A. M.

abhijit.comp@coep.ac.in

Credits:

xv6 book by Cox, Kaashoek, Morris

Notes by Prof. Sorav Bansal

Introduction to xv6

Structure of xv6 code

Compiling and executing xv6 code

About xv6

- Unix Like OS
- Multi tasking, Single user
- On x86 processor
- Supports some system calls
- Small code, 7 to 10k
- Meant for learning OS concepts
- **No** : demand paging, no copy-on-write fork, no shared-memory, fixed size stack for user programs

Use cscope and ctags with VIM

- Go to folder of xv6 code and run
`cscope -q *. [chS]`
- Also run
`ctags *. [chS]`
- Now download the file
http://cscope.sourceforge.net/cscope_maps.vim
as `.cscope_maps.vim` in your `~` folder
- And add line "`source ~/.cscope_maps.vim`" in your
`~/.vimrc` file
- Read this tutorial
http://cscope.sourceforge.net/cscope_vim_tutorial.html

Use call graphs (using doxygen)

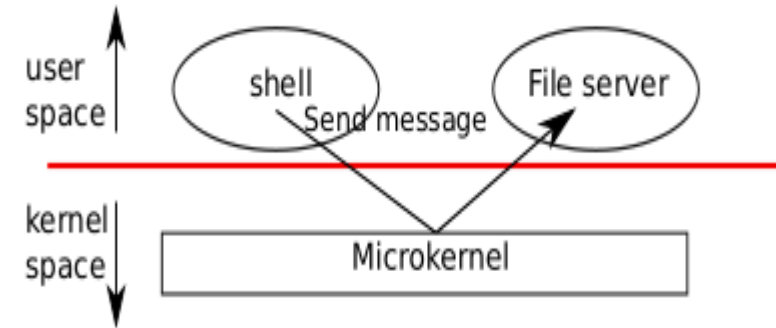
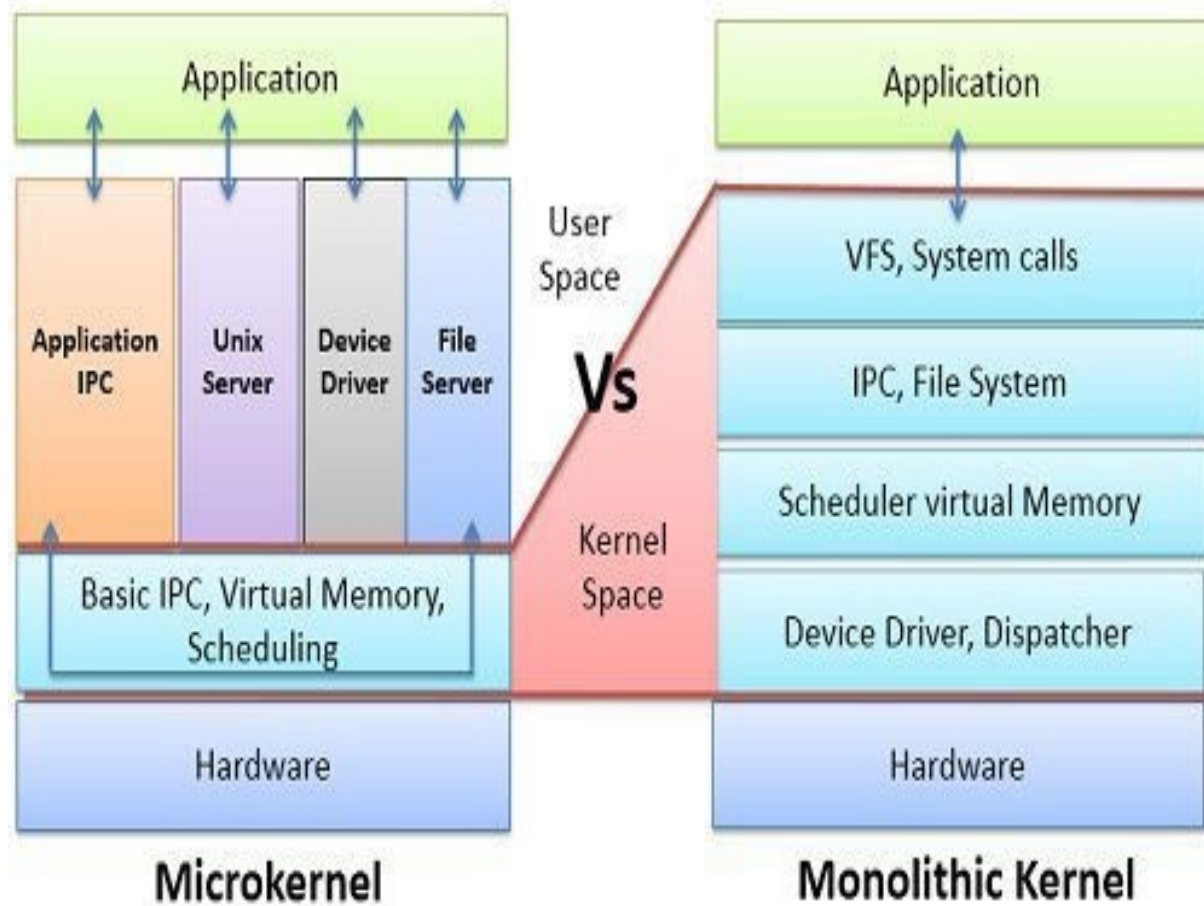
- **Doxygen – a documentation generator.**
- **Can also be used to generate “call graphs” of functions**
- **Download xv6**
- **Install doxygen on your Ubuntu machine.**
- **cd to xv6 folder**
- **Run “doxygen -g doxyconfig”**
 - **This creates the file “doxyconfig”**

Use call graphs (using doxygen)

- Create a folder “doxygen”
- Open “doxyconfig” file and make these changes.

```
PROJECT_NAME           = "XV6"
OUTPUT_DIRECTORY       = ./doxygen
CREATE_SUBDIRS         = YES
EXTRACT_ALL            = YES
EXCLUDE                = usertests.c cat.c yes.c echo.c
                        forktest.c grep.c init.c kill.c ln.c ls.c mkdir.c rm.c sh.c
                        stressfs.c wc.c zombie.c
CALL_GRAPH             = YES
CALLER_GRAPH           = YES
```
- Now run “doxygen doxyconfig”
- Go to “doxygen”/html and open “firefox index.html” --> See call graphs in files -> any file

Xv6 follows monolithic kernel approach



qemu

- A virtual machine manager, like Virtualbox
- Qemu provides us
 - BIOS
 - Virtual CPU, RAM, Disk controller, Keyboard controller
 - IOAPIC, LAPIC
- Qemu runs xv6 using this command

```
qemu -serial mon:stdio -drive
file=fs.img,index=1,media=disk,format=raw -drive
file=xv6.img,index=0,media=disk,format=raw -smp 2 -
m 512
```
- Invoked when you run “make qemu”

qemu

- **Understanding qemu command**
 - **-serial mon:stdio**
 - the window of xv6 is also multiplexed in your normal terminal.
 - Run “make qemu”, then Press “Ctrl-a” and “c” in terminal and you get qemu prompt
 - **-drive file=fs.img,index=1,media=disk,format=raw**
 - Specify the hard disk in “fs.img”, accessible at first slot in IDE(or SATA, etc), as a “disk” , with “raw” format
 - **-smp 2**
 - Two cores in SMP mode to be simulated
 - **-m 512**
 - Use 512 MB ram

About files in XV6 code

- `cat.c` `echo.c` `forktest.c` `grep.c`
`init.c` `kill.c` `ln.c` `ls.c` `mkdir.c`
`rm.c` `sh.c` `stressfs.c` `usertests.c`
`wc.c` `yes.c` `zombie.c`
 - User programs for testing xv6
- **Makefile**
 - To compile the code
- **dot-bochsrc**
 - For running with emulator bochs

About files in XV6 code

- `bootasm.S` `entryother.S` `entry.S`
`initcode.S` `swtch.S` `trapasm.S`
`usys.S`
 - Kernel code written in Assembly. Total 373 lines
- `kernel.ld`
 - Instructions to Linker, for linking the kernel properly
- `README` `Notes` `LICENSE`
 - Misc files

Using Makefile

- **make qemu**
 - Compile code and run using “qemu” emulator
- **make xv6.pdf**
 - Generate a PDF of xv6 code
- **make mkfs**
 - Create the mkfs program
- **make clean**
 - Remove all intermediary and final build files

Files generated by Makefile

- **.o files**

- Compiled from each .c file
- No need of separate instruction in Makefile to create .o files
- `_%: %.o $(ULIB)` line is sufficient to build each .o for a `_xyz` file

-

Files generated by Makefile

- **asm files**

- Each of them has an equivalent object code file or C file. For example

```
bootblock: bootasm.S bootmain.c
```

```
    $(CC) $(CFLAGS) -fno-pic -O -nostdinc -I. -c  
bootmain.c
```

```
    $(CC) $(CFLAGS) -fno-pic -nostdinc -I. -c  
bootasm.S
```

```
    $(LD) $(LDFLAGS) -N -e start -Ttext 0x7C00 -o  
bootblock.o bootasm.o bootmain.o
```

```
    $(OBJDUMP) -S bootblock.o > bootblock.asm
```

```
    $(OBJCOPY) -S -O binary -j .text bootblock.o  
bootblock
```

```
    ./sign.pl bootblock
```

Files generated by Makefile

- `_ln, _ls, etc`
 - Executable user programs
 - Compilation process is explained after few slides

Files generated by Makefile

- **xv6.img**

- Image of xv6 created

`xv6.img: bootblock kernel`

```
        dd if=/dev/zero of=xv6.img  
count=10000
```

```
        dd if=bootblock of=xv6.img  
conv=notrunc
```

```
        dd if=kernel of=xv6.img seek=1  
conv=notrunc
```


Files generated by Makefile

- **bootblock**

```
bootblock: bootasm.S bootmain.c
```

```
    $(CC) $(CFLAGS) -fno-pic -O -nostdinc -I.  
-c bootmain.c
```

```
    $(CC) $(CFLAGS) -fno-pic -nostdinc -I. -c  
bootasm.S
```

```
    $(LD) $(LDFLAGS) -N -e start -Ttext  
0x7C00 -o bootblock.o bootasm.o bootmain.o
```

```
    $(OBJDUMP) -S bootblock.o > bootblock.asm
```

```
    $(OBJCOPY) -S -O binary -j .text  
bootblock.o bootblock
```

```
    ./sign.pl bootblock
```

Files generated by Makefile

kernel

```
kernel: $(OBJS) entry.o entryother initcode  
kernel.ld
```

```
$(LD) $(LDFLAGS) -T kernel.ld -  
o kernel entry.o $(OBJS) -b binary  
initcode entryother
```

```
$(OBJDUMP) -S kernel >  
kernel.asm
```

```
$(OBJDUMP) -t kernel | sed  
'1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d'  
> kernel.sym
```

Files generated by Makefile

- **fs.img**

- A disk image containing user programs and README

```
fs.img: mkfs README $(UPROGS)
```

```
./mkfs fs.img README $(UPROGS)
```

- **.sym files**

- Symbol tables of different programs

- E.g. for file “kernel”

```
$(OBJDUMP) -t kernel | sed '1,/SYMBOL  
TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Size of xv6 C code

- **wc *[ch] | sort -n**
 - **10595 34249 278455 total**
 - **Out of which**
 - **738 4271 33514 dot-bochsrc**
- **wc cat.c echo.c forktest.c grep.c init.c kill.c
ln.c ls.c mkdir.c rm.c sh.c stressfs.c
usertests.c wc.c yes.c zombie.c**
 - **2849 6864 51993 total**
- **So total code is $10595 - 2849 - 738 = 7008$ lines**

List of commands to try (in given order)

usertests # Runs lot of tests and takes upto 10 minutes to run

stressfs # opens , reads and writes to files in parallel

ls # out put is filetype, inode number, type

cat README

ls;ls

cat README | grep BUILD

echo hi there

echo hi there | grep hi

echo "hi there

List of commands to try (in this order)

echo README | grep Wa

echo README | grep Wa |
grep ty # does not work

cat README | grep Wa |
grep bl # works

ls > out # takes time!

mkdir test

cd test

ls # fails

ls ../ # works from inside test

cd # fails

cd / # works

wc README

rm out

ls . test # listing both
directories

In cat xyz; ls

rm xyz; ls

User Libraries: Used to link user land programs

- **Ulib.c**
 - Strcpy, strcmp, strlen, memset, strchr, stat, atoi, memmove
 - Stat uses open()
- **Usys.S -> compiles into usys.o**
 - Assembly code file. Basically converts all calls like open() (e.g. used in ulib.c) into assembly code using “int” instruction.

Run following command see the last 4 lines in the output

```
objdump -d usys.o
```

```
00000048 <open>:
```

| | | | |
|-----|-------------------|-----|-------------|
| 48: | b8 0f 00 00 00 00 | mov | \$0xf, %eax |
| 4d: | cd 40 | int | \$0x40 |
| 4f: | c3 | ret | |

User Libraries: Used to link user land programs

- **printf.c**
 - Code for printf()!
 - Interesting to read this code.
 - Uses variable number of arguments. Normal technique in C is to use va_args library, but here it uses pointer arithmetic.
 - Written using two more functions: printint() and putc() - both call write()
 - Where is code for write()?

User Libraries: Used to link user land programs

- **umalloc.c**
 - This is an implementation of malloc() and free()
 - Almost same as the one done in “The C Programming Language” by Kernighan and Ritchie
 - Uses sbrk() to get more memory from xv6 kernel

Understanding the build process in more details

- **Run**

`make qemu | tee make-output.txt`

- **You will get all compilation commands in
make-output.txt**

Compiling user land programs

Normally when you compile a program on Linux

You compile it for the same 'target' machine (= CPU + OS)

The compiler itself runs on the same OS

To compile a user land program for xv6, we don't have a compiler on xv6,

So we compile the programs (using make, cc) on Linux , for xv6

Obviously they can't link with the standard libraries on Linux

Compiling user land programs

```
ULIB = ulib.o usys.o printf.o umalloc.o
```

```
_%: %.o $(ULIB)
```

```
$(LD) $(LDFLAGS) -N -e main -Ttext 0 -o $@ $^
```

```
$(OBJDUMP) -S $@ > $*.asm
```

```
$(OBJDUMP) -t $@ | sed '1,/SYMBOL TABLE/d;  
s/ .* / /; /^$$/d' > $*.sym
```

`$@` is the name of the file being generated

`$^` is dependencies . i.e. `$(ULIB)` and `%.o` in this case

Compiling user land programs

```
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing  
-O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-  
pointer -fno-stack-protector -fno-pie -no-pie -c -o  
cat.o cat.c
```

```
ld -m elf_i386 -N -e main -Ttext 0 -o _cat cat.o  
ulib.o usys.o printf.o umalloc.o
```

```
objdump -S _cat > cat.asm
```

```
objdump -t _cat | sed '1,/SYMBOL TABLE/d;  
s/ .* / /; /^$/d' > cat.sym
```

Compiling user land programs

Mkfs is compiled like a Linux program !

gcc -Werror -Wall -o mkfs mkfs.c

How to read kernel code ?

- **Understand the data structures**
 - Know each global variable, typedefs, lists, arrays, etc.
 - Know the purpose of each of them
- **While reading a code path, e.g. `exec()`**
 - Try to 'locate' the key line of code that does major work
 - Initially (but not forever) ignore the 'error checking' code
- **Keep summarising what you have read**
 - Remembering is important !
- **To understand kernel code, you should be good with concepts in OS , C, assembly, hardware**

Pre-requisites for reading the code

- **Understanding of core concepts of operating systems**
 - Memory Management, processes, fork-exec, file systems, synchronization, x86 architecture, calling convention , computer organization
- **2 approaches:**
 - 1) Read OS basics first, and then start reading xv6 code
 - **Good approach, but takes more time !**
 - 2) Read some basics, read xv6, repeat
 - **Gives a headstart, but you will always have gaps in your understanding of the code, until you are done with everything**
 - **We normally follow this approach**
- Good knowledge of C, pointers, **function pointers particularly**
 - Data structures: doubly linked lists, queues, structures and pointers