

Processes

Abhijit A M

abhijit.comp@coep.ac.in

Process related data structures in kernel code

- Kernel needs to maintain following types of data structures for managing processes
 - List of all processes
 - Memory management details for each, files opened by each etc.
 - Scheduling information about the process
 - Status of the process
 - List of processes “waiting” for different events to occur,
 - Etc.

Process Control Block

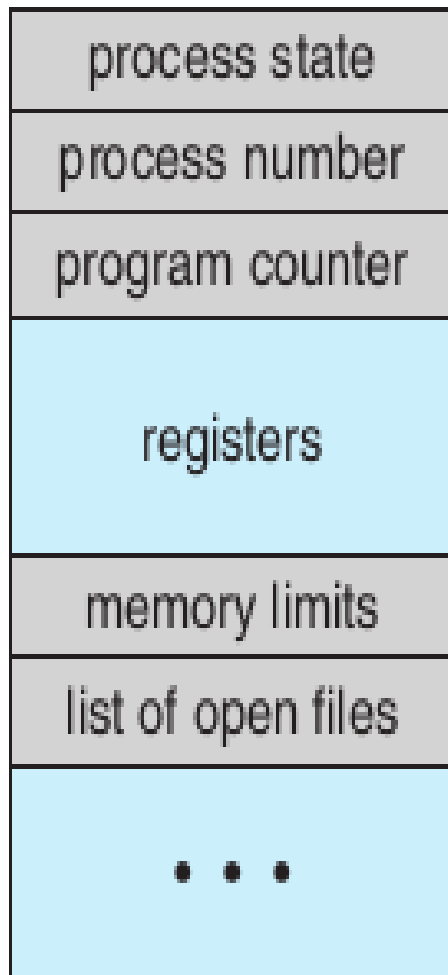


Figure 3.3 Process control block (PCB).

- A record representing a process in operating system's data structures
- OS maintains a “list” of PCBs, one for each process
- Called “struct task_struct” in Linux kernel code and “struct proc” in xv6 code

Fields in PCB

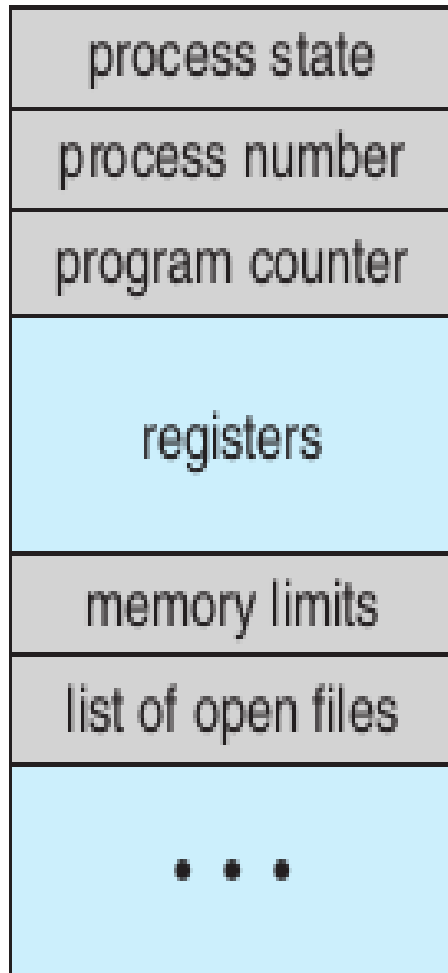


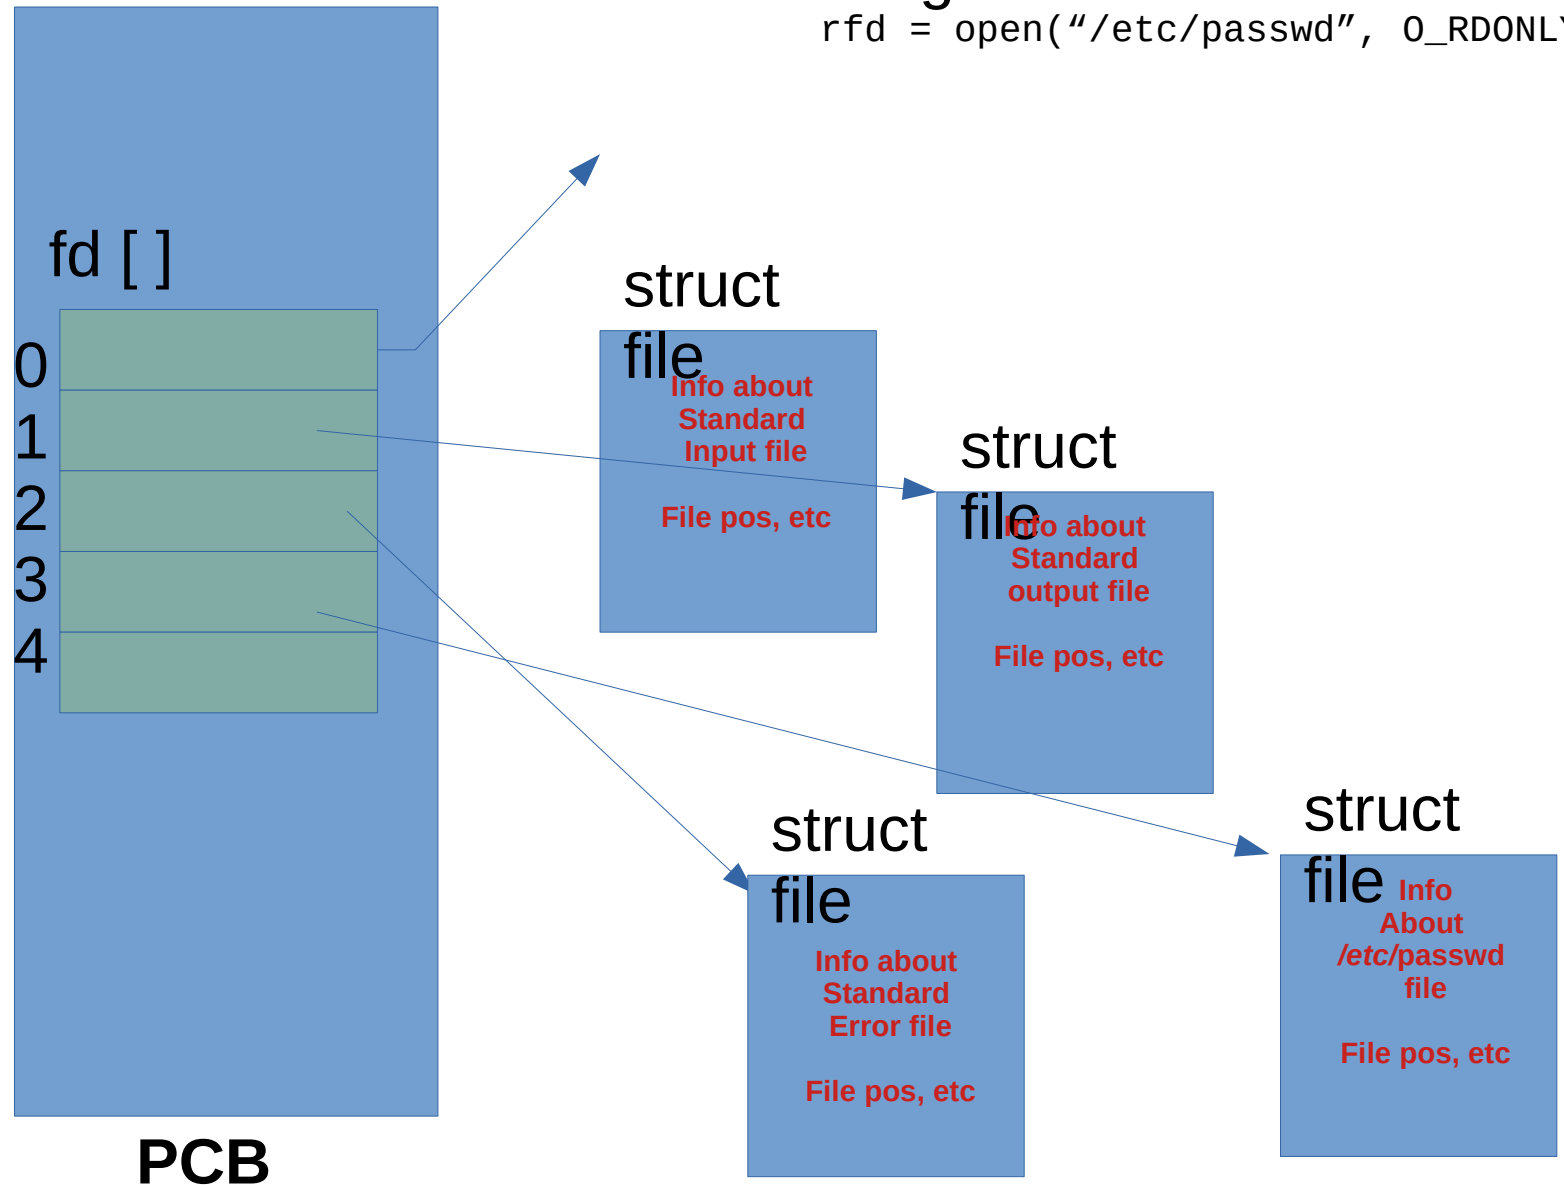
Figure 3.3 Process control block (PCB).

- Process ID (PID)
- Process State
- Program counter
- Registers
- Memory limits of the process
- Accounting information
- I/O status
- Scheduling information
- array of file descriptors (list of open files)
- ...etc

List of open files

Program did

```
rfd = open("/etc/passwd", O_RDONLY)
```



List of open files

- The PCB contains an array of pointers, called file descriptor array (fd[]), pointers to structures representing files
- When open() system call is made
 - A new file structure is created and relevant information is stored in it
 - Smallest available of fd [] pointers is made to point to this new struct file
 - The index of this fd [] pointer is returned by open
- When subsequent calls are made to read(fd,) or write(fd, ...) , etc.
 - The kernel gets the “fd” as an index in the fd[] array and is able to locate the file structure for that file

// XV6 Code : Per-process state

```
enum procstate { UNUSED, EMBRYO, SLEEPING,  
RUNNABLE, RUNNING, ZOMBIE };
```

```
struct proc {
```

```
    uint sz;                // Size of process memory (bytes)
```

```
    pde_t* pgdir;           // Page table
```

```
    char *kstack;           // Bottom of kernel stack for this
```

```
process
```

```
    enum procstate state;    // Process state
```

```
    int pid;                 // Process ID
```

```
    struct proc *parent;     // Parent process
```

```
    struct trapframe *tf;    // Trap frame for current syscall
```

```
    struct context *context; // swtch() here to run process
```

```
    void *chan;              // If non-zero, sleeping on chan
```

```
    int killed;              // If non-zero, have been killed
```

```
    struct file *ofile[NOFILE]; // Open files
```

```
    struct inode *cwd;        // Current directory
```

```
    char name[16];           // Process name (debugging)
```

```
};
```

```
struct {
```

```
    struct spinlock lock;
```

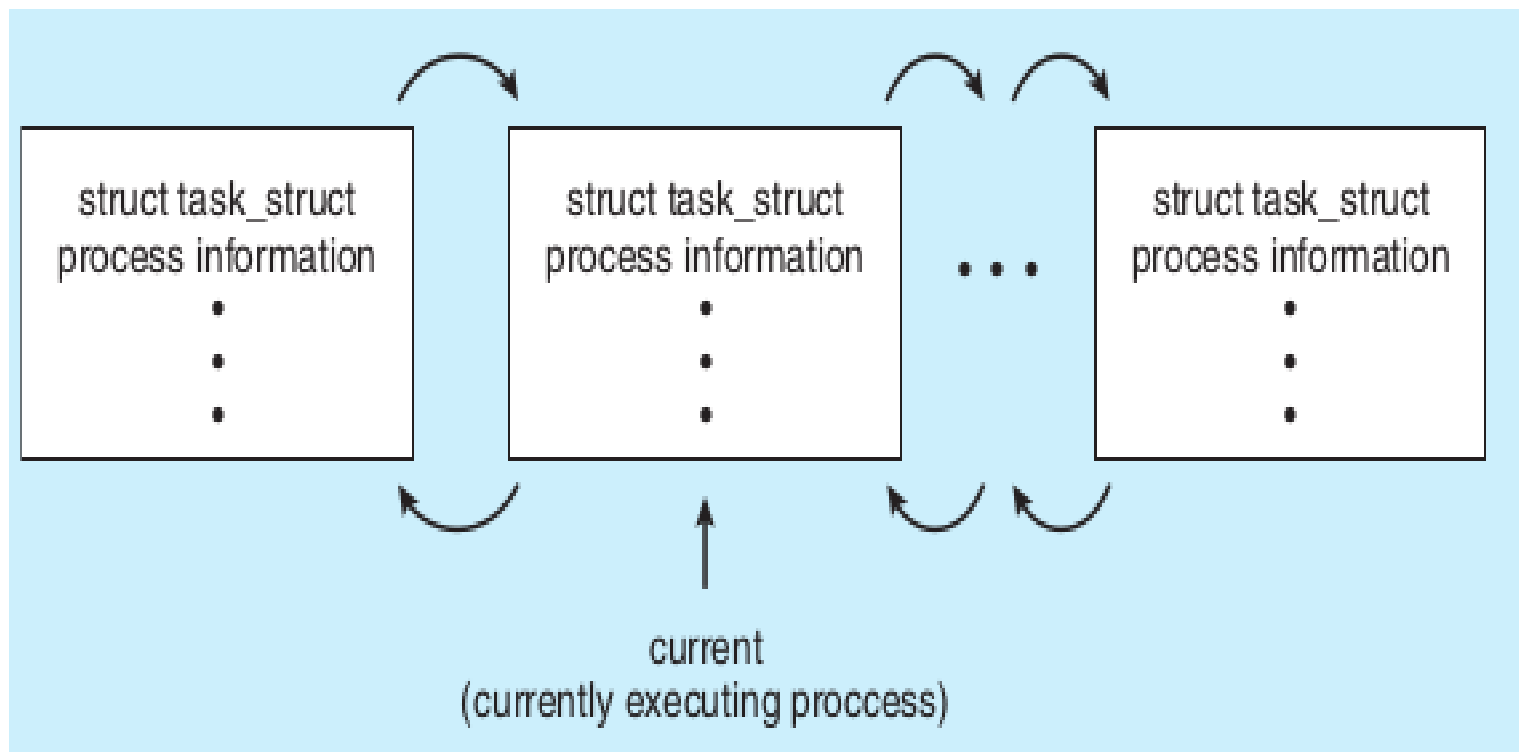
```
    struct proc proc[NPROC];
```

```
} ptable;
```

```
struct file {  
    enum { FD_NONE,  
FD_PIPE, FD_INODE } type;  
    int ref; // reference count  
    char readable;  
    char writable;  
    struct pipe *pipe;  
    struct inode *ip;  
    uint off;  
};
```

Process Queues/Lists inside OS

- Different types of queues/lists can be maintained by OS for the processes
 - A queue of processes which need to be scheduled
 - A queue of processes which have requested input/output to a device and hence need to be put on hold/wait
 - List of processes currently running on multiple CPUs
 - Etc.



// Linux data structure

```
struct task_struct {
```

```
    long state; /*state of the process */
```

```
    struct sched_entity se; /* scheduling information */
```

```
    struct task_struct *parent; /*this process's parent */
```

```
    struct list_head children; /*this process's children */
```

```
    struct files_struct *files; /* list of open files */
```

```
    struct mm_struct *mm; /*address space */
```

```
struct list_head {  
    struct list_head  
    *next, *prev;  
};
```




Figure 3.4 The ready queue and wait queues.

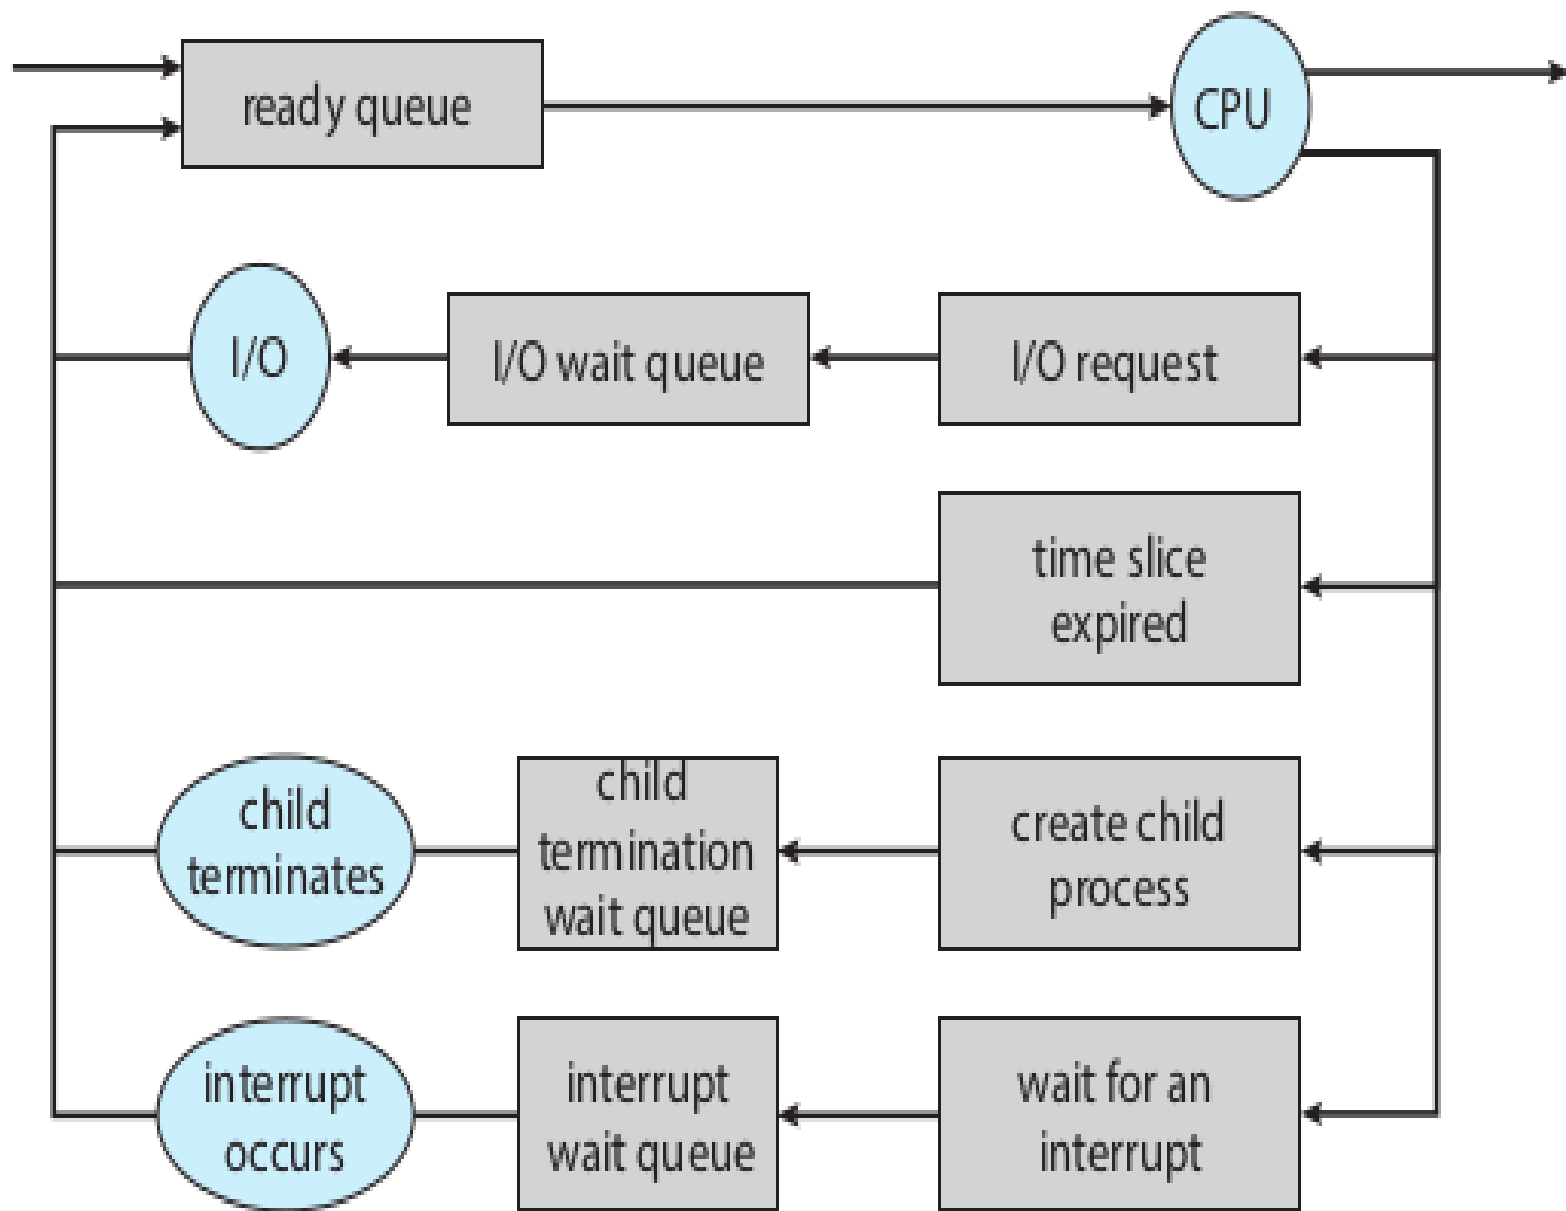


Figure 3.5 Queueing-diagram representation of process scheduling.

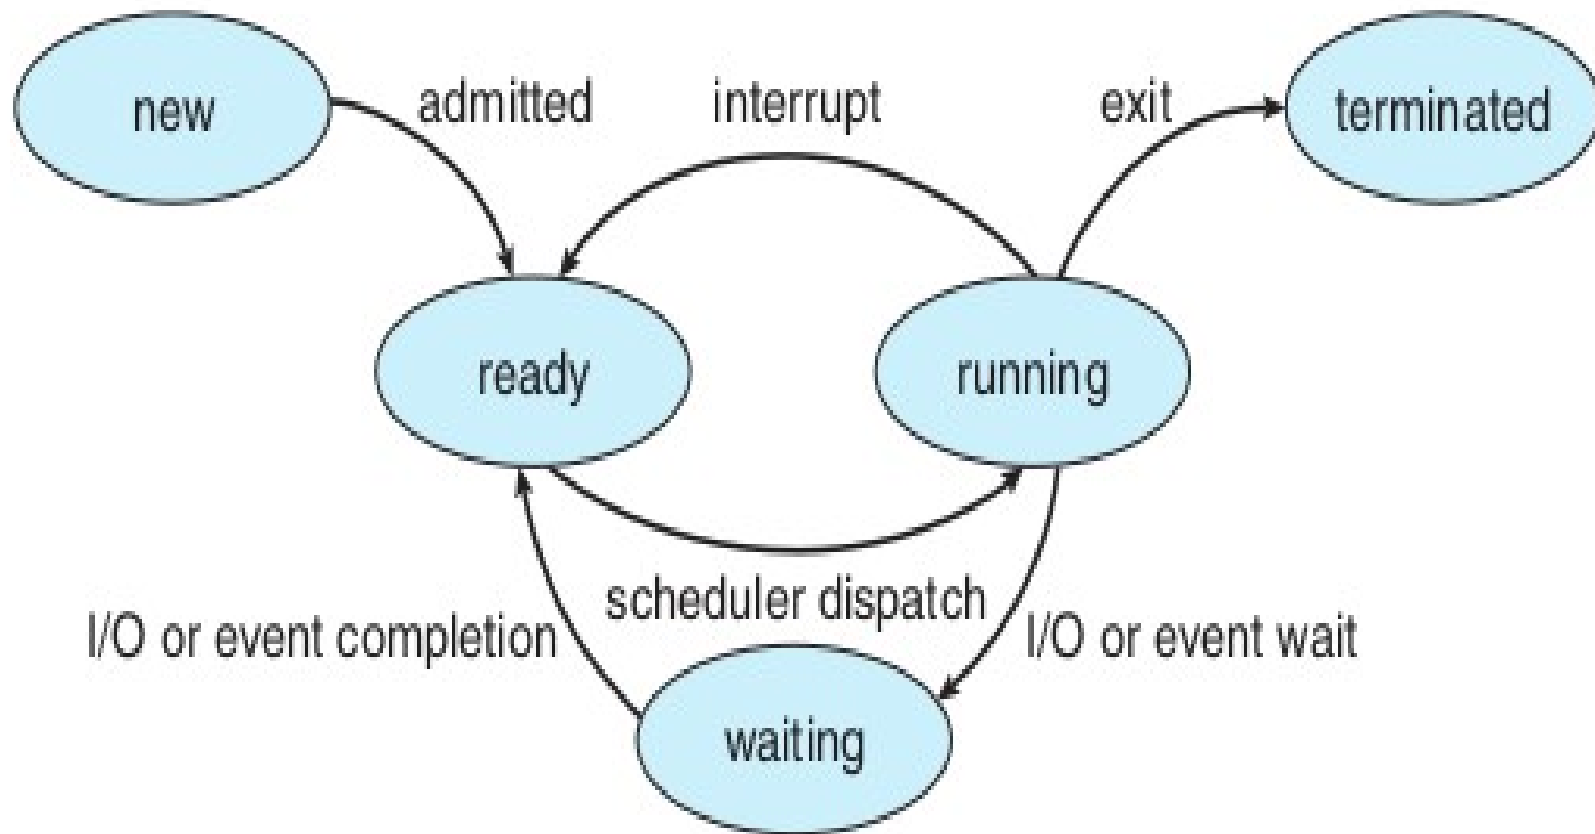


Figure 3.2 Diagram of process state.

Conceptual diagram

“Giving up” CPU by a process or blocking

```
int main() {
    i = j + k;
    scanf("%d", &k);
}

int scanf(char *x, ...) {
    ...
    read(0, ..., ...);
}

int read(int fd, char *buf, int len) {
    ...
    __asm__ { "int 0x80..." }
    ...
}
```

OS Syscall

```
sys_read(int fd, char *buf, int len) {
    file f = current->fdarray[fd];
    int offset = f->position;
    ...
    disk_read(... , offset, ...);
    // Do what now?
    //asynchronous read
    //Interrupt will occur when the disk read is
    complete
    // Move the process from ready queue to a
    wait queue and call scheduler!
    // This is called "blocking"
    Return the data read ;
}

disk_read(...., offset, .... ) {
    __asm__("outb PORT ..");
    return;
}
```

“Giving up” CPU by a process or blocking

The relevant code in xv6 is in

`Sleep()`

The wakeup code is in `wakeup()` and `wakeup1()`

To be seen later

Context Switch

- Context
 - Execution context of a process
 - CPU registers, process state, memory management information, all configurations of the CPU that are specific to execution of a process/kernel
- Context Switch
 - Change the context from one process/OS to OS/another process
 - Need to save the old context and load new context
 - Where to save? --> PCB of the process

Context Switch

- Is an overhead
- No useful work happening while doing a context switch
- Time can vary from hardware to hardware
- Special instructions may be available to save a set of registers in one go




Figure 3.6 Diagram showing context switch from process to process.

Pecularity of context switch

- When a process is running, the function calls work in LIFO fashion
 - Made possible due to calling convention
- When an interrupt occurs
 - It can occur anytime
 - Context switch can happen in the middle of execution of any function
- After context switch
 - One process takes place of another
 - This “switch” is obviously not going to happen using calling convention, as no “call” is happening
 - Code for context switch must be in assembly!

NEXT: XV6 code overview

1. Understanding how traps are handled
2. How timer interrupt goes to scheduler
3. How scheduling takes place
4. How a “blocking” system call (e.g. `read()`) “blocks”