

Creation of first process by kernel

Why first process needs 'special' treatment?

- Normally process is created using `fork()`
 - and typically followed by a call to `exec()`
- Fork will use the PCB of existing process to create a new process
 - as a clone
- The first process has nothing to copy from!
- So it's PCB needs to "built" by kernel code

Why first process needs 'special' treatment?

- **XV6 approach**
 - Create the process as if it was created by “fork”
 - Ensure that the process starts in a call to “exec”
 - Let “Exec” do the rest of the JOB as expected
 - In this case exec() will call
 - `exec(“/init”, NULL);`
- **See the code of init.c**
 - opens console() device for I/O; dups 0 on 1 and 2!
 - **Same device file for I/O**
 - forks a process and execs (“sh”) on it.
 - Itself keeps waiting for zombie processes

Why first process needs 'special' treatment?

- **What needs to be done ?**
 - Build struct proc by hand
 - How data structures (proc, stack, etc) are hand-crafted so that when kernel returns, the process starts in code of init

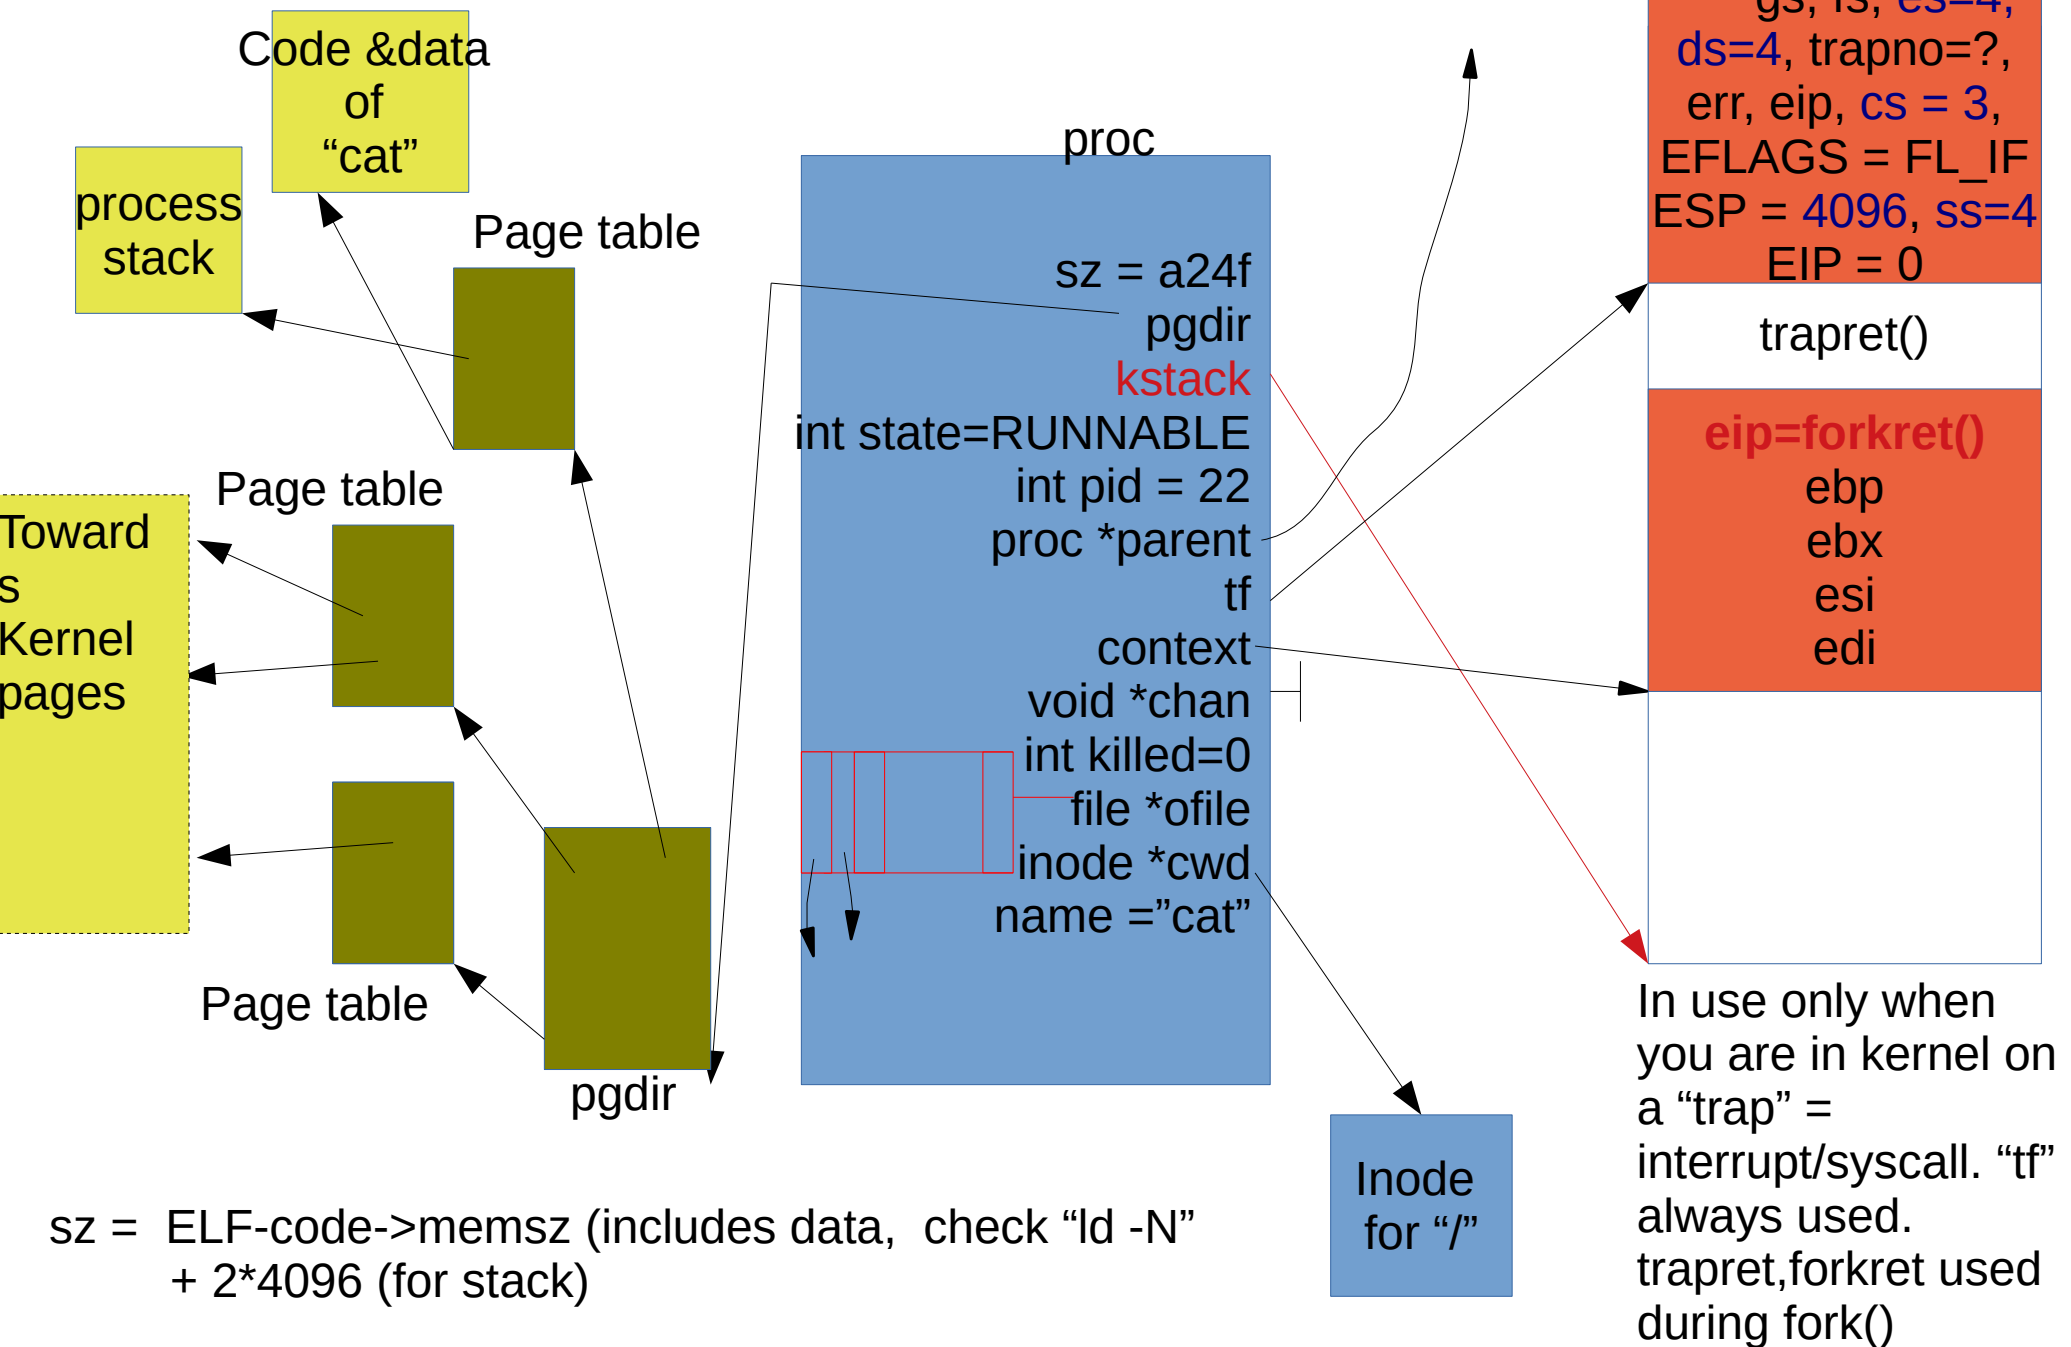
Imp Concepts

- **A process has two stacks**
 - user stack: used when user code is running
 - kernel stack: used when kernel is running on behalf of a process
- **Note: there is a third stack also!**
 - The kernel stack used by the scheduler itself
 - Not a per process stack

Imp Concepts

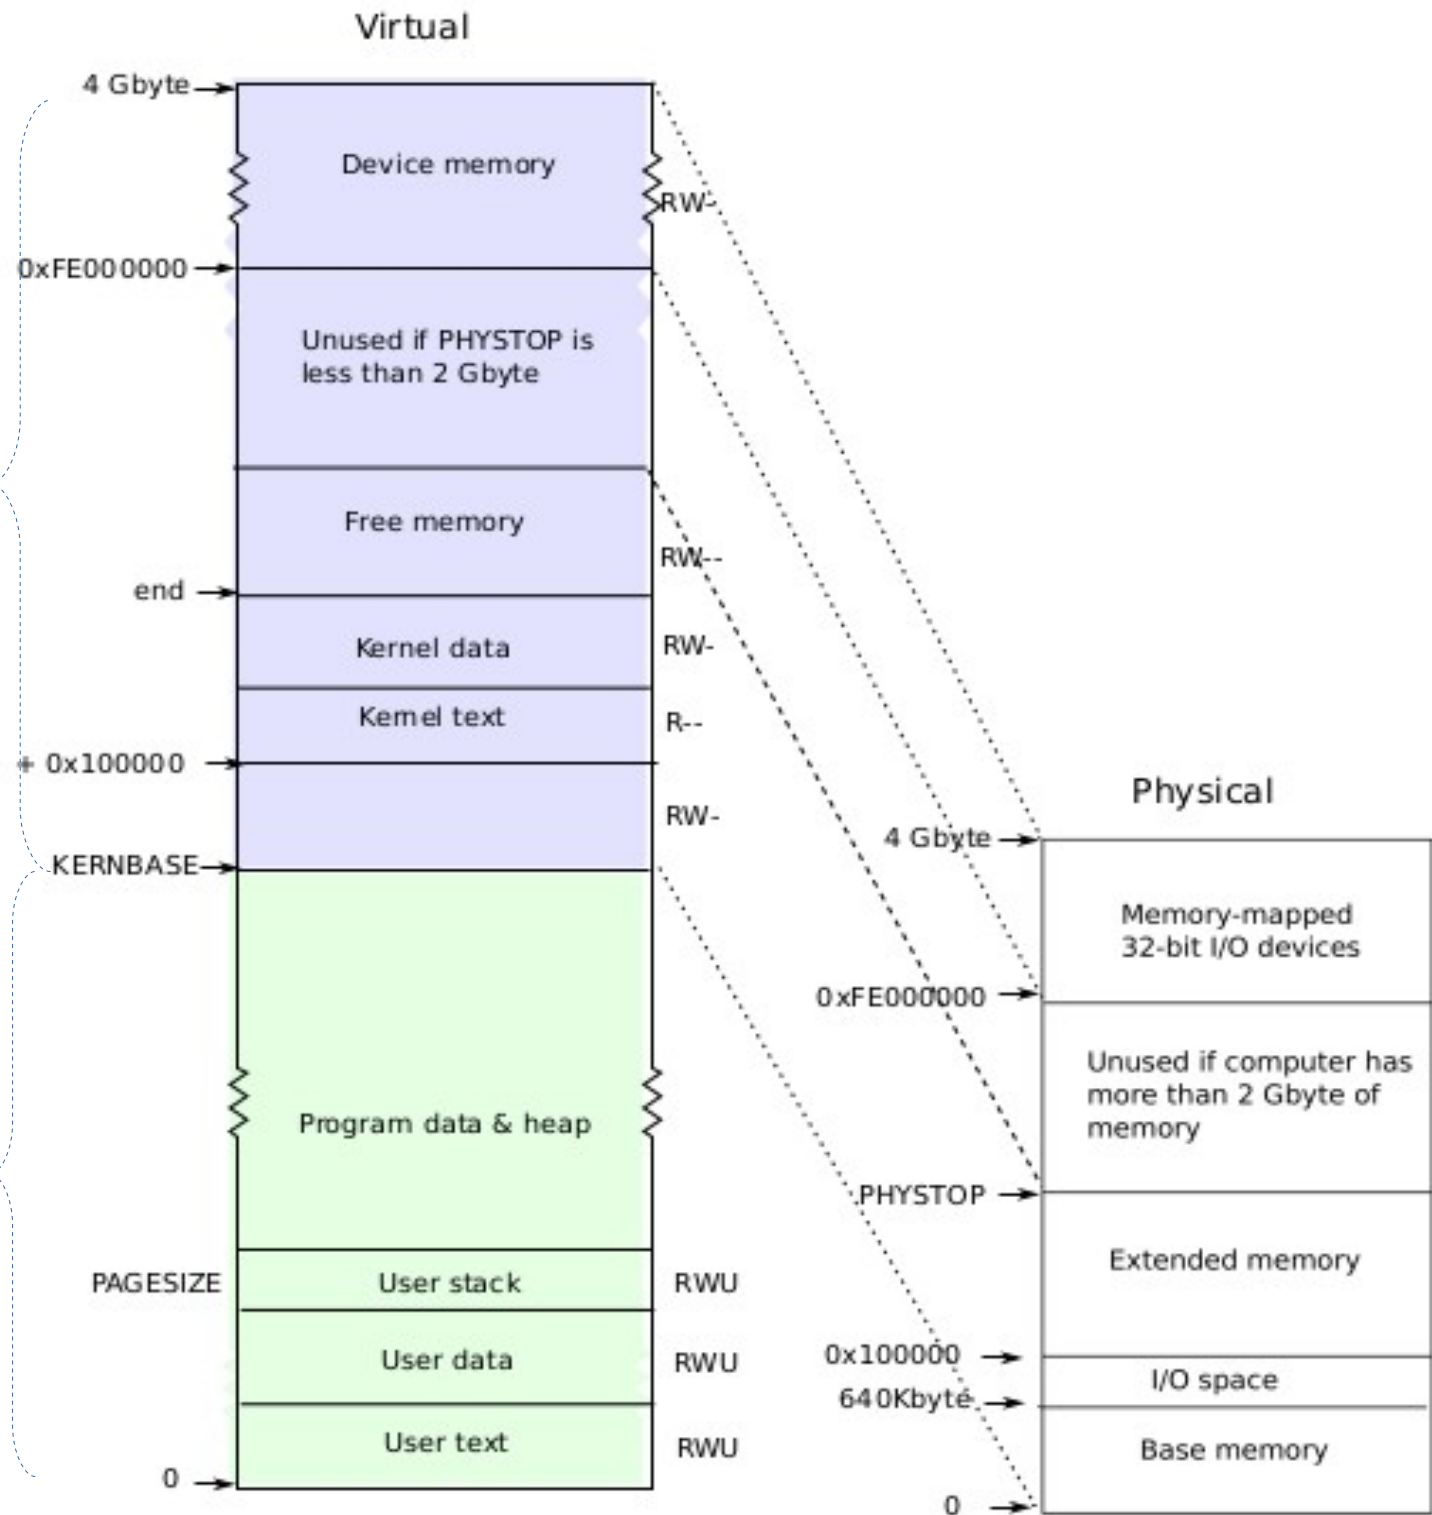
```
struct proc {  
    uint sz;                // Size of process memory (bytes)  
    pde_t* pgdir;           // Page table  
    char *kstack;           // Bottom of kernel stack for this process  
    enum procstate state;    // Process state  
    int pid;                // Process ID  
    struct proc *parent;     // Parent process  
    struct trapframe *tf;    // Trap frame for current syscall  
    struct context *context; // swtch() here to run process  
    void *chan;              // If non-zero, sleeping on chan  
    int killed;              // If non-zero, have been killed  
    struct file *ofile[NOFILE]; // Open files  
    struct inode *cwd;        // Current directory  
    char name[16];           // Process name (debugging)  
};
```

struct proc diagram: Very imp!



**setupkvm()
does this mapping**

**These mappings
need to be
created per
process**




Memory Layout of a user process After exec()

Note the argc, argv on stack

stack is just one page.

size of text and data is derived from ELF file



main()->userinit()

Creating first process by hand

- **Code of the first process**
 - **initcode.S and init.c**
 - **init.c is compiled into “/init” file**
 - **During make !**
 - **Trick:**
 - **Use initcode.S to “exec(“/init”)”**
 - **And let exec() do rest of the job**
 - **But before you do exec()**
 - **Process must exist as if it was forked() and running**

main()->userinit() Creating first process by hand

```
void
```

```
userinit(void)
```

```
{
```

```
    struct proc *p;
```

```
    extern char _binary_initcode_start[], _binary_initcode_size[];
```

```
    // Abhijit: obtain proc 'p', with stack initialized
```

```
    // and trapframe created and eip set to 'forkret'
```

```
    p = allocproc();
```

```
    // let's see what allocproc() does
```

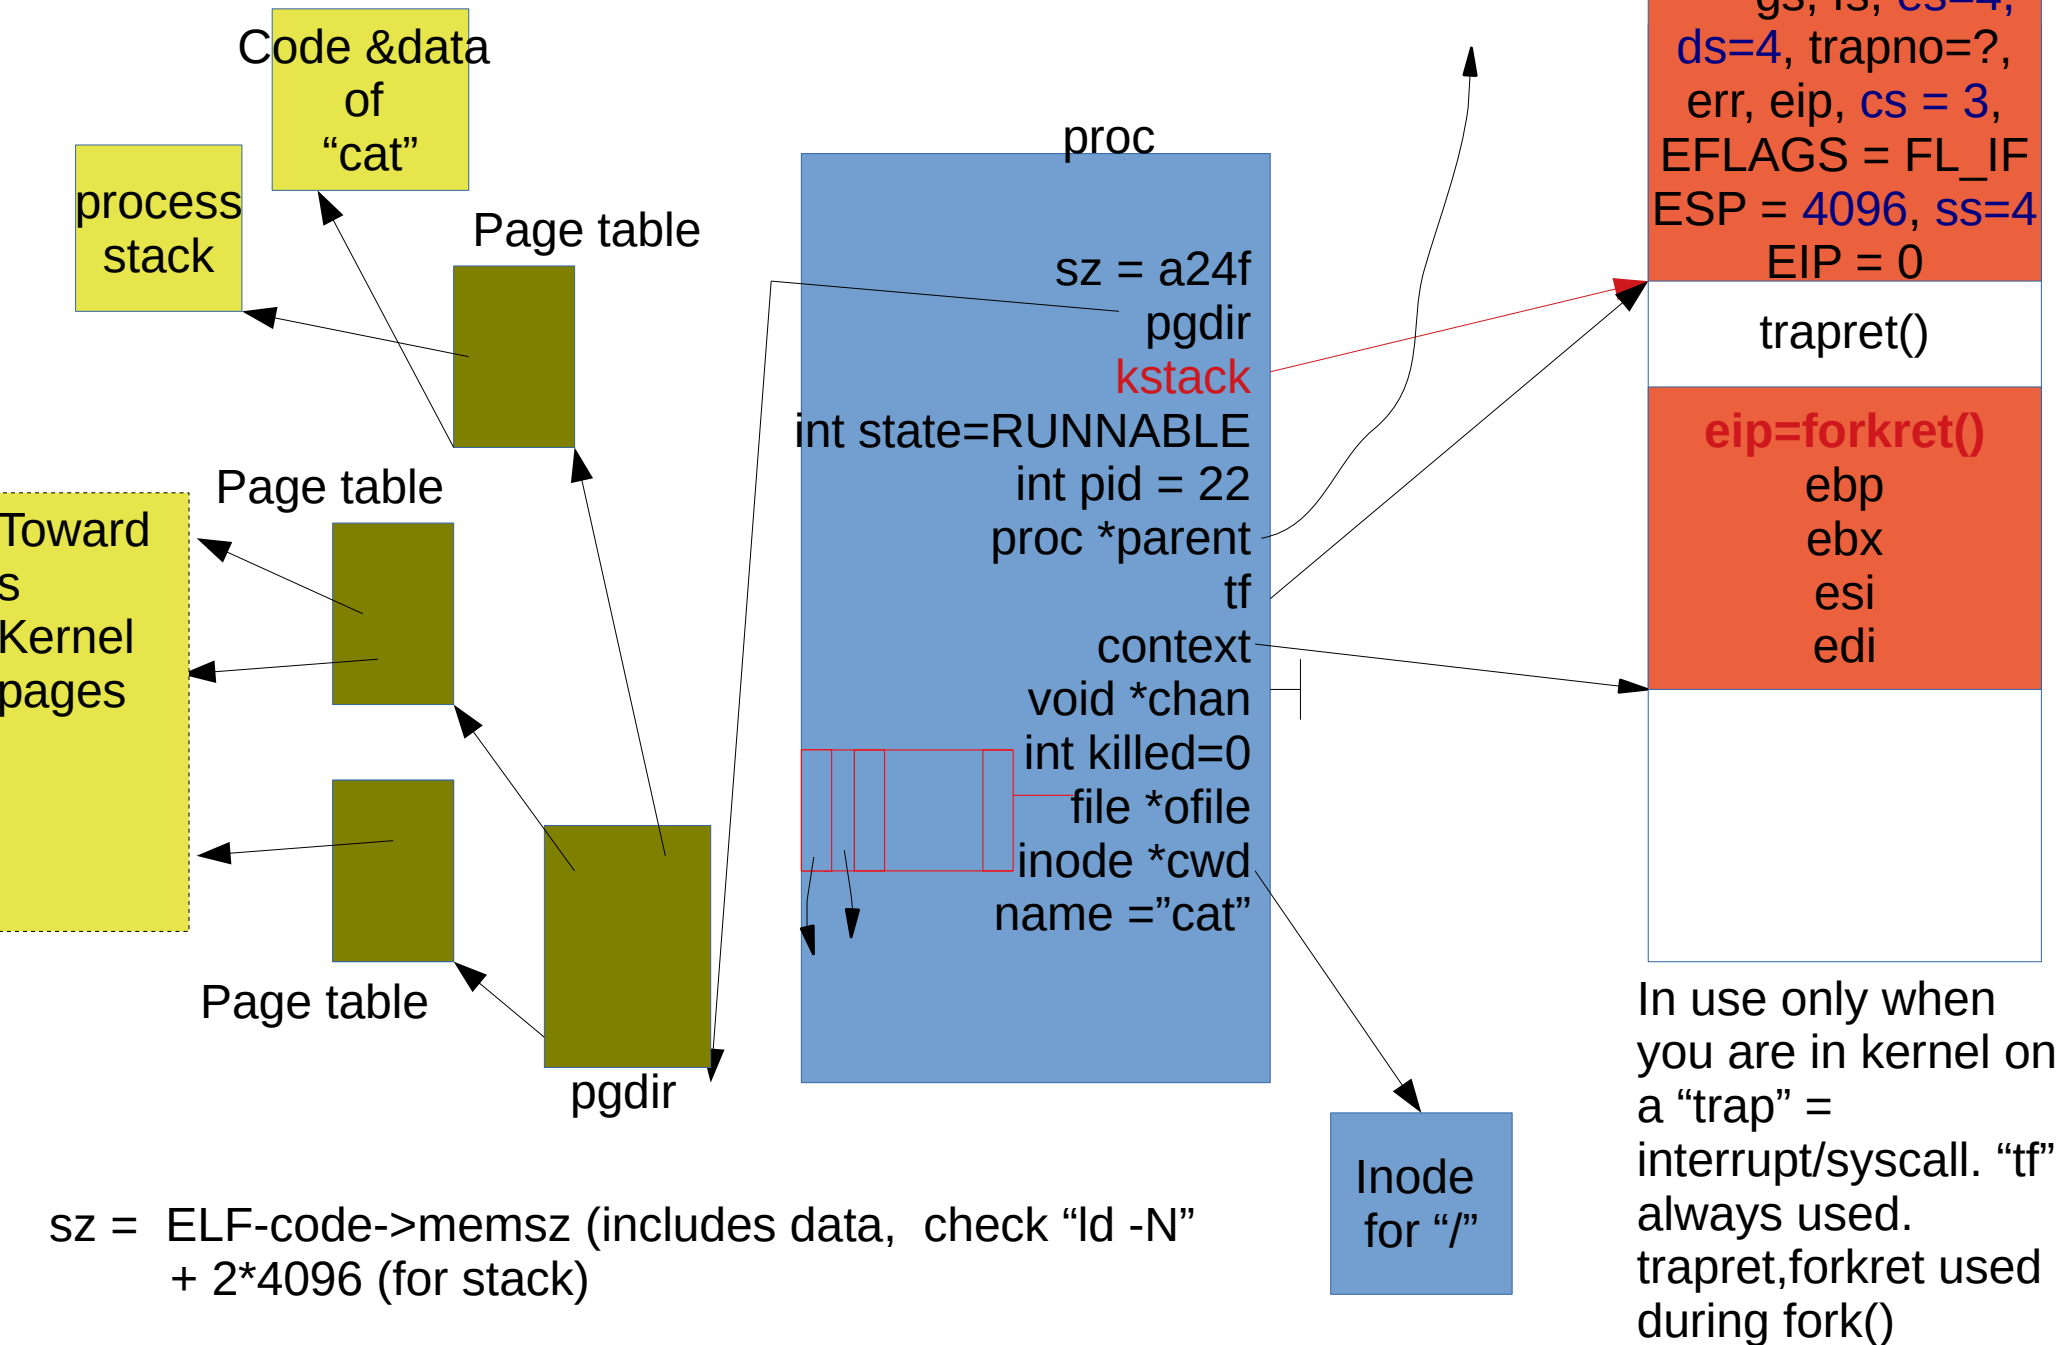
First process creation

Let's revisit struct proc

// Per-process state

```
struct proc {  
    uint sz;                // Size of process memory (bytes)  
    pde_t* pgdir;           // Page table  
    char *kstack;           // Bottom of kernel stack for this process  
    enum procstate state;    // Process state. allocated, ready to run, running,  
wait-                       // ing for I/O, or exiting.  
    int pid;                // Process ID  
    struct proc *parent;     // Parent process  
    struct trapframe *tf;    // Trap frame for current syscall  
    struct context *context; // swtch() here to run process. Process's context  
    void *chan;              // If non-zero, sleeping on chan. More when we discuss  
sleep, wakeup  
    int killed;              // If non-zero, have been killed  
    struct file *ofile[NOFILE]; // Open files, used by open(), read(),...  
    struct inode *cwd;        // Current directory, changed with "chdir()"   
    char name[16];           // Process name (for debugging)  
};
```

struct proc diagram



allocproc()

```
static struct proc*  
allocproc(void)  
{  
    struct proc *p;  
    char *sp;  
    acquire(&ptable.lock);  
    for(p = ptable.proc; p <  
        &ptable.proc[NPROC]; p++)  
        if(p->state == UNUSED)  
            goto found;  
    release(&ptable.lock);  
    return 0;
```

found:

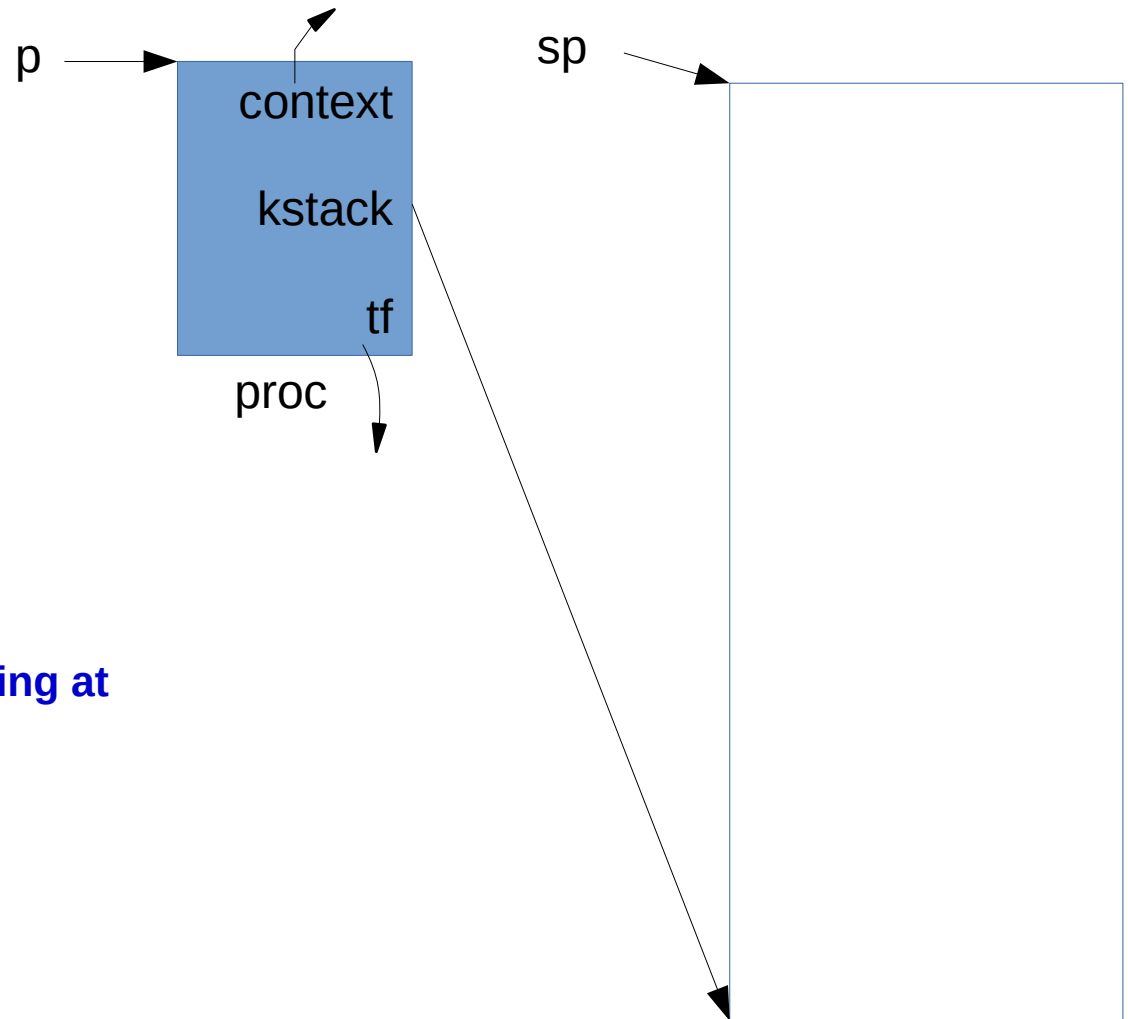
p->state = EMBRYO;

p->pid = nextpid++;

release(&ptable.lock);

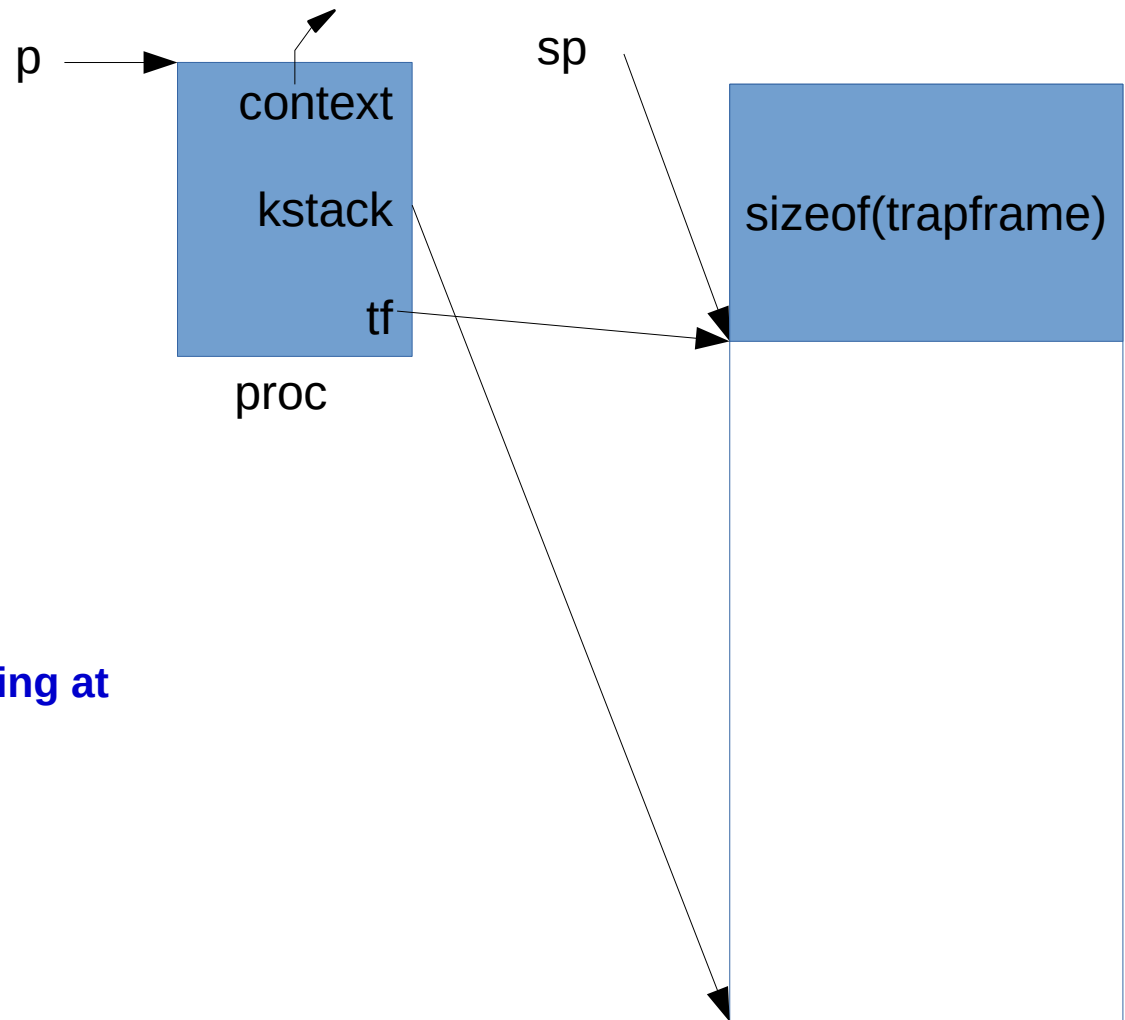
allocproc() setting up stack

```
if((p->kstack = kalloc()) == 0){  
    p->state = UNUSED;  
    return 0;  
}  
sp = p->kstack + KSTACKSIZE;  
// Abhijit KSTACKSIZE = PGSIZE  
// Leave room for trap frame.  
sp -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
// Set up new context to start executing at  
forkret,  
// which returns to trapret.  
sp -= 4;  
*(uint*)sp = (uint)trapret;  
sp -= sizeof *p->context;  
p->context = (struct context*)sp;  
memset(p->context, 0, sizeof *p->context);  
p->context->eip = (uint)forkret;
```



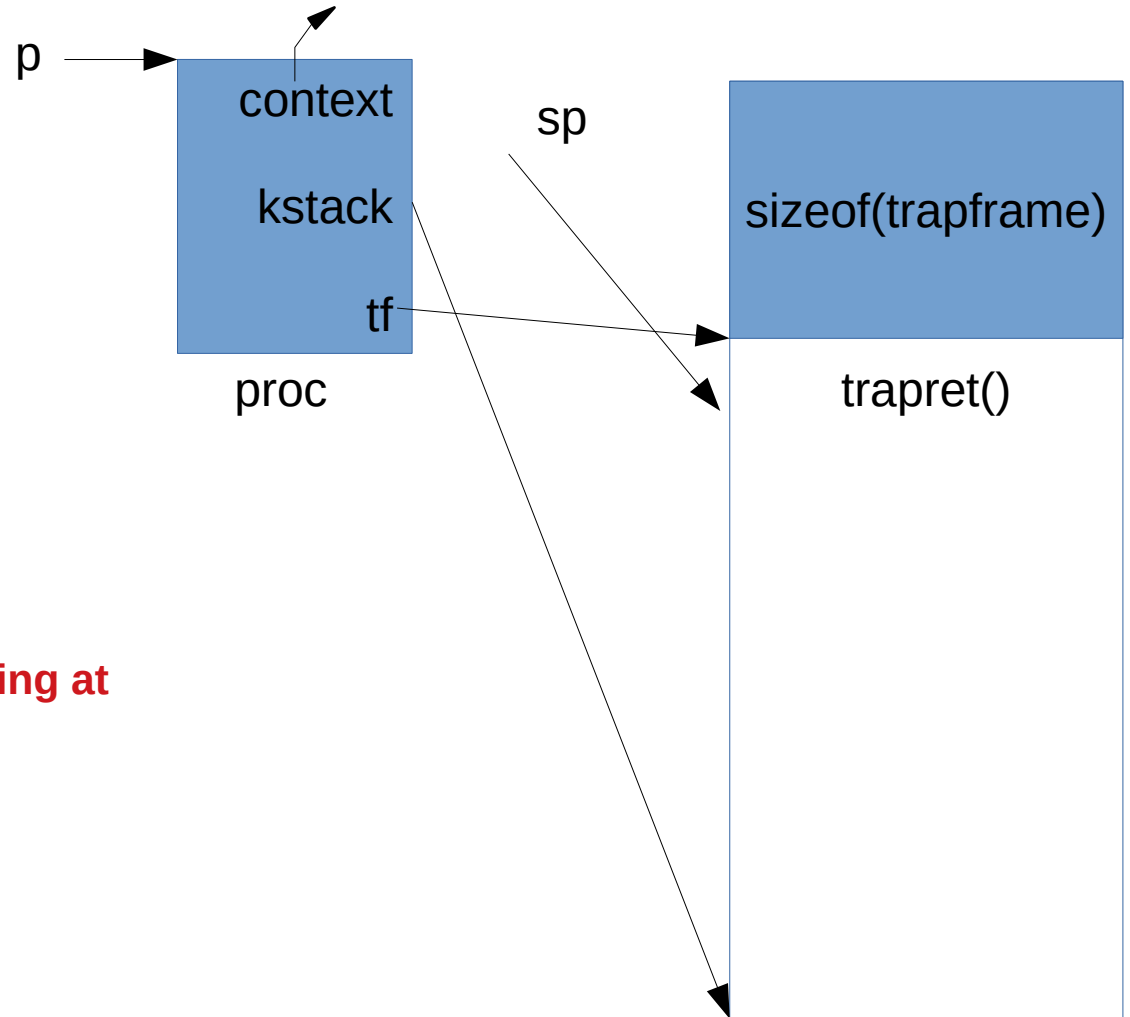
allocproc() setting up stack

```
if((p->kstack = kalloc()) == 0){  
    p->state = UNUSED;  
    return 0;  
}  
sp = p->kstack + KSTACKSIZE;  
// Abhijit KSTACKSIZE = PGSIZE  
// Leave room for trap frame.  
sp -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
// Set up new context to start executing at  
forkret,  
// which returns to trapret.  
sp -= 4;  
*(uint*)sp = (uint)trapret;  
sp -= sizeof *p->context;  
p->context = (struct context*)sp;  
memset(p->context, 0, sizeof *p->context);  
p->context->eip = (uint)forkret;
```



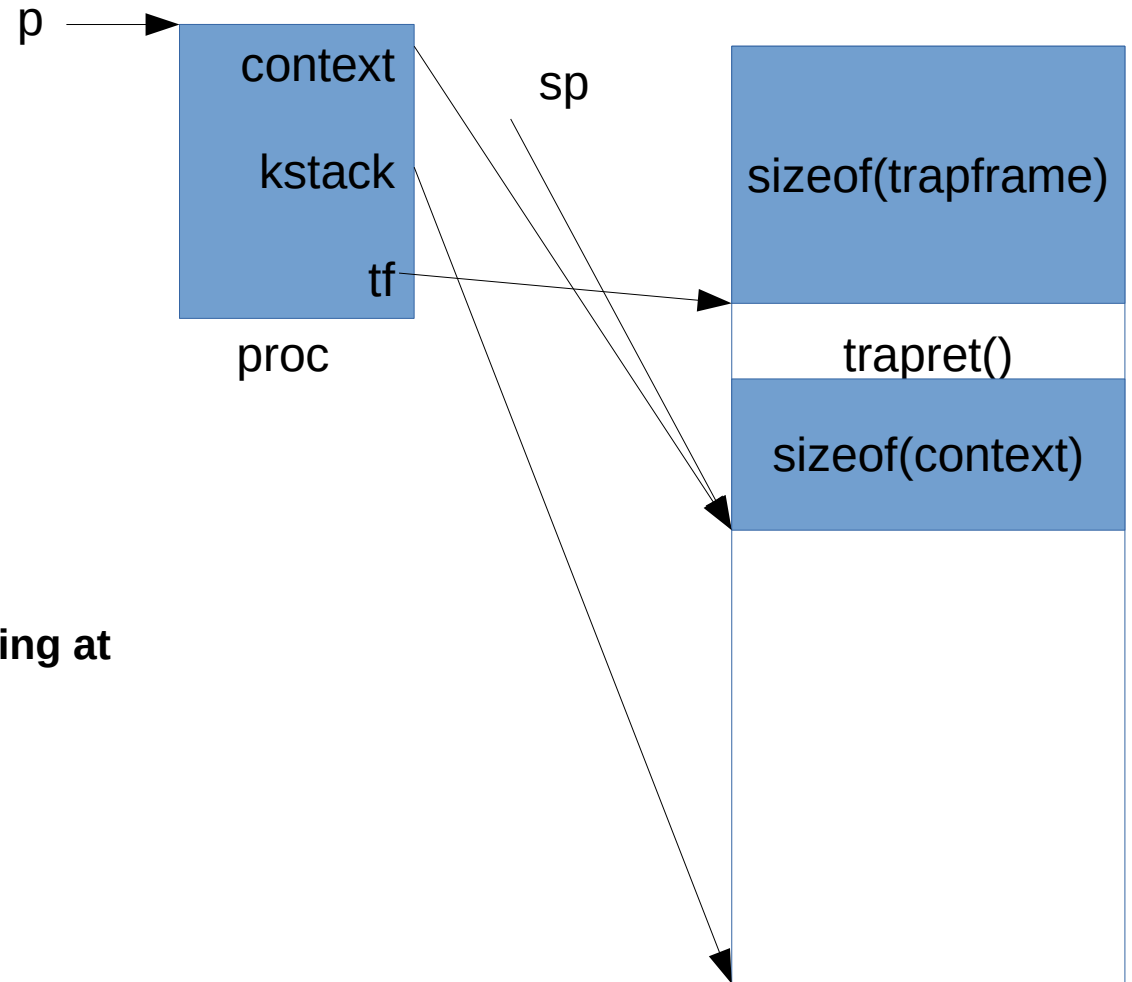
allocproc() setting up stack

```
if((p->kstack = kalloc()) == 0){  
    p->state = UNUSED;  
    return 0;  
}  
sp = p->kstack + KSTACKSIZE;  
// Abhijit KSTACKSIZE = PGSIZE  
// Leave room for trap frame.  
sp -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
// Set up new context to start executing at  
forkret,  
// which returns to trapret.  
sp -= 4;  
*(uint*)sp = (uint)trapret;  
sp -= sizeof *p->context;  
p->context = (struct context*)sp;  
memset(p->context, 0, sizeof *p->context);  
p->context->eip = (uint)forkret;
```



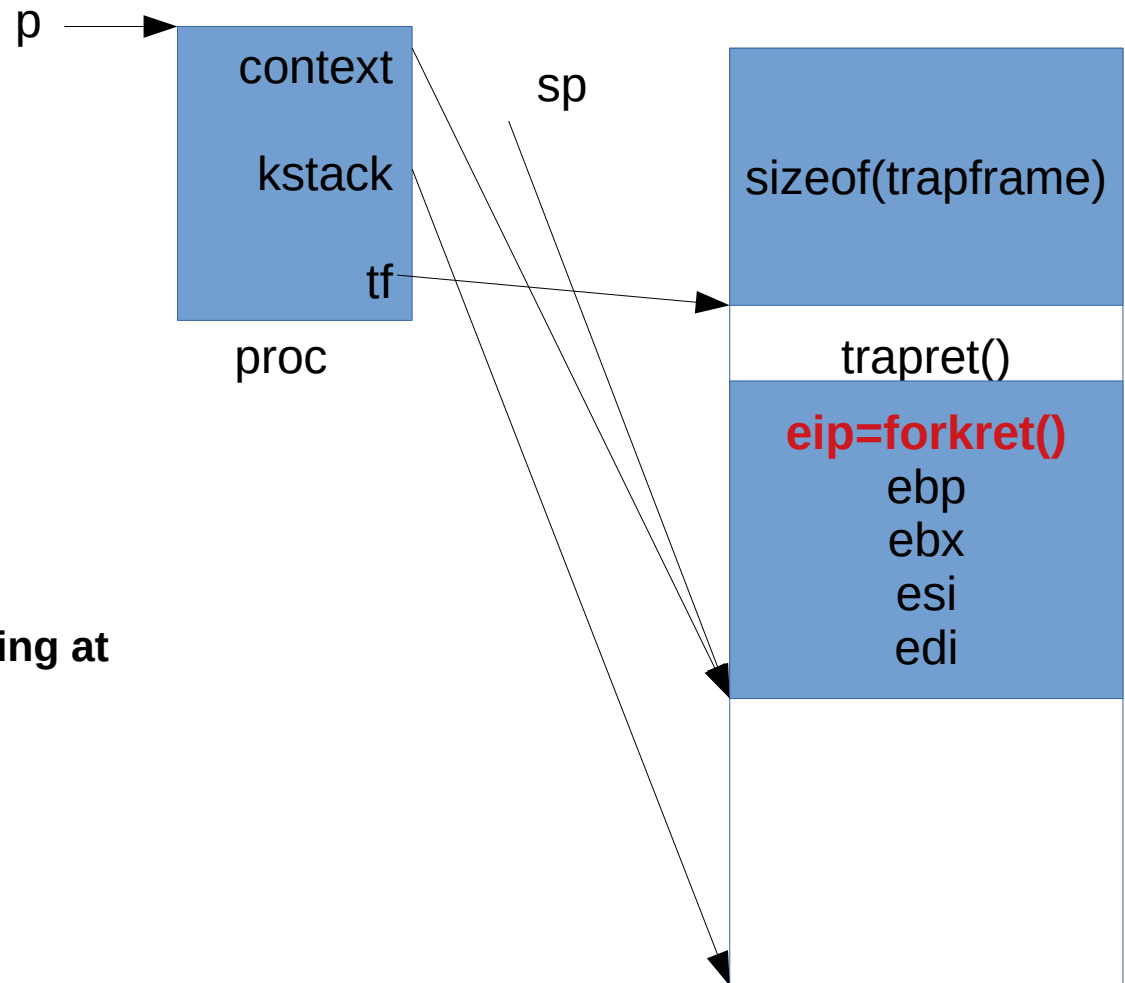
allocproc() setting up stack

```
if((p->kstack = kalloc()) == 0){  
    p->state = UNUSED;  
    return 0;  
}  
sp = p->kstack + KSTACKSIZE;  
// Abhijit KSTACKSIZE = PGSIZE  
// Leave room for trap frame.  
sp -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
// Set up new context to start executing at  
forkret,  
// which returns to trapret.  
sp -= 4;  
*(uint*)sp = (uint)trapret;  
sp -= sizeof *p->context;  
p->context = (struct context*)sp;  
memset(p->context, 0, sizeof *p->context);  
p->context->eip = (uint)forkret;
```



allocproc() setting up stack

```
if((p->kstack = kalloc()) == 0){  
    p->state = UNUSED;  
    return 0;  
}  
sp = p->kstack + KSTACKSIZE;  
// Abhijit KSTACKSIZE = PGSIZE  
// Leave room for trap frame.  
sp -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
// Set up new context to start executing at  
forkret,  
// which returns to trapret.  
sp -= 4;  
*(uint*)sp = (uint)trapret;  
sp -= sizeof *p->context;  
p->context = (struct context*)sp;  
memset(p->context, 0, sizeof *p->context);  
p->context->eip = (uint)forkret;
```

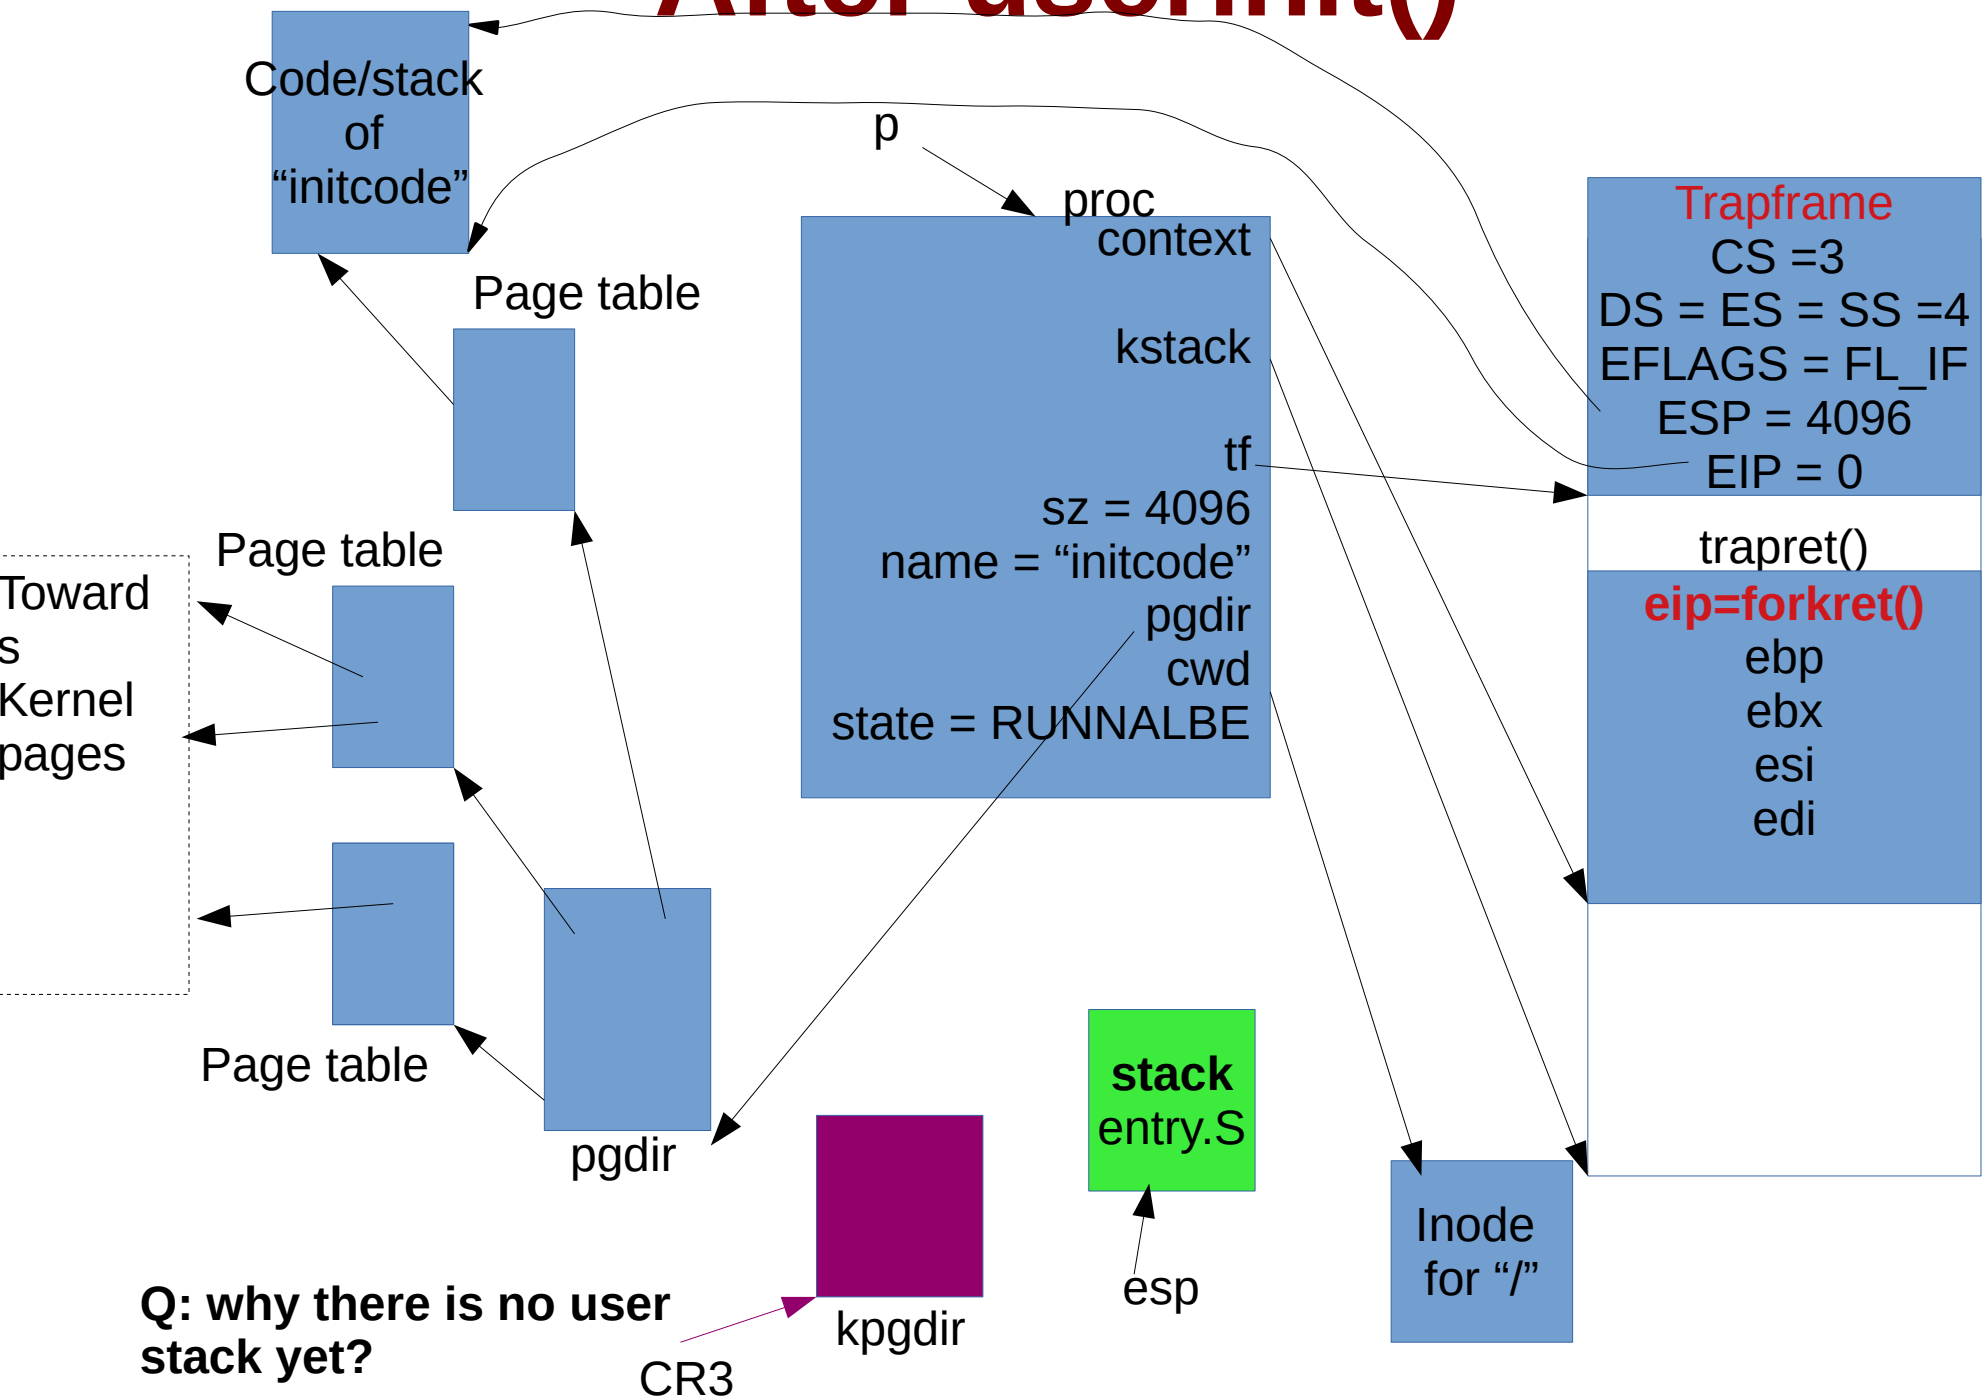


Next in userinit()

```
initproc = p;  
if((p->pgdir = setupkvm()) == 0)  
    panic("userinit: out of  
memory?");  
inituvm(p->pgdir,  
_binary_initcode_start,  
(int)_binary_initcode_size);  
p->sz = PGSIZE;  
memset(p->tf, 0, sizeof(*p->tf));  
p->tf->cs = (SEG_UCODE << 3) |  
DPL_USER;  
p->tf->ds = (SEG_UDATA << 3) |  
DPL_USER;  
p->tf->es = p->tf->ds;  
p->tf->ss = p->tf->ds;
```

```
p->tf->eflags = FL_IF;  
p->tf->esp = PGSIZE;  
p->tf->eip = 0; // beginning of  
initcode.S  
  
safestrcpy(p->name, "initcode",  
sizeof(p->name));  
  
p->cwd = namei("/");  
  
acquire(&ptable.lock);  
  
p->state = RUNNABLE;  
  
release(&ptable.lock);
```

After userinit()



main()->mpmain()

```
static void
mpmain(void)
{
    cprintf("cpu%d: starting %d\n",
    cpuid(), cpuid());
    idtinit();    // load idt register
    xchg(&(mycpu()->started), 1); //
    tell_startothers() we're up
    scheduler();  // start running
    processes
}
```

- **Load IDT register**
 - Copy from `idt[]` array into IDTR
- **Call scheduler()**
 - One process has already been made runnable
 - Let's enter scheduler now

Before reading scheduler(): Note

- The **esp** is still pointing to the **stack** which was allocated in **entry.S** !
 - this is the kernel only stack
 - Not the per process kernel stack.
- **CR3** points to **kpgdir**
- **Struct cpu[]** has been setup up already
 - apicid – in mpinit()
 - segdesc gdt – in seginit()
 - started – in mpmain()
- **Fields in cpu[] not yet set**
 - context * scheduler --> will be setup in sched()
 - taskstate ts --> large structure, only parts used in switchvm()
 - ncli, intena --> used while locking
 - proc *proc -> set during scheduler()

scheduler()

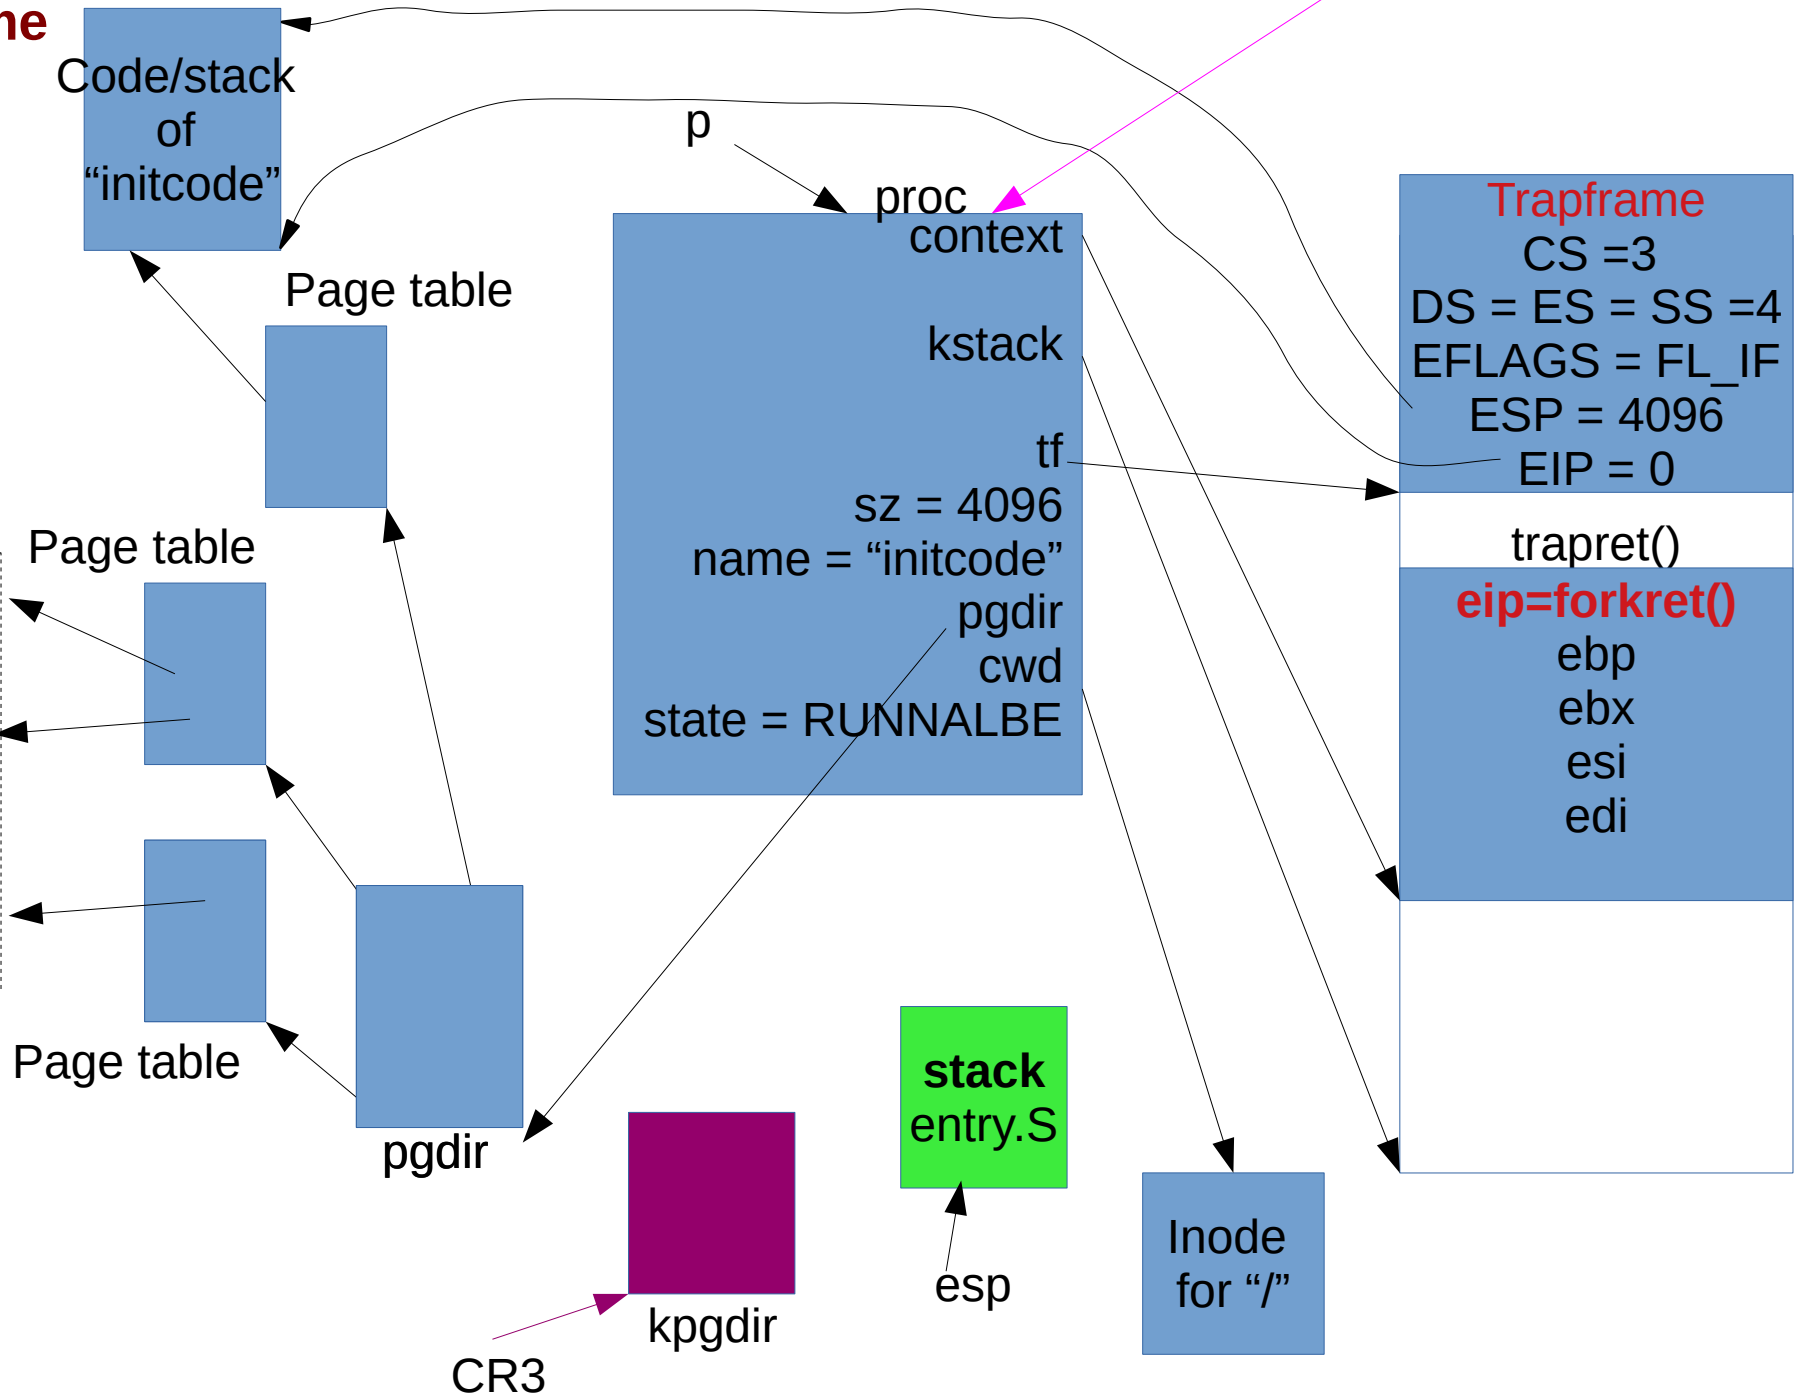
```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        sti();
        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
```


**scheduler()
called first
time**

proc
cpu
*c

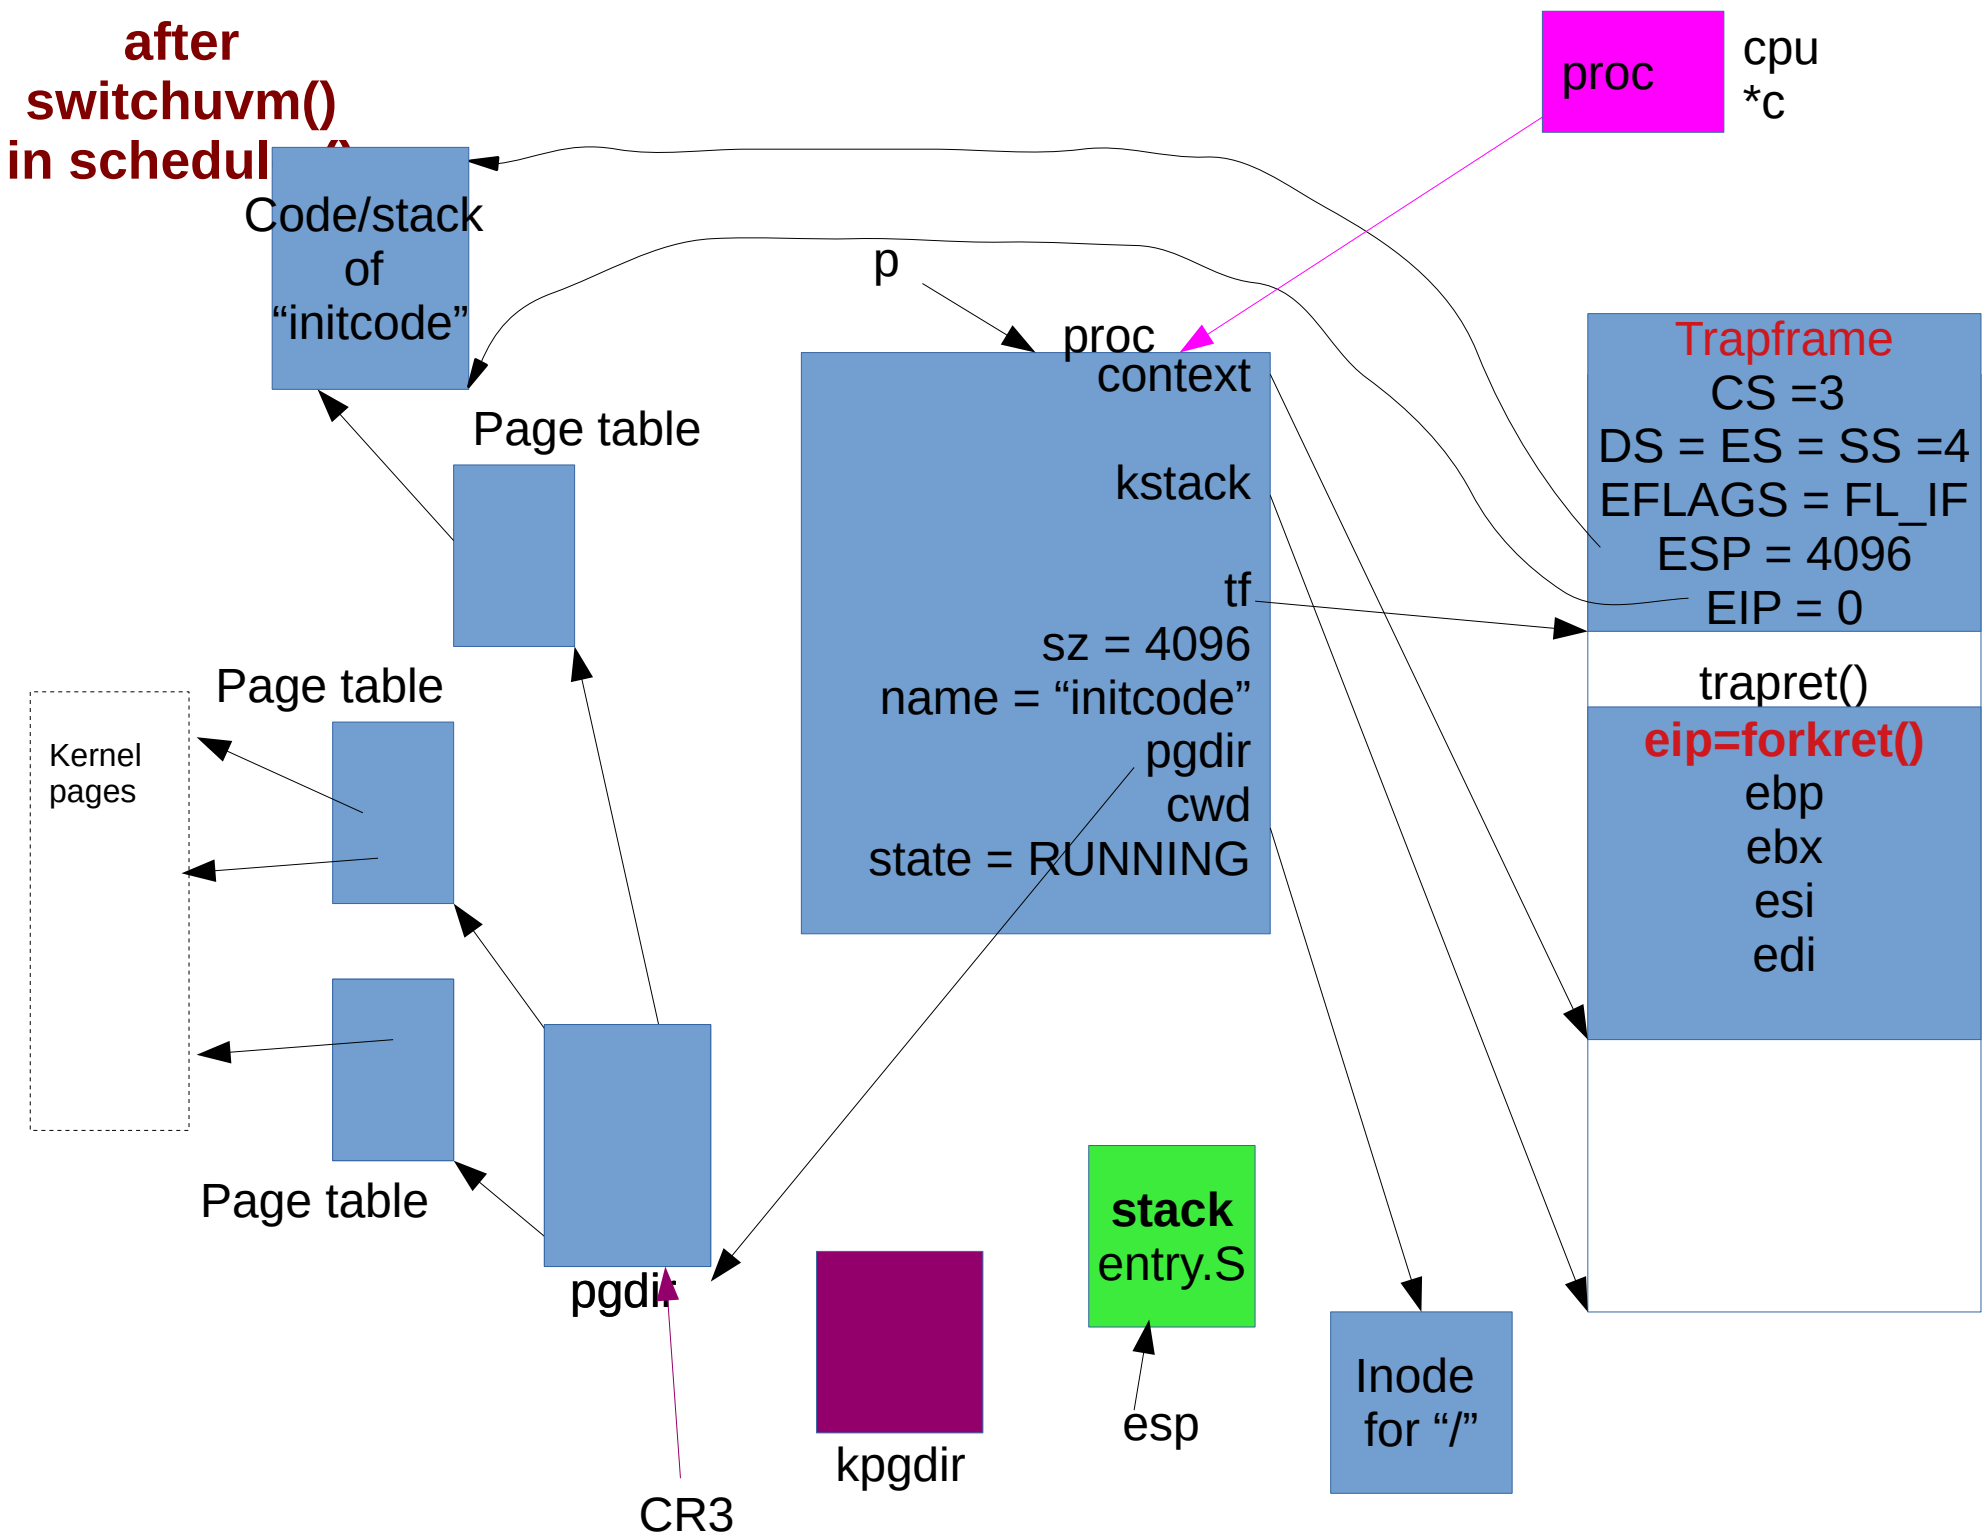
Toward
S
Kernel
pages



scheduler()

```
acquire(&ptable.lock);  
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){  
    if(p->state != RUNNABLE)  
        continue;  
  
    // Switch to chosen process. It is the process's job  
    // to release ptable.lock and then reacquire it  
    // before jumping back to us.  
    c->proc = p;  
    switchvm(p);  
    p->state = RUNNING;
```

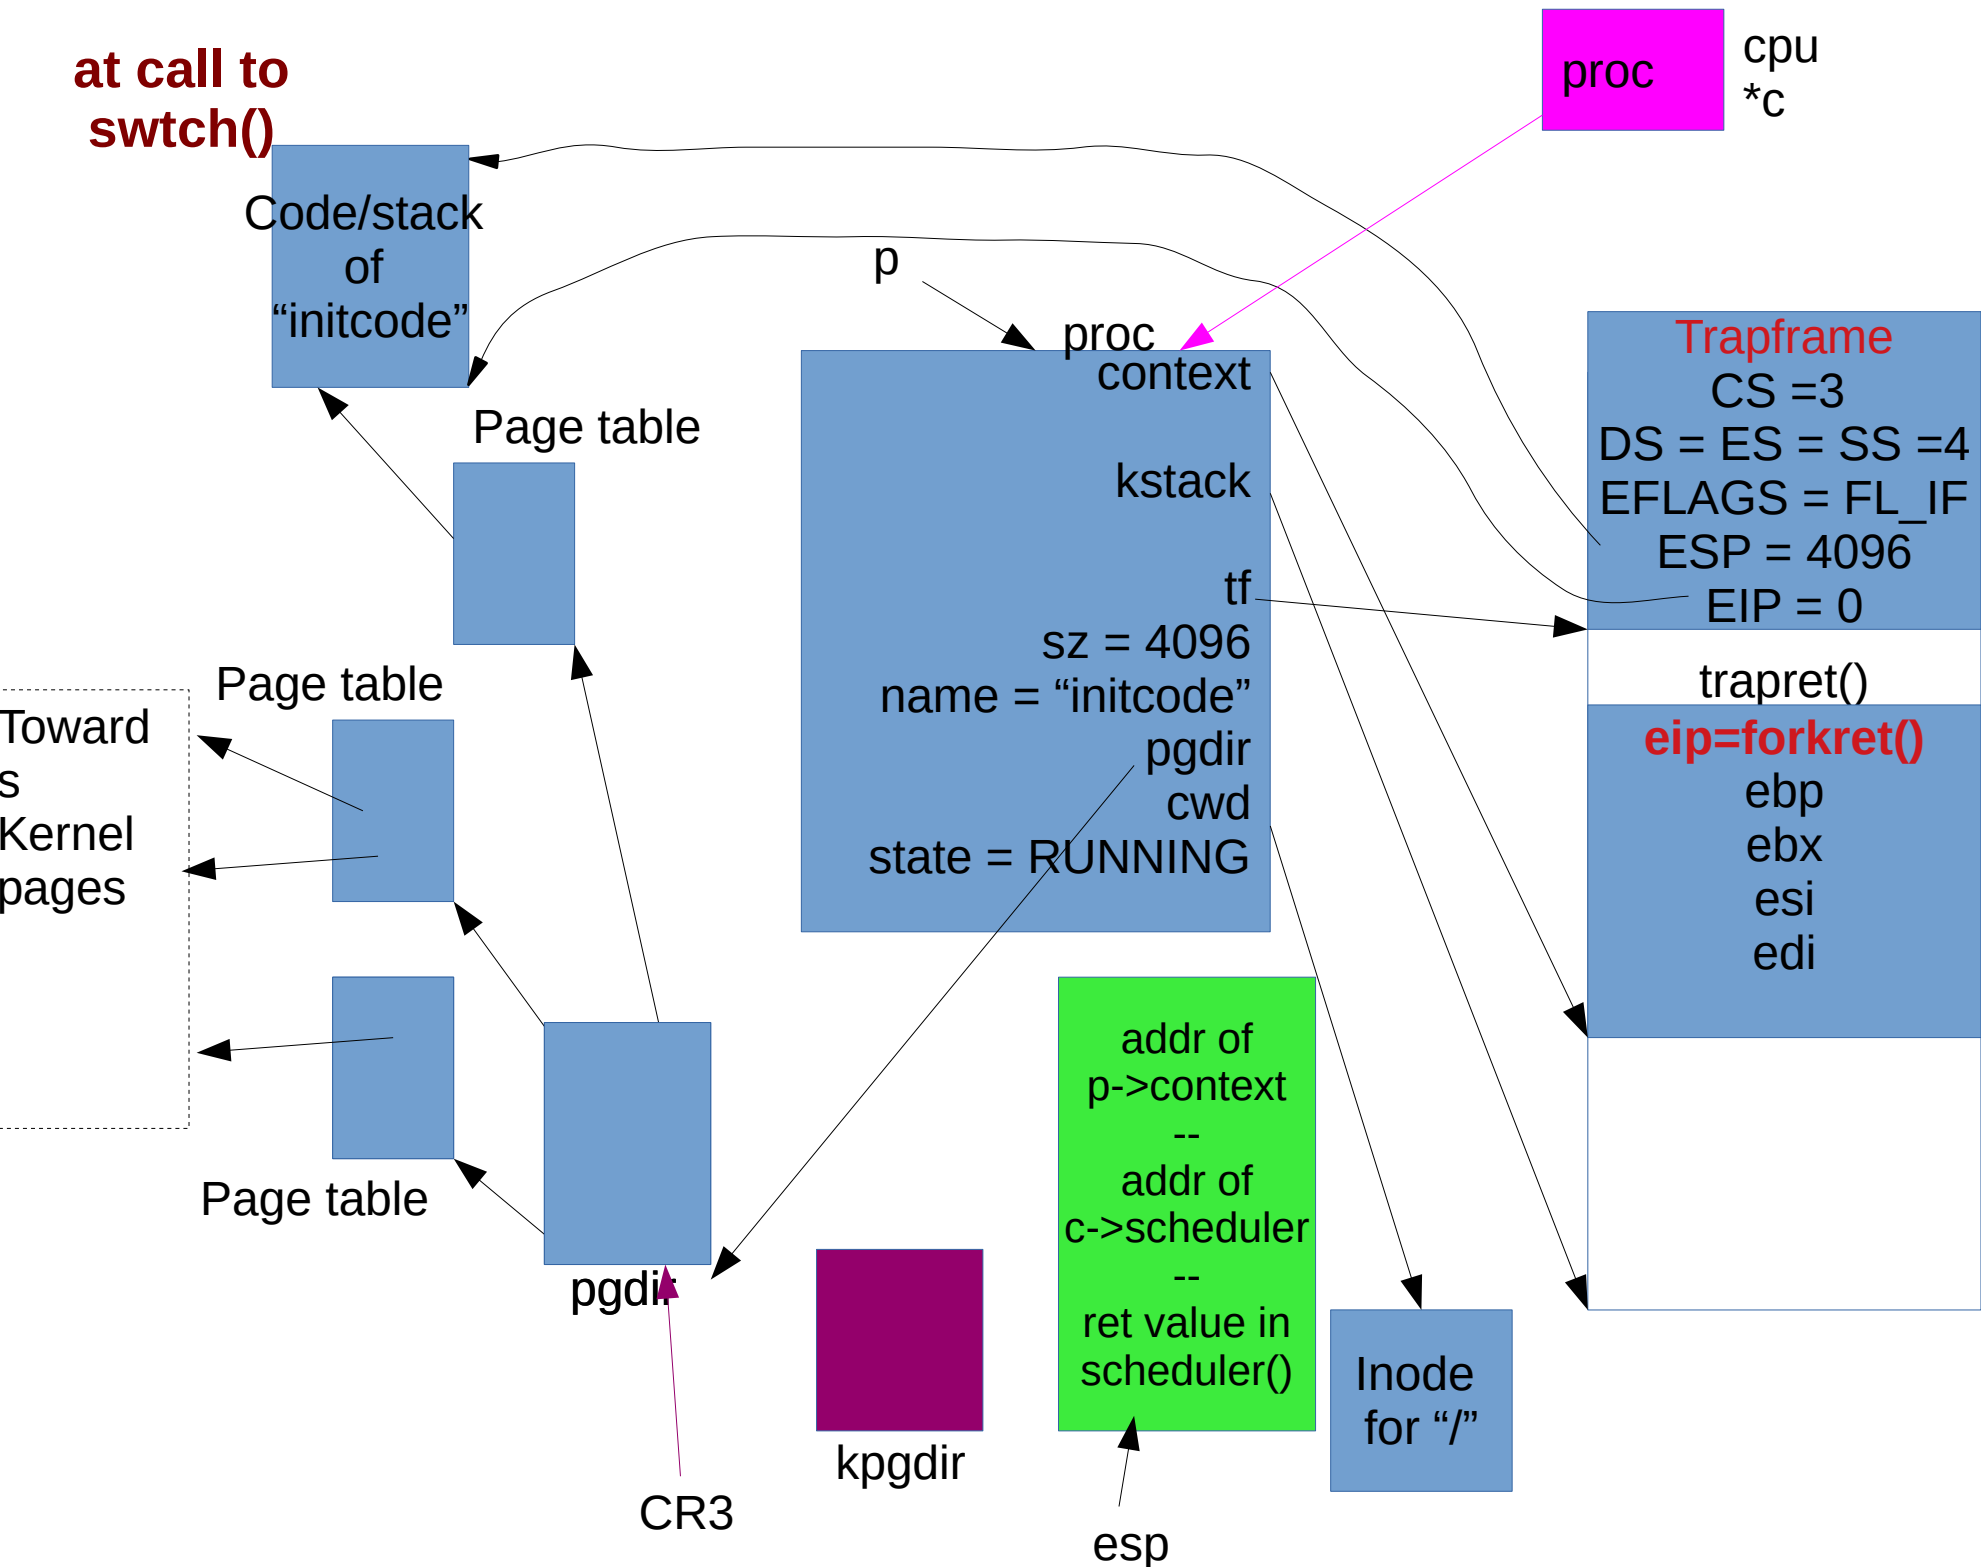
after
switchvm()
in schedul



scheduler()

```
acquire(&ptable.lock);  
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){  
    if(p->state != RUNNABLE)  
        continue;  
  
    // Switch to chosen process. It is the process's job  
    // to release ptable.lock and then reacquire it  
    // before jumping back to us.  
    c->proc = p;  
    switchvm(p);  
    p->state = RUNNING  
    swtch(&(c->scheduler), p->context);  
    ;
```

at call to
switch()



swtch

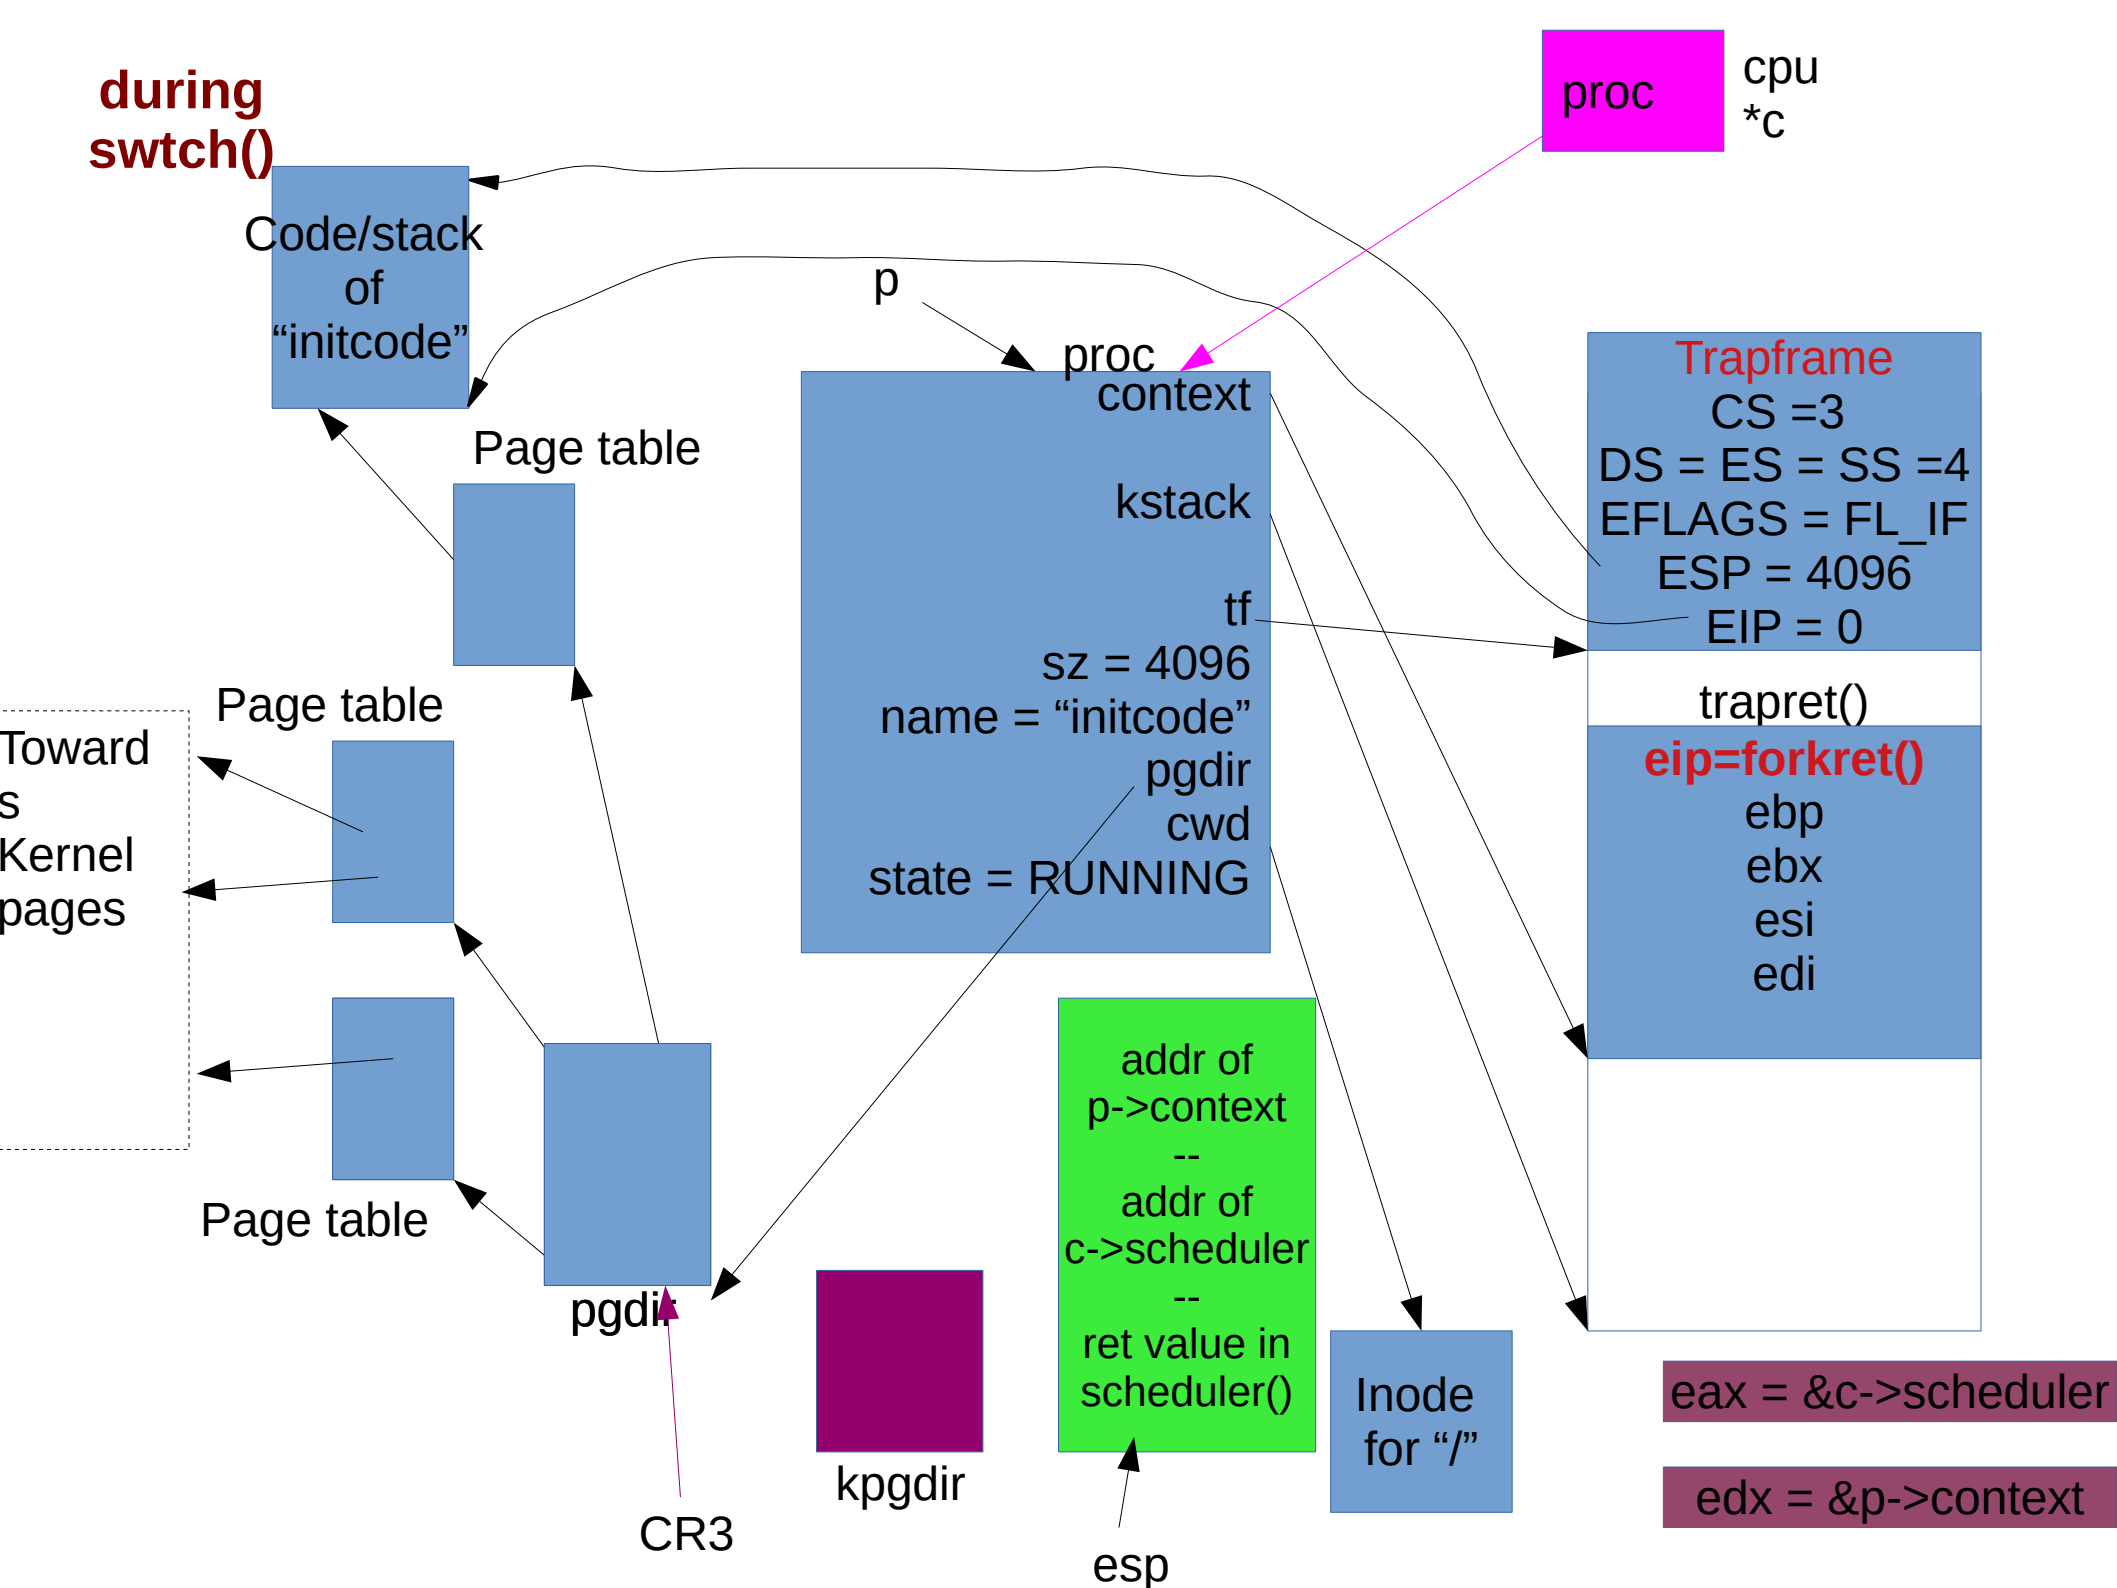
swtch:

#Abhijit: swtch was called through a function call.

#So %eip was saved on stack already

movl 4(%esp), %eax # Abhijit: eax = old

movl 8(%esp), %edx # Abhijit: edx = new



swtch

swtch:

#Abhijit: swtch was called through a function call.

#So %eip was saved on stack already

movl 4(%esp), %eax # Abhijit: eax = old

movl 8(%esp), %edx # Abhijit: edx = new

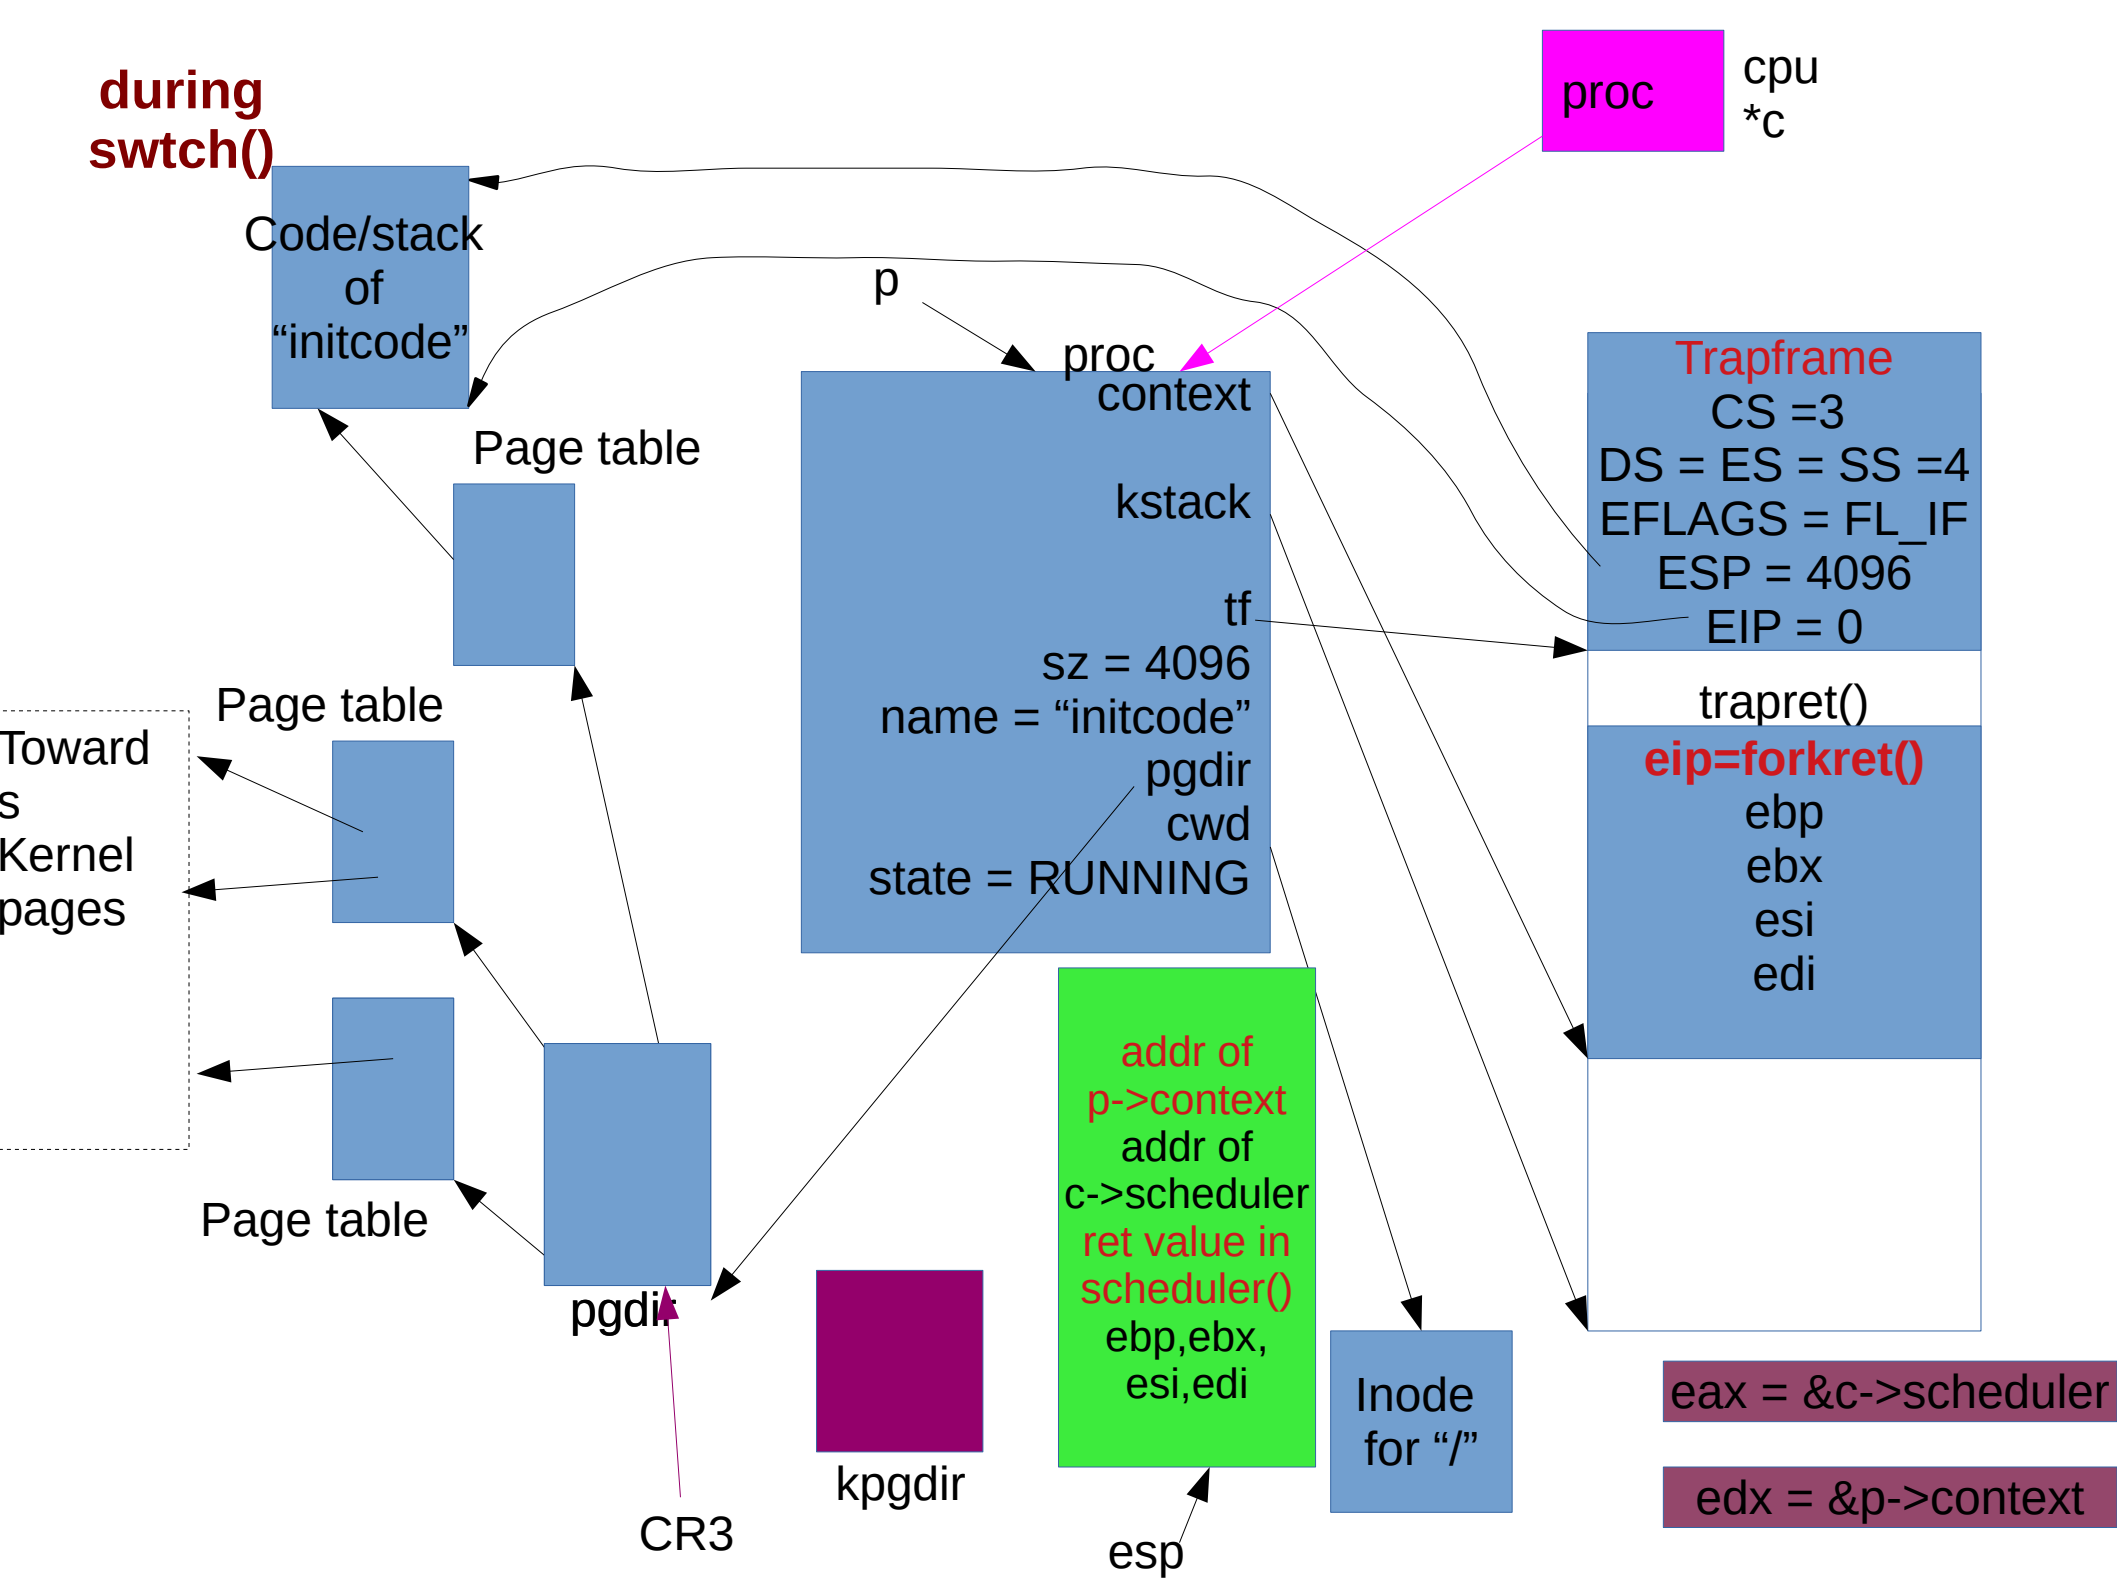
Save old callee-saved registers

pushl %ebp

pushl %ebx

pushl %esi

pushl %edi # Abhijit: esp = esp + 16



swtch

swtch:

#Abhijit: swtch was called through a function call.

#So %eip was saved on stack already

movl 4(%esp), %eax # Abhijit: eax = old

movl 8(%esp), %edx # Abhijit: edx = new

Save old callee-saved registers

pushl %ebp

pushl %ebx

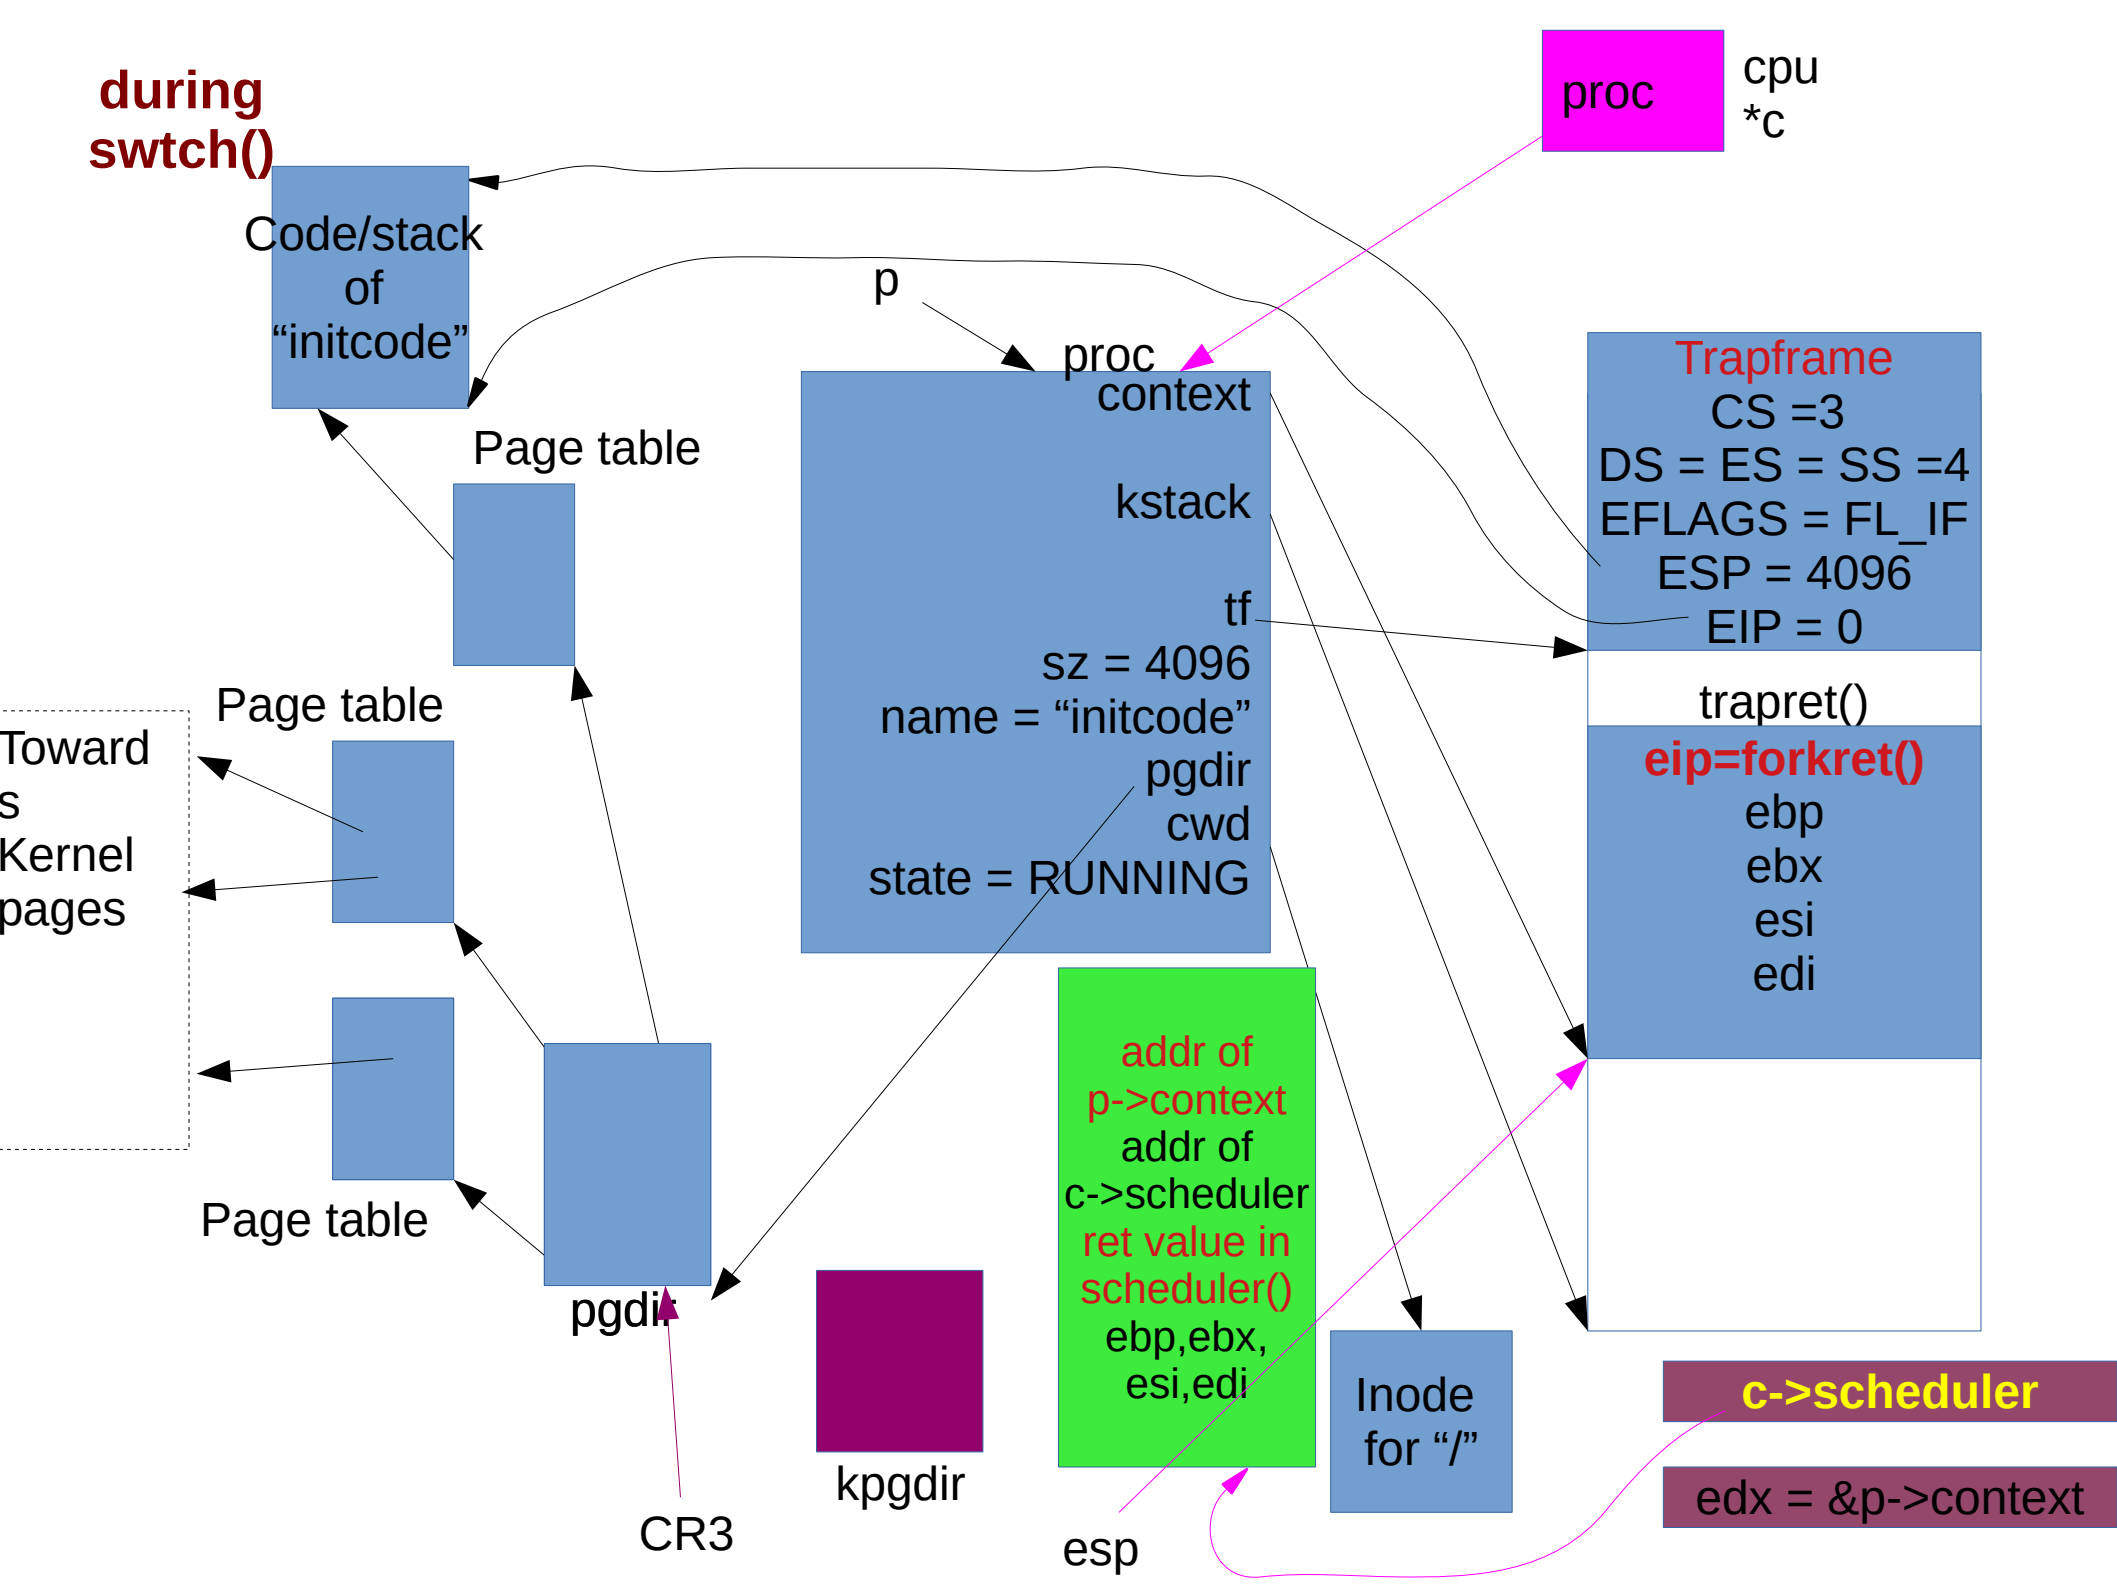
pushl %esi

pushl %edi # Abhijit: esp = esp + 16

Switch stacks

movl %esp, (%eax) # Abhijit: *old = updated old stack

movl %edx, %esp # Abhijit: esp = new



swtch

swtch:

#Abhijit: swtch was called through a function call.

#So %eip was saved on stack already

movl 4(%esp), %eax # Abhijit: eax = old

movl 8(%esp), %edx # Abhijit: edx = new

Save old callee-saved registers

pushl %ebp

pushl %ebx

pushl %esi

pushl %edi # Abhijit: esp = esp + 16

Switch stacks

movl %esp, (%eax) # Abhijit: *old = updated old stack

movl %edx, %esp # Abhijit: esp = new

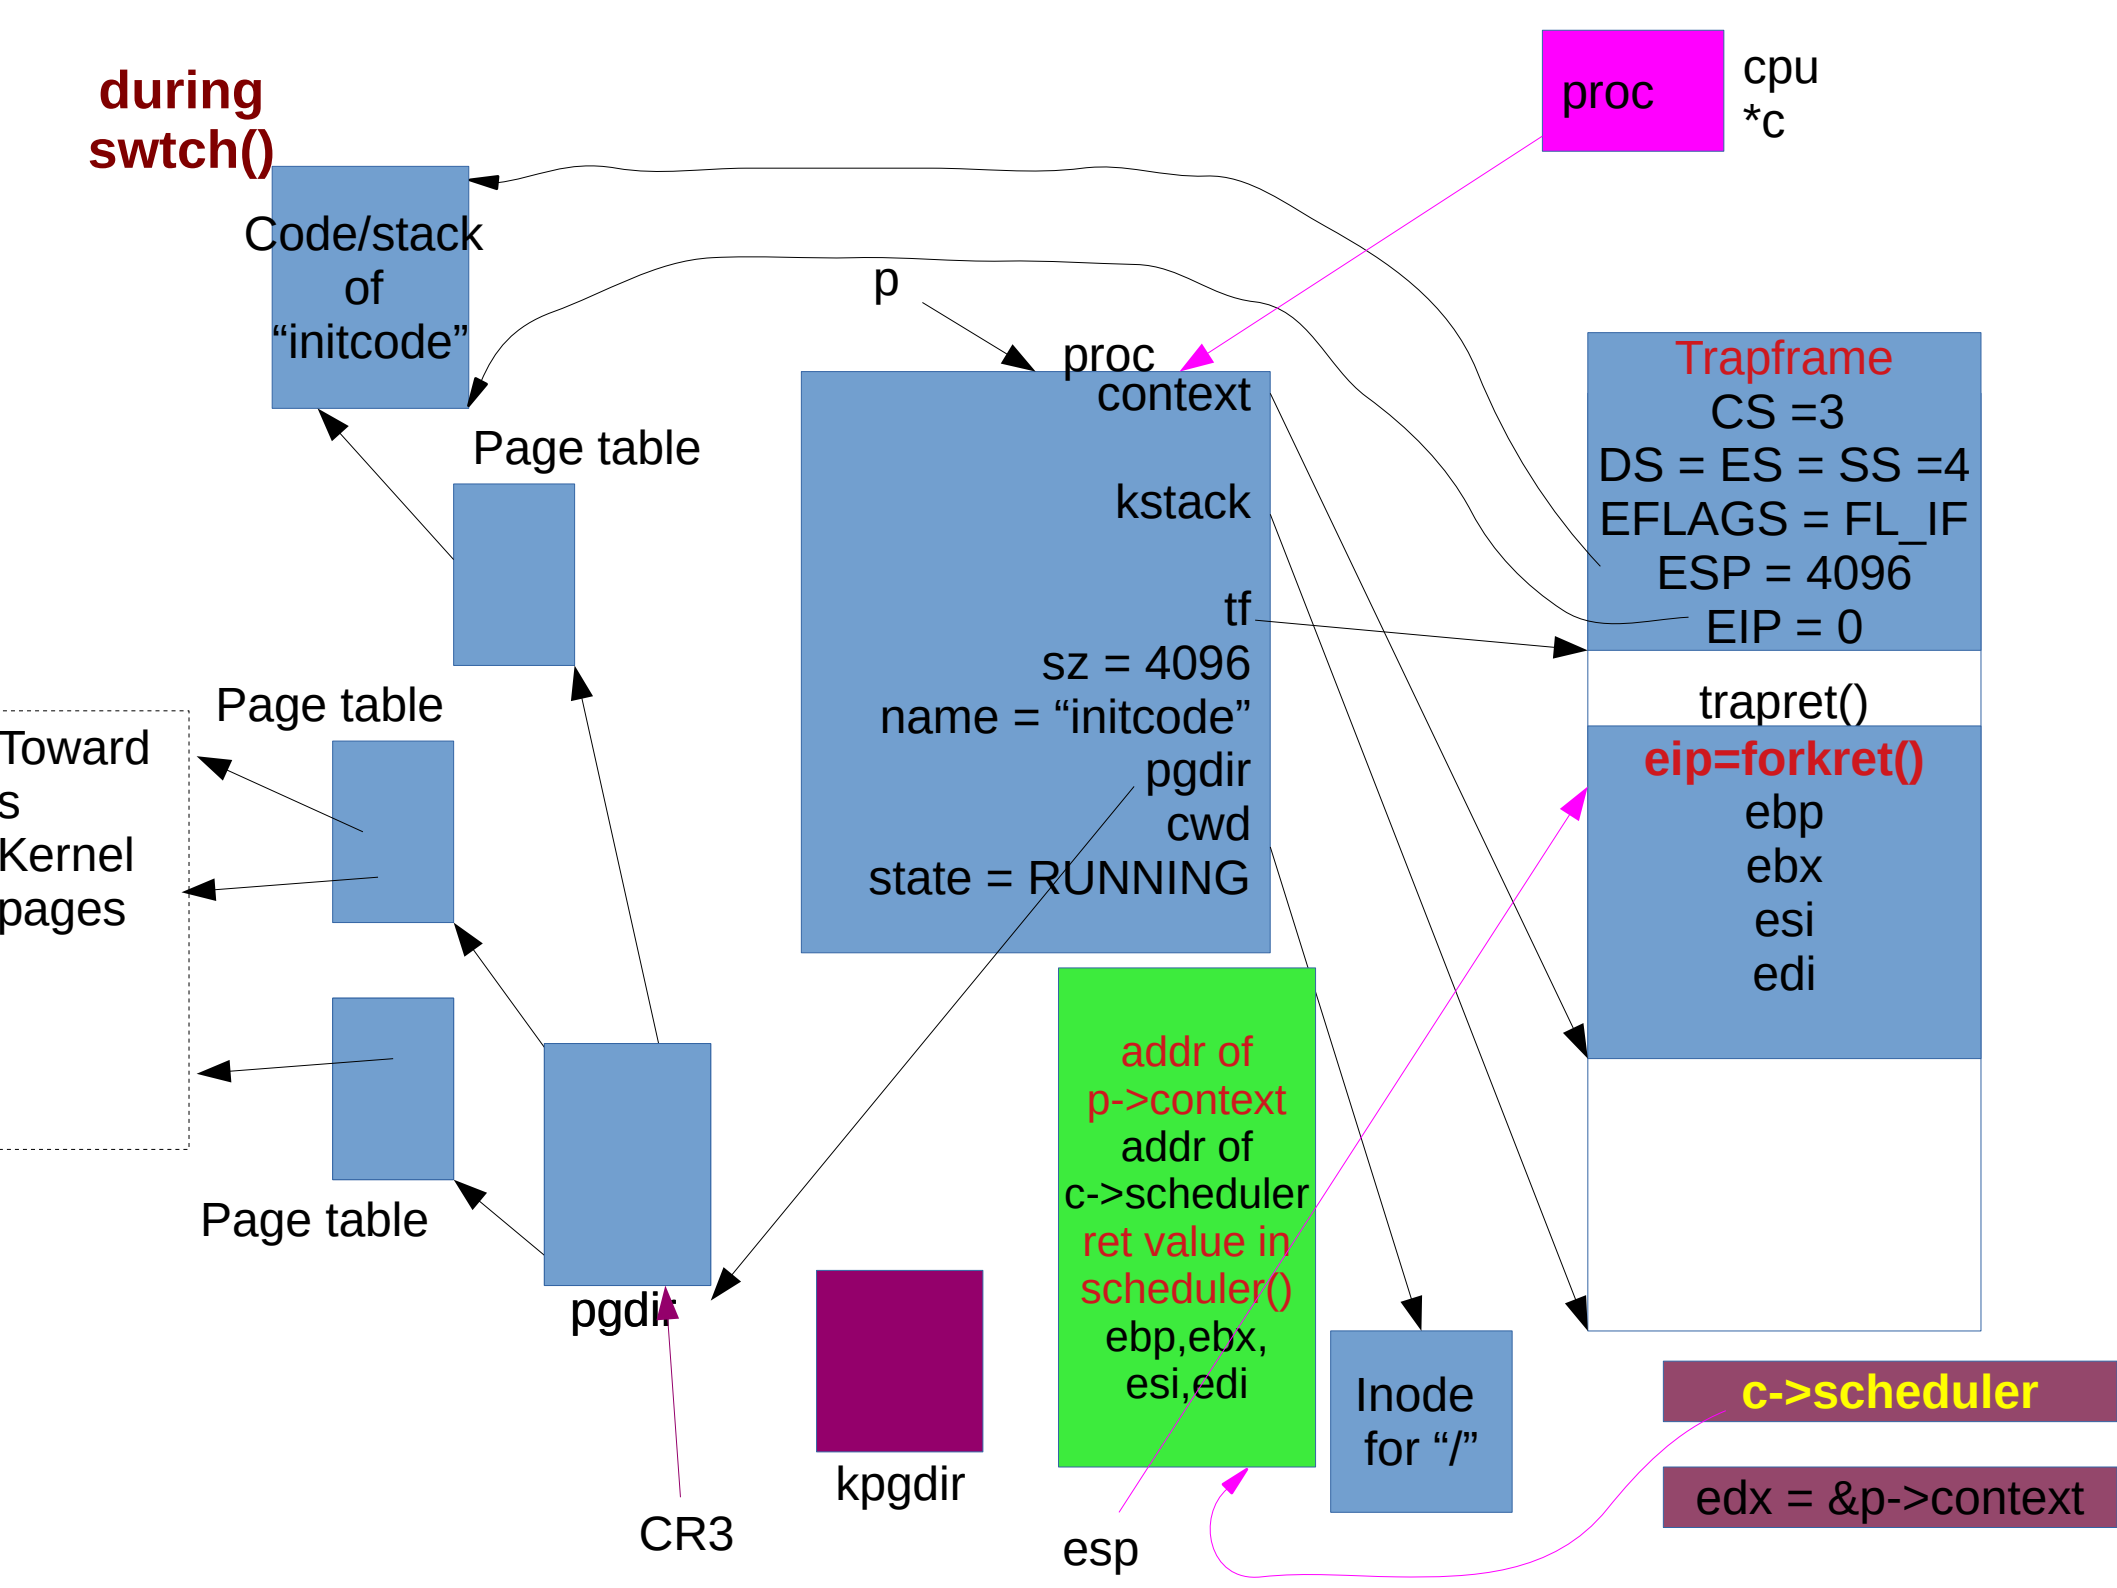
Load new callee-saved registers

popl %edi

popl %esi

popl %ebx

popl %ebp # Abhijit: newesp = newesp - 16, context restored



swtch:

#Abhijit: swtch was called through a function call. **swtch**

#So %eip was saved on stack already

movl 4(%esp), %eax # Abhijit: eax = old

movl 8(%esp), %edx # Abhijit: edx = new

Save old callee-saved registers

pushl %ebp

pushl %ebx

pushl %esi

pushl %edi # Abhijit: esp = esp + 16

Switch stacks

movl %esp, (%eax) # Abhijit: *old = updated old stack

movl %edx, %esp # Abhijit: esp = new

Load new callee-saved registers

popl %edi

popl %esi

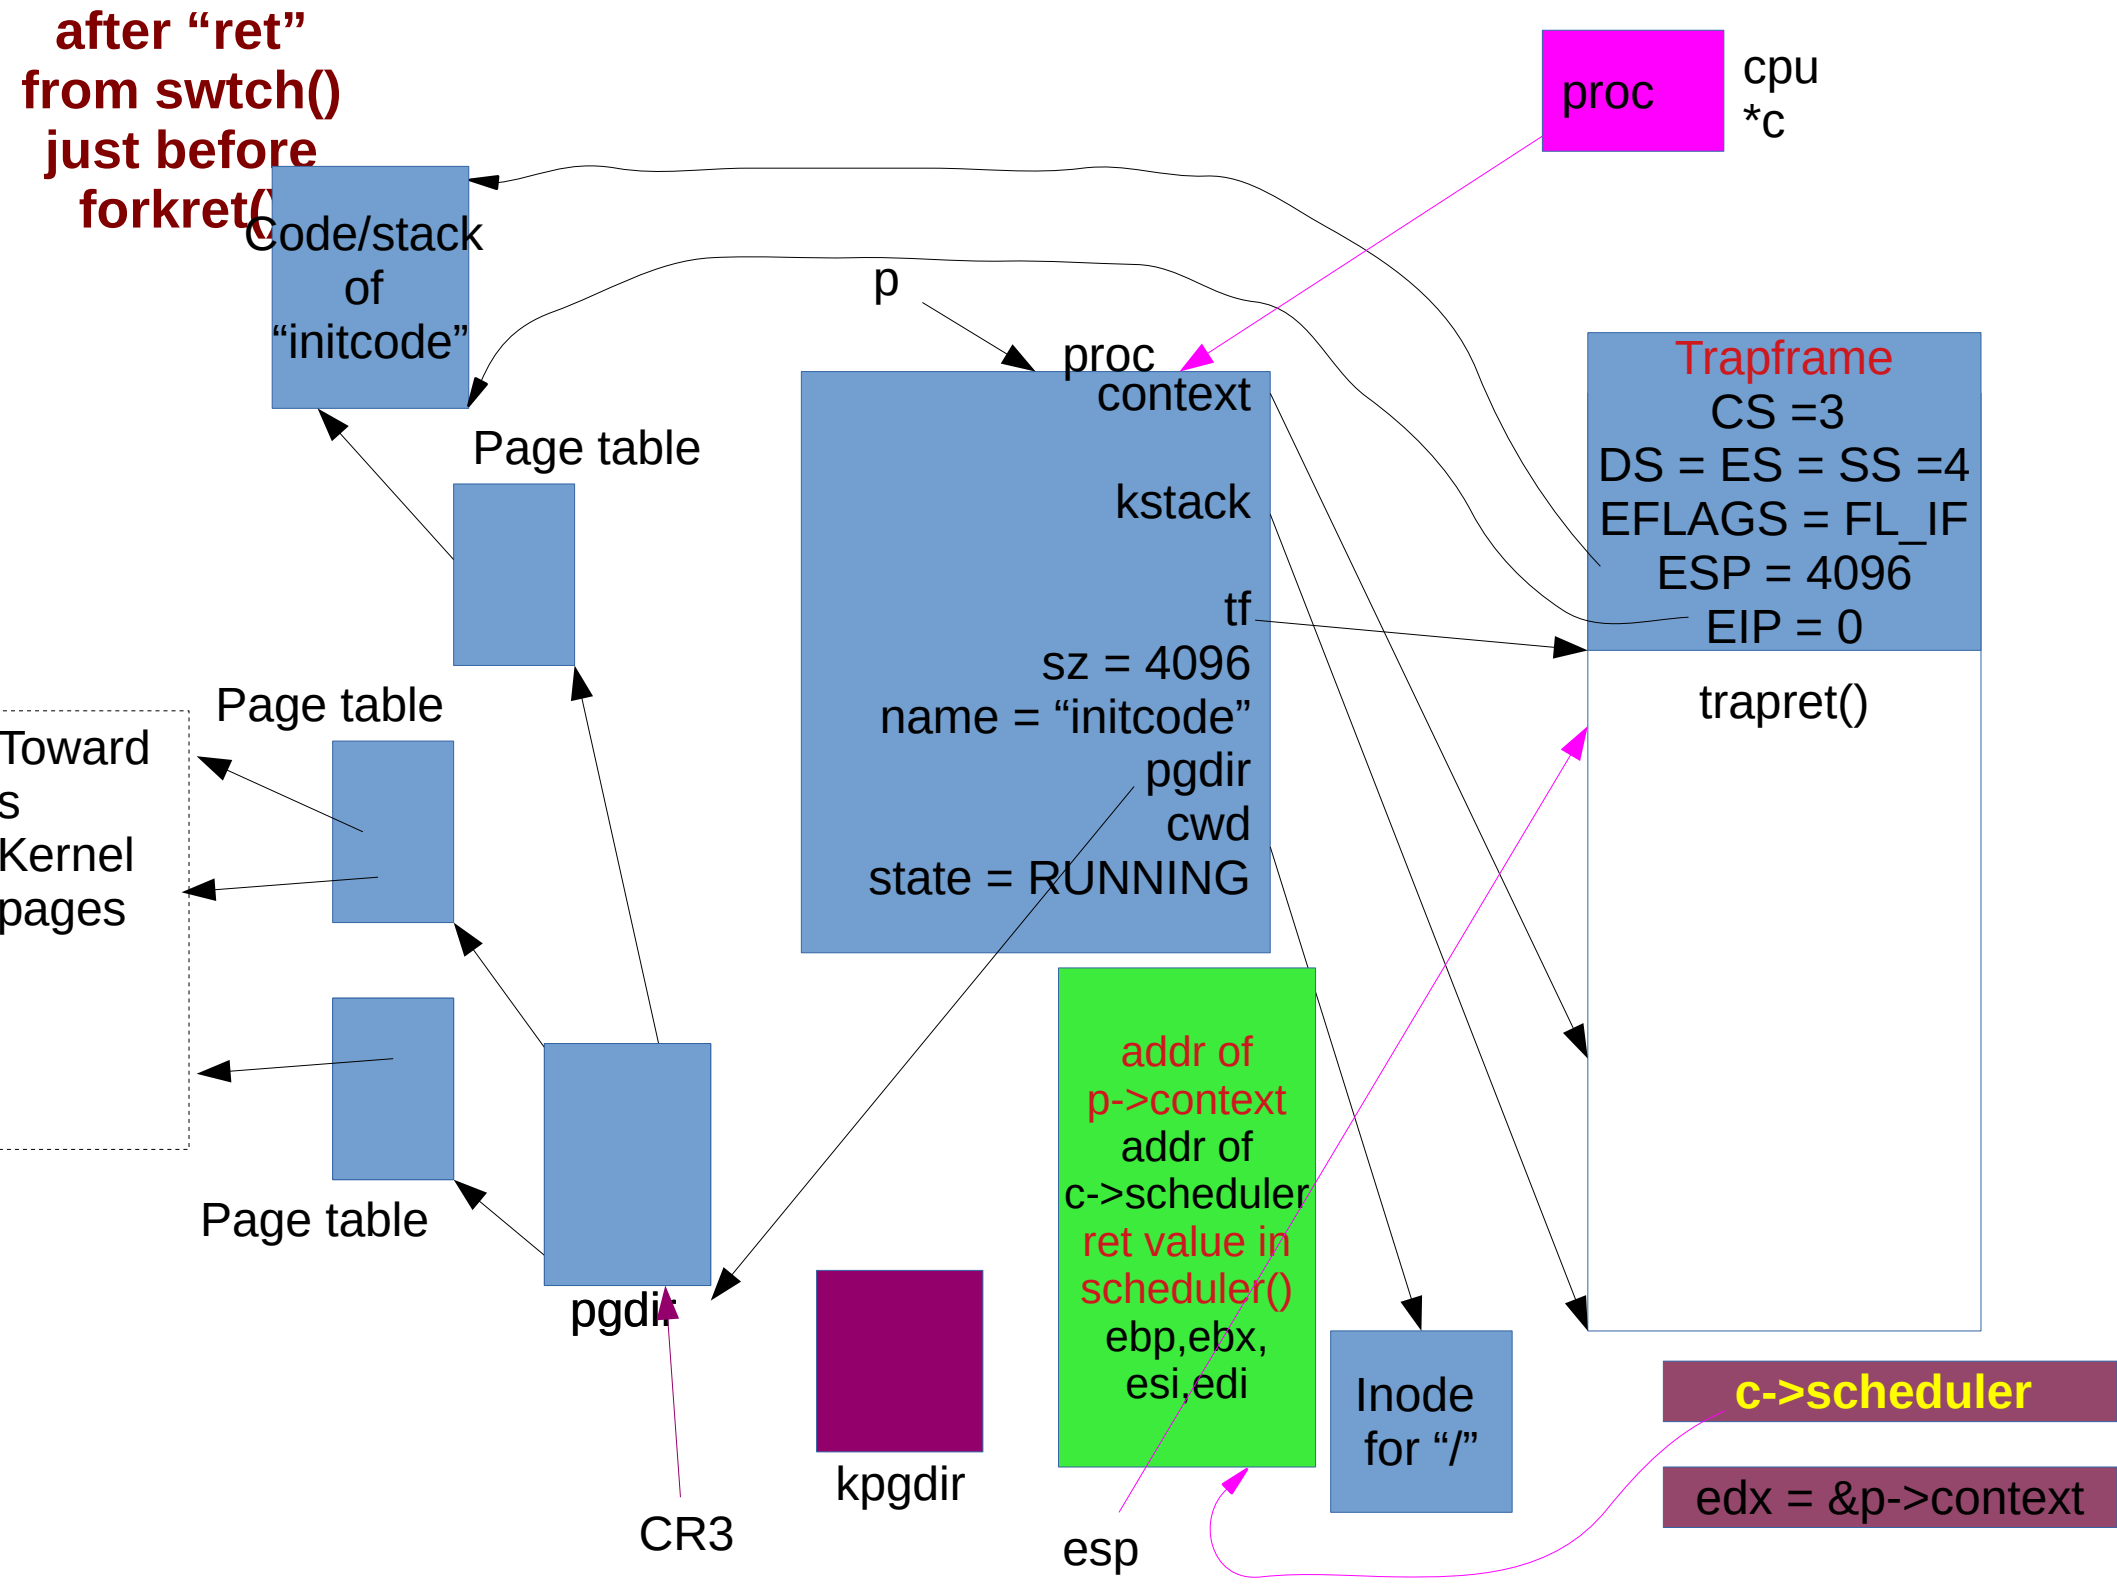
popl %ebx

popl %ebp # Abhijit: newesp = newesp - 16, context restored

ret # Abhijit: will pop from esp now -> function where to

return.

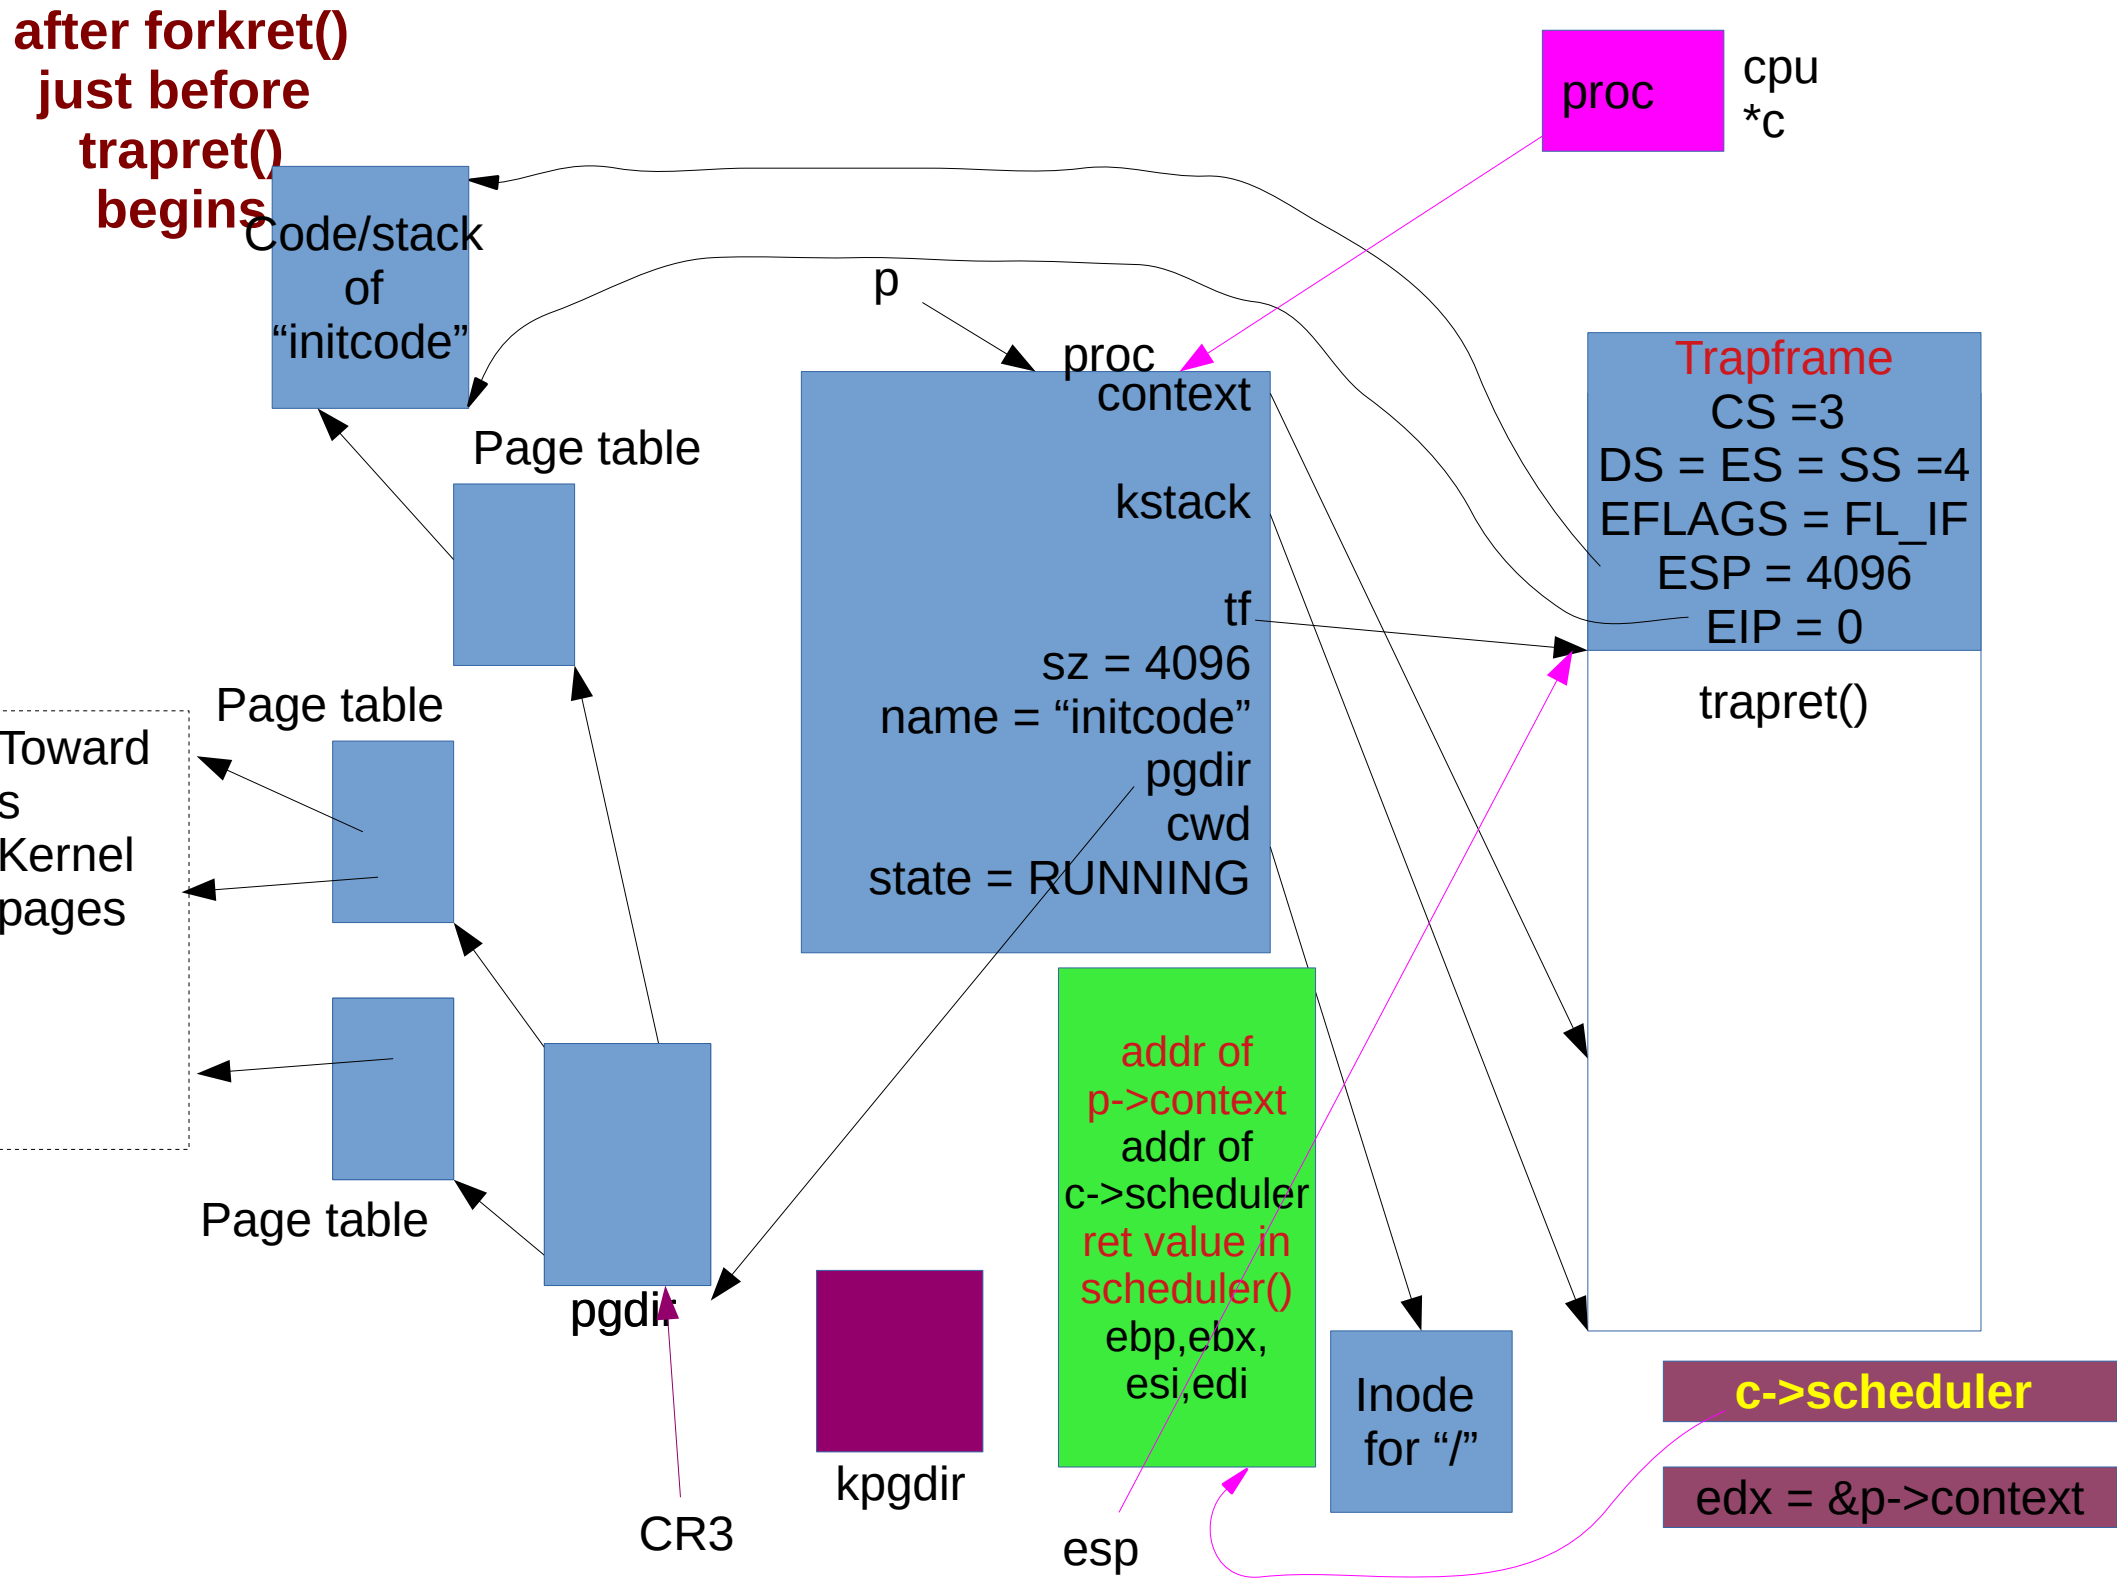
after "ret"
from switch()
just before
forkret()



After switch()

- **Process is running in forkret()**
- **c->cscheduler has saved the old kernel stack**
 - with the context of p, return value in scheduler, ebp, ebx, esi, edi on stack
 - remember {edi, esi, ebx, ebp, ret-value } = context
 - The c->scheduler is pointing to old context
- **CR3 is pointing to process pgdir**

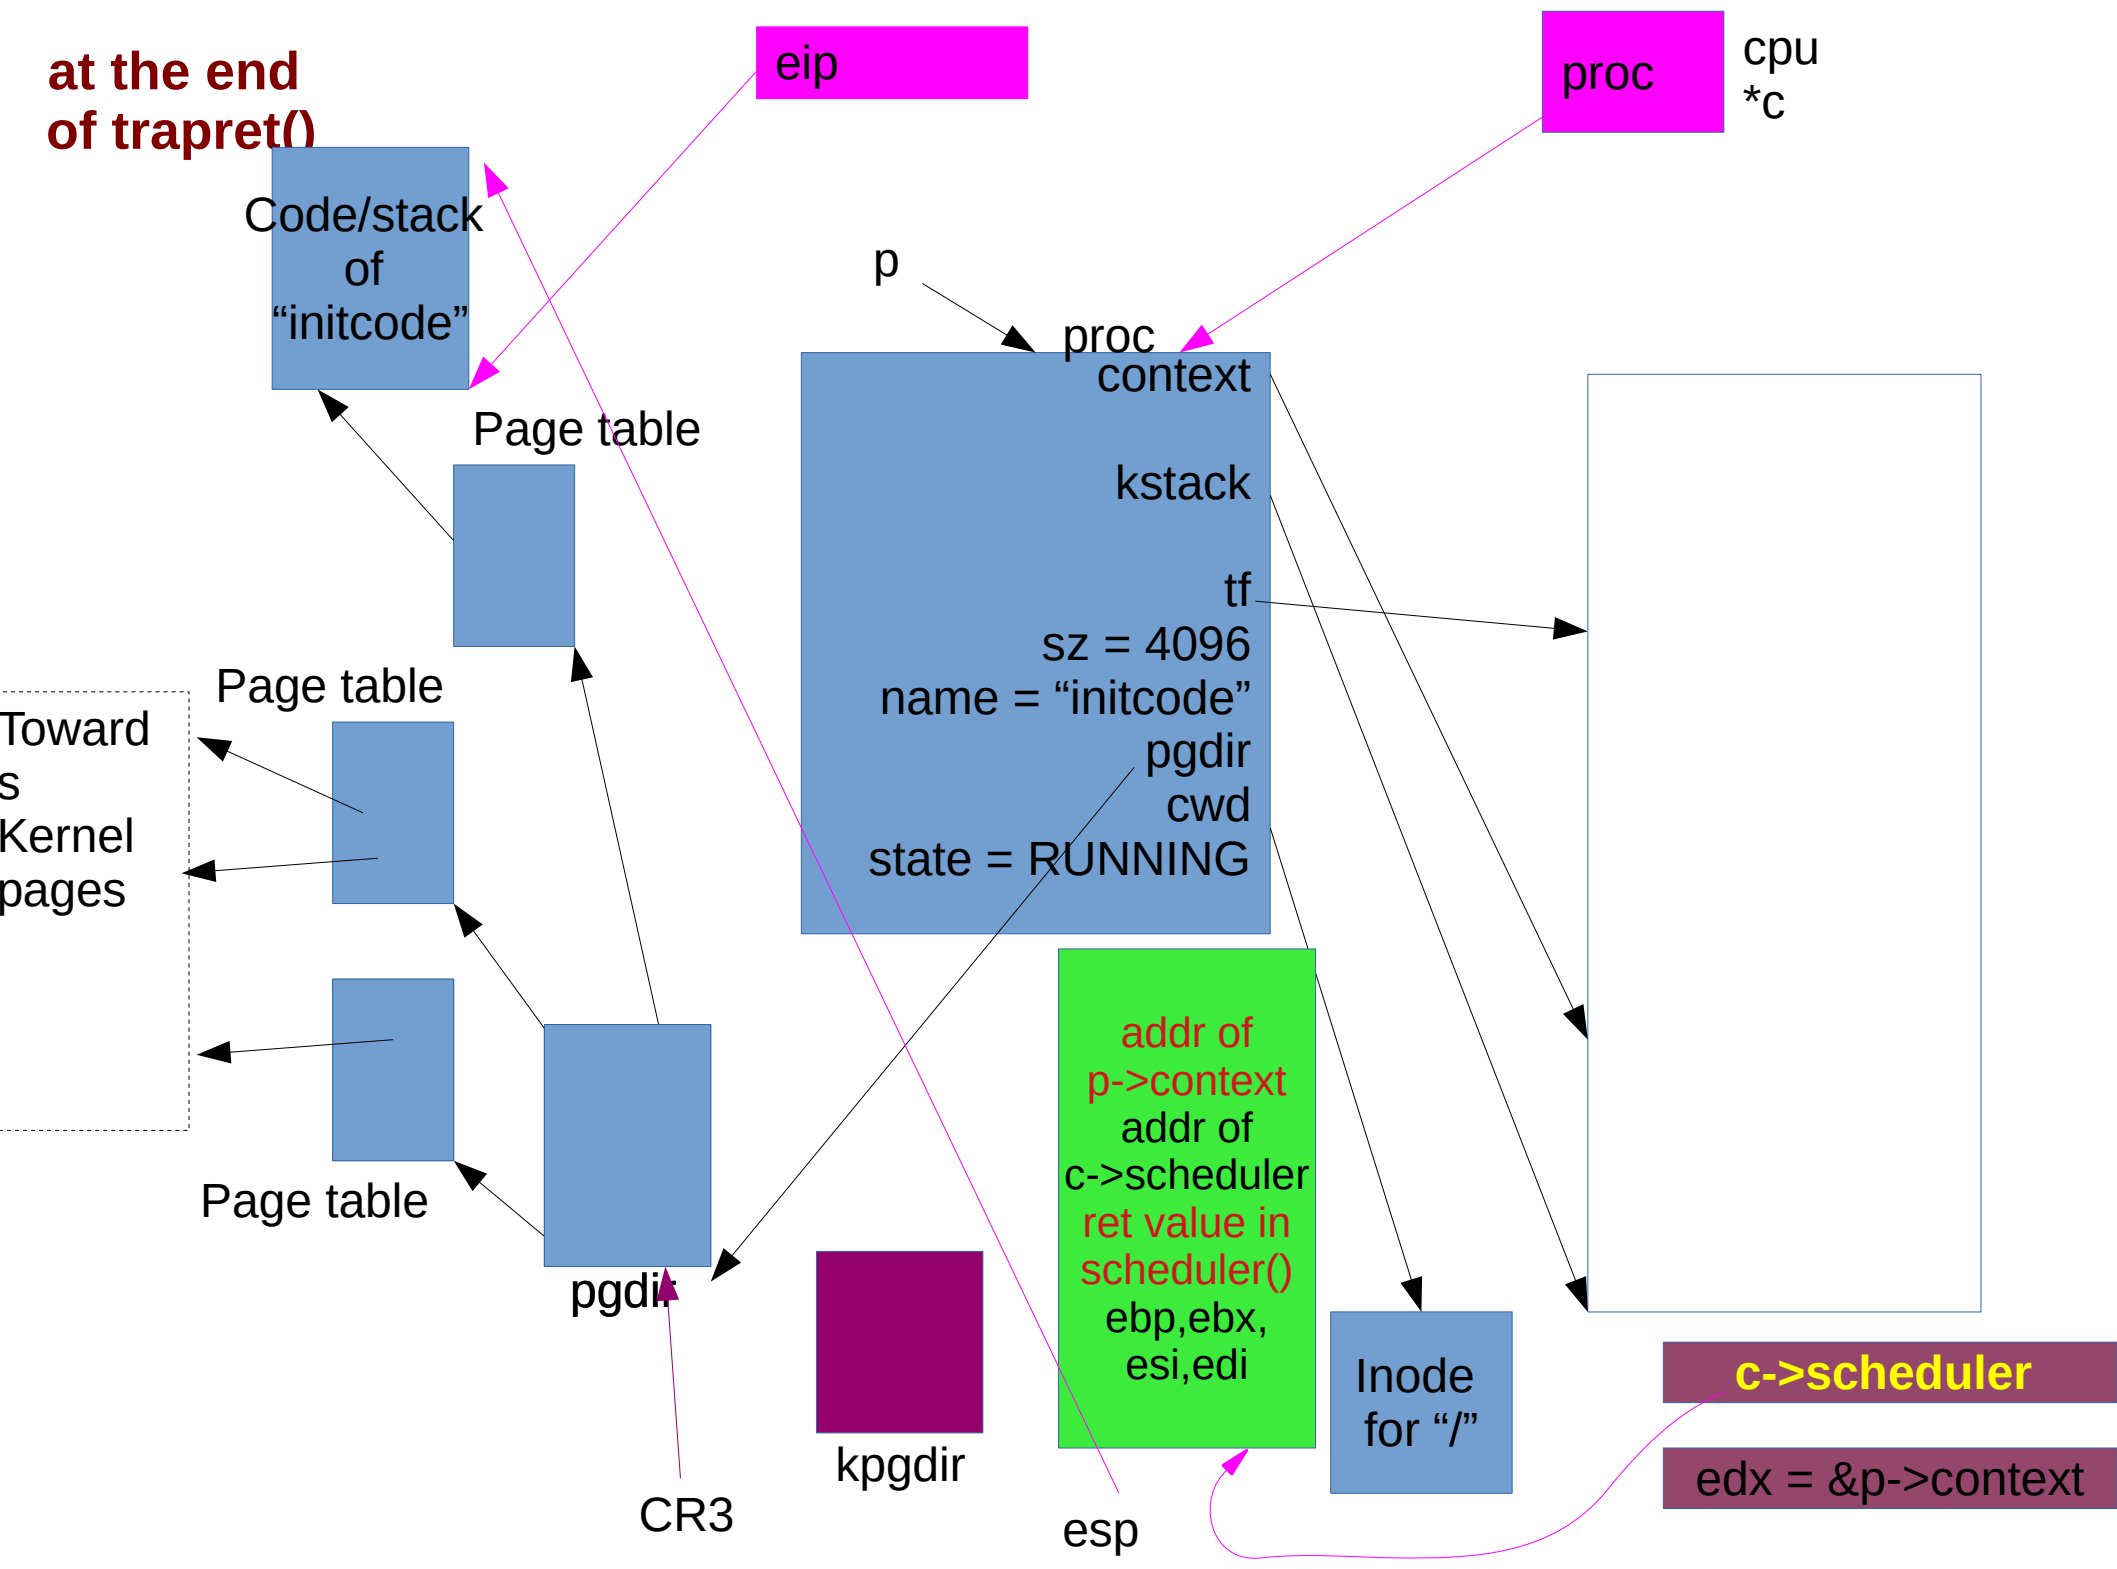
after forkret()
just before
trapret()
begins



After iret in trapret

- **The CS, EIP, ESP will be changed**
 - to values already stored on trapframe
 - this is done by iret
- **Hence after this user code will run**
 - On user stack!
- **Hence code of *initcode* will run now**

at the end
of trapret()



initcode

```
# char init[] = "/init\0";
```

```
init:
```

```
    .string "/init\0"
```

```
# char *argv[] = { init, 0 };
```

```
.p2align 2
```

```
argv:
```

```
    .long init
```

```
    .long 0
```

```
start:
```

```
    pushl $argv
```

```
    pushl $init
```

```
    pushl $0 // where caller pc  
would be
```

```
    movl $SYS_exec, %eax
```

```
    int $T_SYSCALL
```

```
# for(;;) exit();
```

```
exit:
```

```
    movl $SYS_exit, %eax
```

```
    int $T_SYSCALL
```

```
    jmp exit
```

0x24 = addr of argv
0x1c = addr of init
0x0

esp

00000000 <start>:

0:	68 24 00 00 00	push \$0x24
5:	68 1c 00 00 00	push \$0x1c
a:	6a 00	push \$0x0
c:	b8 07 00 00 00	mov \$0x7,%eax
11:	cd 40	int \$0x40

00000013 <exit>:

13:	b8 02 00 00 00	mov \$0x2,%eax
18:	cd 40	int \$0x40
1a:	eb f7	jmp 13 <exit>

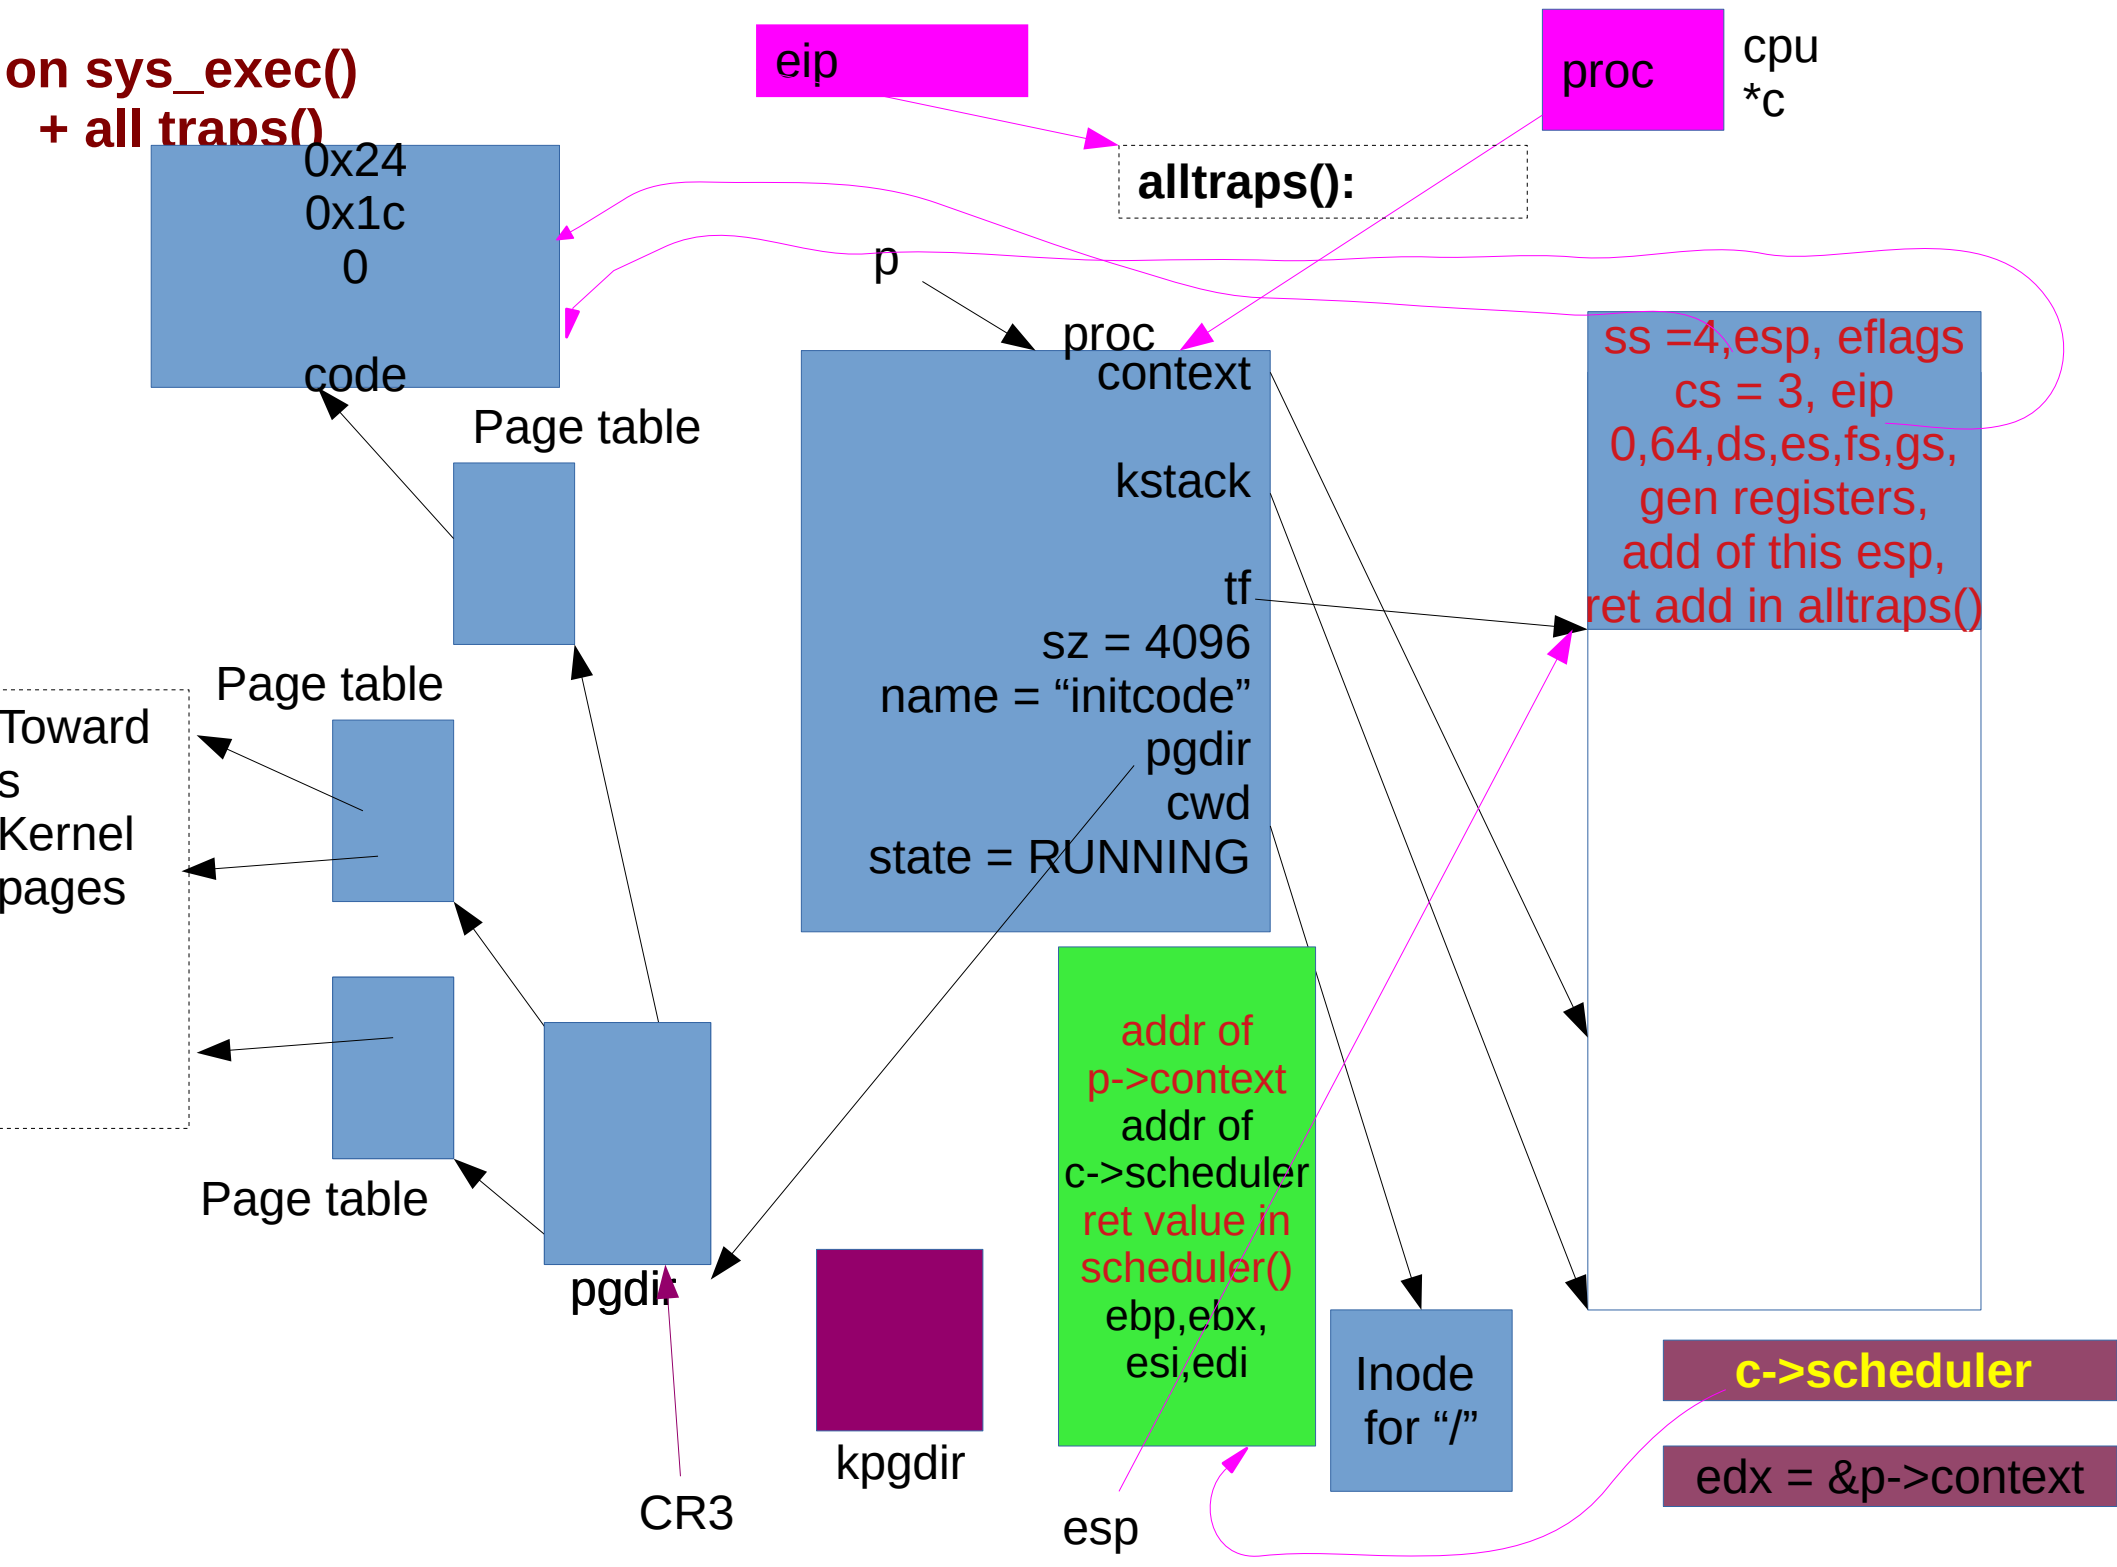
0000001c <init>:

"/init\0"

00000024 <argv>:

1c 00
00 00

on sys_exec()
+ all traps()



Understanding fork() and exec()

**First, revising some concepts already learnt
then code of fork(), exec()**

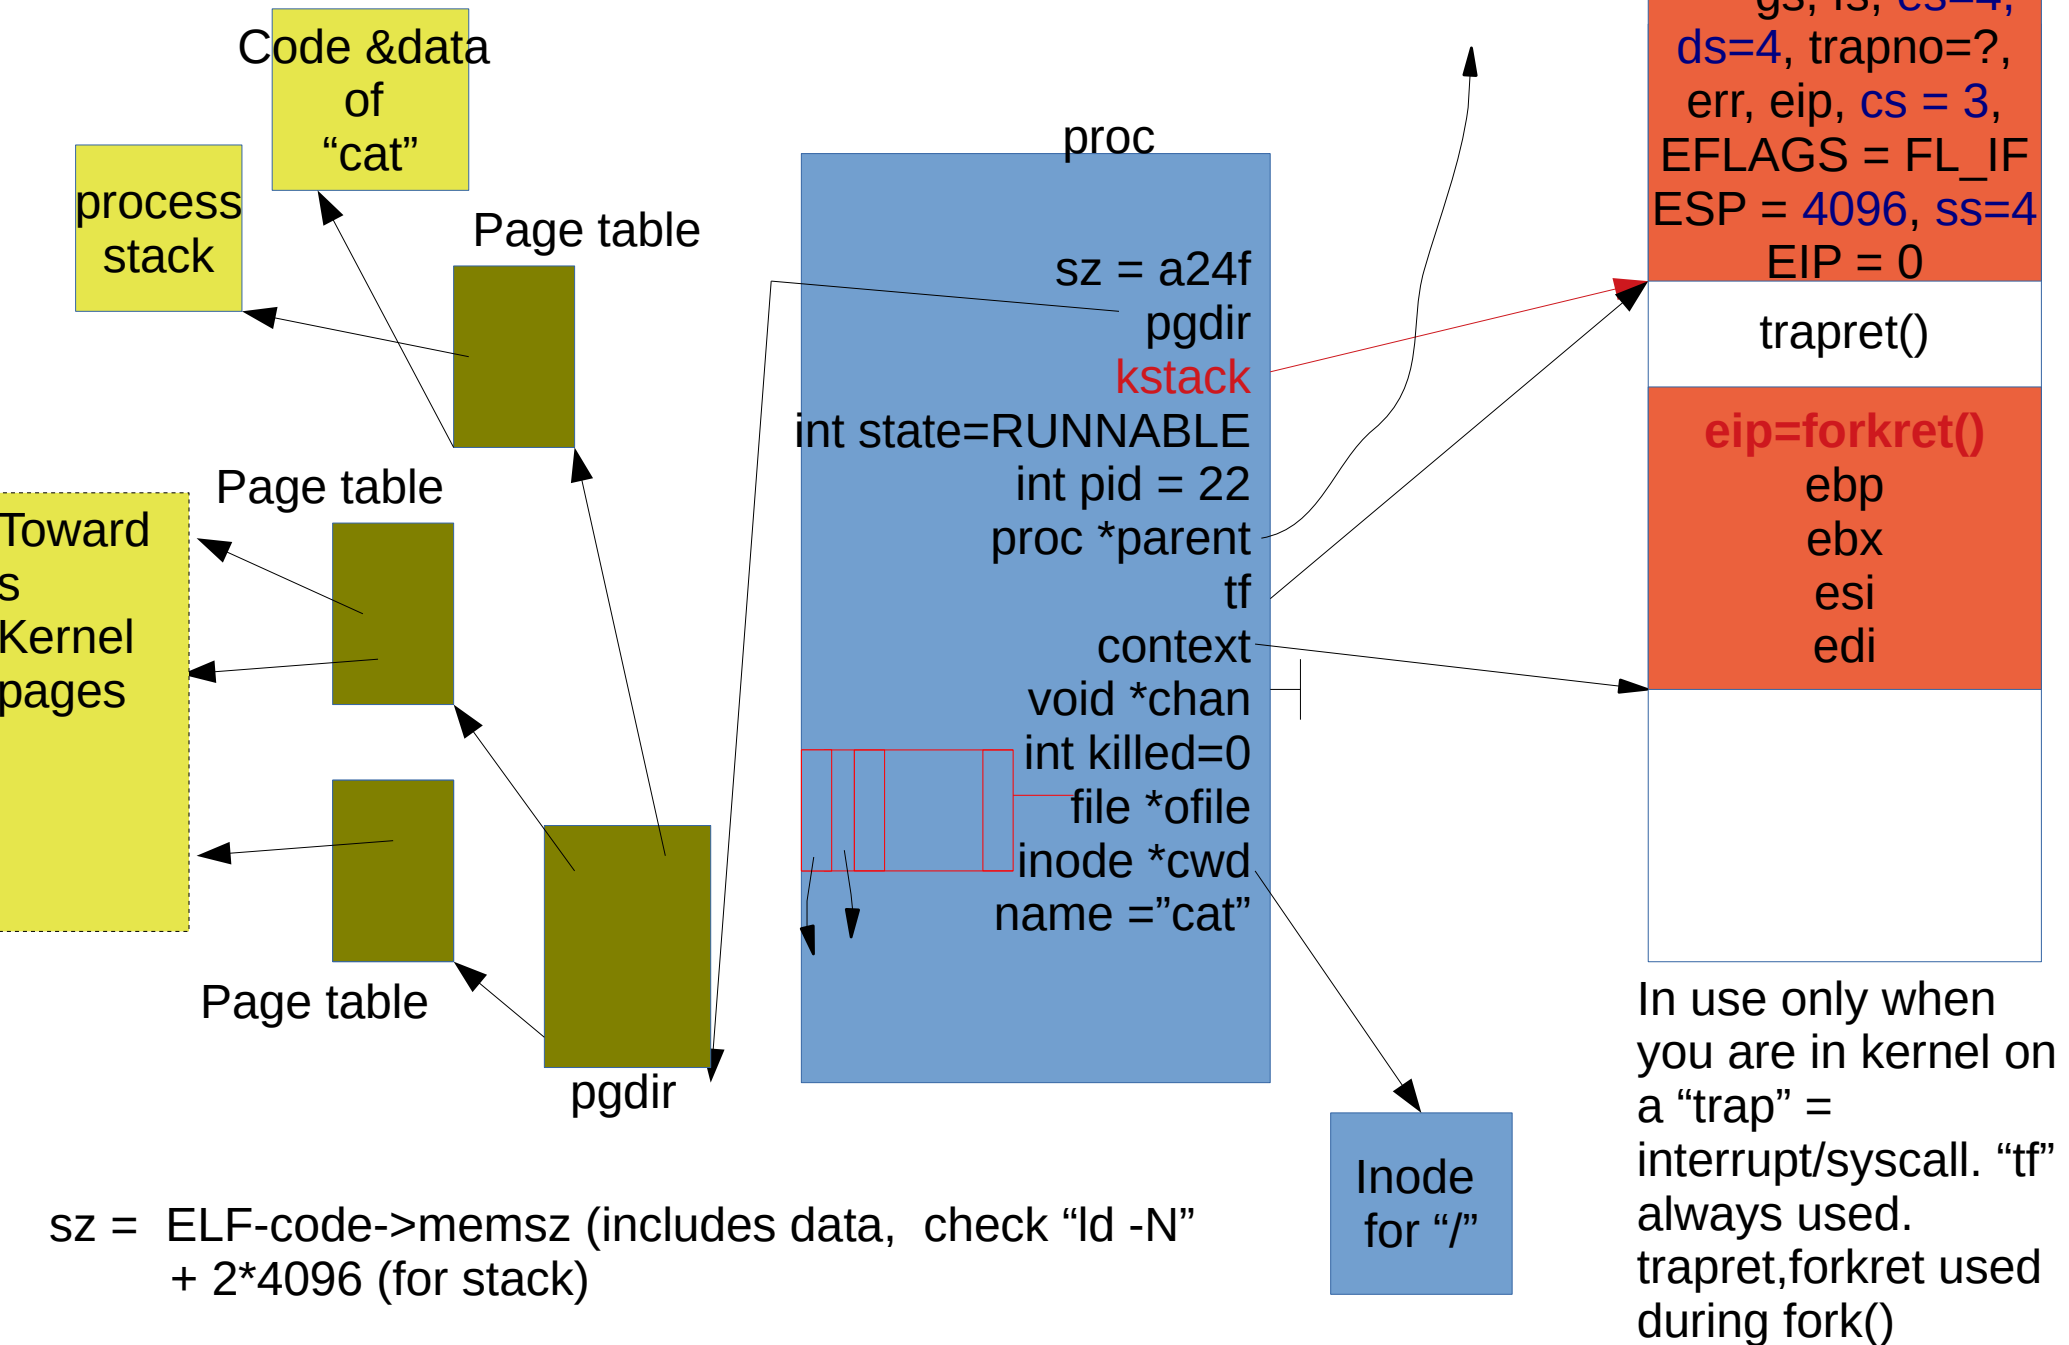
First process creation

Let's revisit struct proc

// Per-process state

```
struct proc {  
    uint sz;                // Size of process memory (bytes)  
    pde_t* pgdir;           // Page table  
    char *kstack;           // Bottom of kernel stack for this process  
    enum procstate state;    // Process state. allocated, ready to run, running,  
wait-                       // ing for I/O, or exiting.  
    int pid;                // Process ID  
    struct proc *parent;     // Parent process  
    struct trapframe *tf;    // Trap frame for current syscall  
    struct context *context; // swtch() here to run process. Process's context  
    void *chan;              // If non-zero, sleeping on chan. More when we discuss  
sleep, wakeup  
    int killed;              // If non-zero, have been killed  
    struct file *ofile[NOFILE]; // Open files, used by open(), read(),...  
    struct inode *cwd;        // Current directory, changed with "chdir()"   
    char name[16];           // Process name (for debugging)  
};
```


struct proc diagram



fork()/exec() are syscalls. On every syscall this happens

- Fetch the n'th descriptor from the IDT, where n is the argument of int.
- Check that CPL in %cs is \leq DPL, where DPL is the privilege level in the descriptor.
- Save %esp and %ss in CPU-internal registers, but only if the target segment selector's PL $<$ CPL.
 - Switching from user mode to kernel mode. Hence save user code's SS and ESP
- Load %ss and %esp from a task segment descriptor.
 - Stack changes to kernel stack now. TS descriptor is on GDT, index given by TR register. See switchvm()
- Push %ss. // optional
- Push %esp. // optional (also changes ss,esp using TSS)
- Push %eflags.
- Push %cs.
- Push %eip.
- Clear the IF bit in %eflags, but only on an interrupt.
- Set %cs and %eip to the values in the descriptor.

After “int” ‘s job is done

- IDT was already set, during idtinit()
 - Remember vectors.S – gives jump locations for each interrupt
- “int 64” -> jump to 64th entry in vector table
 - vector64:
 - pushl \$0
 - pushl \$64
 - jmp alltraps
 - So now stack has **ss, esp, eflags, cs, eip, 0 (for error code), 64**
 - Next run alltraps from trapasm.S

alltraps:

```
# Build trap frame.  
pushl %ds  
pushl %es  
pushl %fs  
pushl %gs  
pushal // push all gen purpose  
regs  
# Set up data segments.  
movw $(SEG_KDATA<<3), %ax  
movw %ax, %ds  
movw %ax, %es  
# Call trap(tf), where tf=%esp  
pushl %esp # first arg to trap()  
call trap  
addl $4, %esp
```

- Now stack contains

ss, esp, eflags, cs, eip, 0
(for error code), 64, ds,
es, fs, gs, eax, ecx, edx,
ebx, oesp, ebp, esi, edi

- This is the struct trapframe !
- So the kernel stack now contains the trapframe
- Trapframe is a part of kernel stack

void

trap(struct trapframe *tf)

{

if(tf->trapno == T_SYSCALL){

if(myproc()->killed)

exit();

myproc()->tf = tf;

syscall();

if(myproc()->killed)

exit();

return;

}

switch(tf->trapno){

.....

trap()

- **Argument is trapframe**

- **In alltraps**

- Before “call trap”, there was “push %esp” and stack had the trapframe

- Remember calling convention --> when a function is called, the stack contains the arguments in reverse order (here only 1 arg)

trap()

- **Has a switch**
 - `switch(tf->trapno)`
 - Q: who set this trapno?
- **Depending on the type of trap**
 - Call interrupt handler
- **Timer**
 - `wakeup(&ticks)`
- **IDE: disk interrupt**
 - `Ideintr()`
- **KBD**
 - `Kbdintr()`
- **COM1**
 - `Uatrintr()`
- **If Timer**
 - Call `yield()` -- calls `sched()`
- **If process was killed (how is that done?)**
 - Call `exit()`!

when trap() returns

- **#Back in alltraps**

call trap

addl \$4, %esp

Return falls through to trapret...

.globl trapret

trapret:

popal

popl %gs

popl %fs

popl %es

popl %ds

addl \$0x8, %esp # trapno and errcode

iret

- **Stack had (trapframe)**

- **ss, esp,eflags, cs, eip, 0 (for error code), 64, ds, es, fs, gs, eax, ecx, edx, ebx, oesp, ebp, esi, edi, esp**

- **add \$4 %esp**

- **esp**

- **popal**

- **eax, ecx, edx, ebx, oesp, ebp, esi, edi**

- **Then gs, fs, es, ds**

- **add \$0x8, %esp**

- **0 (for error code), 64**

- **iret**

- **ss, esp,eflags, cs, eip,**

understanding fork()

- **What should fork do?**
 - Create a copy of the existing process
 - child is same as parent, except pid, parent-child relation, return value (pid or 0)
 - Please go through every member of struct proc, understand it's meaning to appreciate what fork() should do
 - create a struct proc, and
 - duplicate pages, page directory, sz, state, trapframe, context, ofile (and files!), cwd, name
 - modify: pid, parent, trapframe, state

understanding fork()

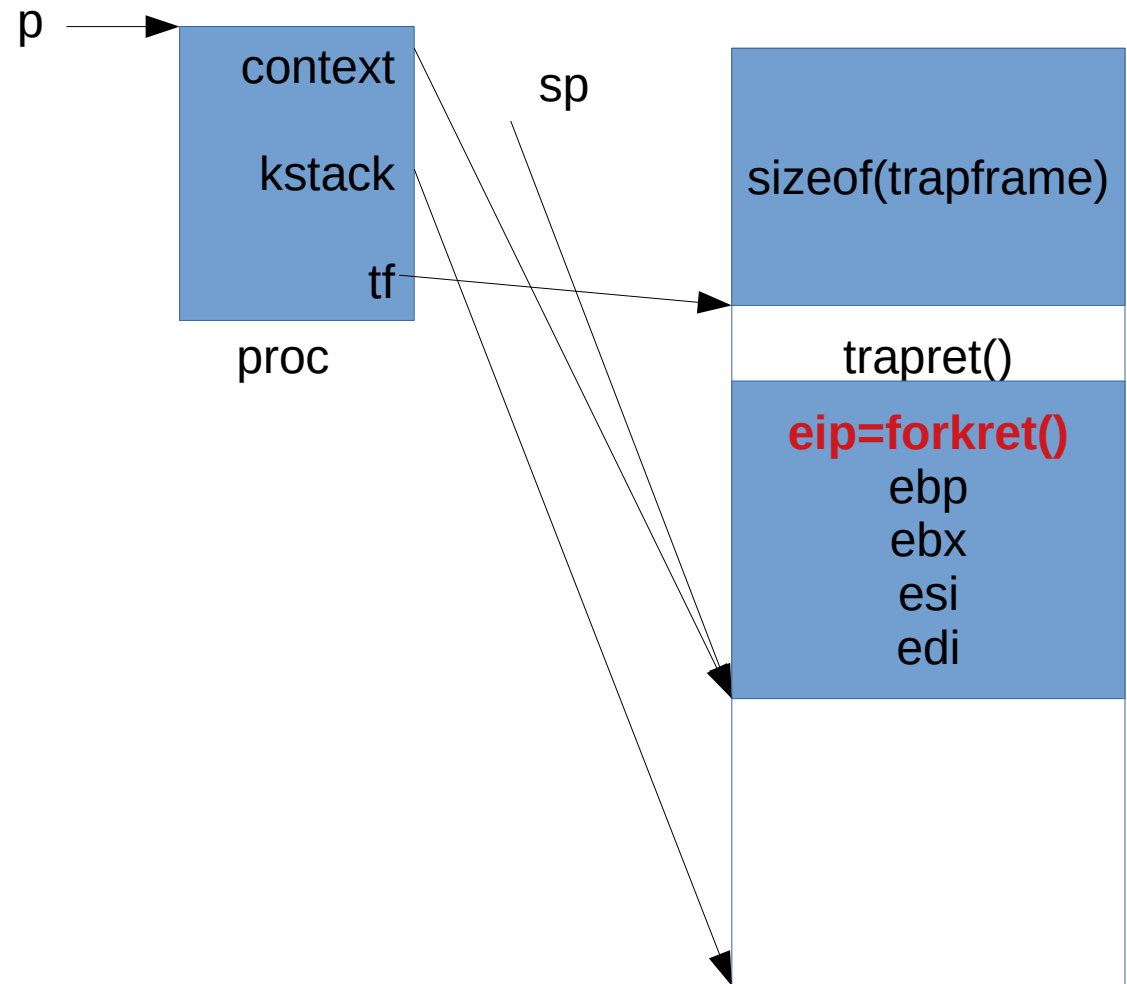
```
int
sys_fork(void)
{
    return fork();
}
```

```
int
fork(void)
{
    int i, pid;
    struct proc *np;
    struct proc *curproc = myproc();

    // Allocate process.
    if((np = allocproc()) == 0){
        return -1;
    }
```

after allocproc()

-- we studied this -- same as creation of first process



understanding fork()

```
// Copy process state from proc.  
if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){  
    kfree(np->kstack);  
    np->kstack = 0;  
    np->state = UNUSED;  
    return -1;  
}  
np->sz = curproc->sz;  
np->parent = curproc;  
*np->tf = *curproc->tf;
```

- **copy the pages, page tables, page directory**
 - no copy on write here!
 - Rewind if operation of copyuvm() fails
- **copy size**
- **set parent of child**
- **copy trapframe (structure is copied)**

```

pde_t*
copyuvm(pde_t *pgdir, uint sz)
{
    pde_t *d; pte_t *pte; uint pa, i, flags;
    char *mem;
    if((d = setupkvm()) == 0)
        return 0;
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
            panic("copyuvm: pte should exist");
        if(!(*pte & PTE_P))
            panic("copyuvm: page not present");
        pa = PTE_ADDR(*pte);
        flags = PTE_FLAGS(*pte);
        if((mem = kalloc()) == 0)
            goto bad;
        memmove(mem, (char*)P2V(pa), PGSIZE);
        if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0) {
            kfree(mem);
            goto bad;
        }
    }
    return d;
bad:
    freevm(d);
    return 0;
}

```

understanding fork()->copyuvm()

- Map kernel pages
- for every page in parent's VM address space
 - allocate a PTE for child
 - set flags
 - copy data
 - map pages in child's page directory/tables

understanding fork()

```
np->tf->eax = 0;
for(i = 0; i < NOFILE; i++)
    if(curproc->ofile[i])
        np->ofile[i] = filedup(curproc-
>ofile[i]);
np->cwd = idup(curproc->cwd);
safestrcpy(np->name, curproc-
>name, sizeof(curproc->name));
pid = np->pid;
acquire(&ptable.lock);
np->state = RUNNABLE;
release(&ptable.lock);
```

- set return value of child to 0
 - eax contains return value, it's on TF
- copy each struct file
- copy current working dir inode
- copy name
- set pid of child
- set child "RUNNABLE"

exec() - different prototype

- **int exec(char*, char**);**
 - usage: to print README and test.txt using “cat”

```
int main(int argc, char *argv[])
{
    char *cmd = "/cat";

    char *argstr[4] = { "/cat", "README",
"test.txt", 0};

    exec(cmd, argstr);
}
```

note: to really run this code in xv6, you need to make changes to Makefile. First, add this program to UPROGS, then write a file test.txt using Linux, and add 'test.txt' to list of files in 'mkfs' target in Makefile

```

int
sys_exec(void)
{
    char *path, *argv[MAXARG];
    int i;
    uint uargv, uarg;
    if(argstr(0, &path) < 0 || argint(1, (int*)&uargv) < 0){
        return -1;
    }
    memset(argv, 0, sizeof(argv));
    for(i=0;; i++){
        if(i >= NELEM(argv))
            return -1;
        if(fetchint(uargv+4*i, (int*)&uarg) < 0)
            return -1;
        if(uarg == 0){
            argv[i] = 0;
            break;
        }
        if(fetchstr(uarg, &argv[i]) < 0)
            return -1;
    }
    return exec(path, argv);
}

```

sys_exec()

- **argstr(n,), argint(n,)**
 - Fetch the n'th argument from *process stack* using p->tf->esp + offset
 - Again: revise calling conventions
 - **0'th argument:** name of executable file
 - **1st Argument:** address of the array of arguments
 - store in *uargv*

```

int sys_exec(void)
{
    char *path, *argv[MAXARG];
    int i;  uint uargv, uarg;
    if(argstr(0, &path) < 0 || argint(1,
(int*)&uargv) < 0){
        return -1;
    }
    memset(argv, 0, sizeof(argv));
    for(i=0;; i++){
        if(i >= NELEM(argv))    return -1;
        if(fetchint(uargv+4*i, (int*)&uarg) < 0)
            return -1;
        if(uarg == 0){
            argv[i] = 0;    break;
        }
        if(fetchstr(uarg, &argv[i]) < 0)
            return -1;
    }
    return exec(path, argv);
}

```

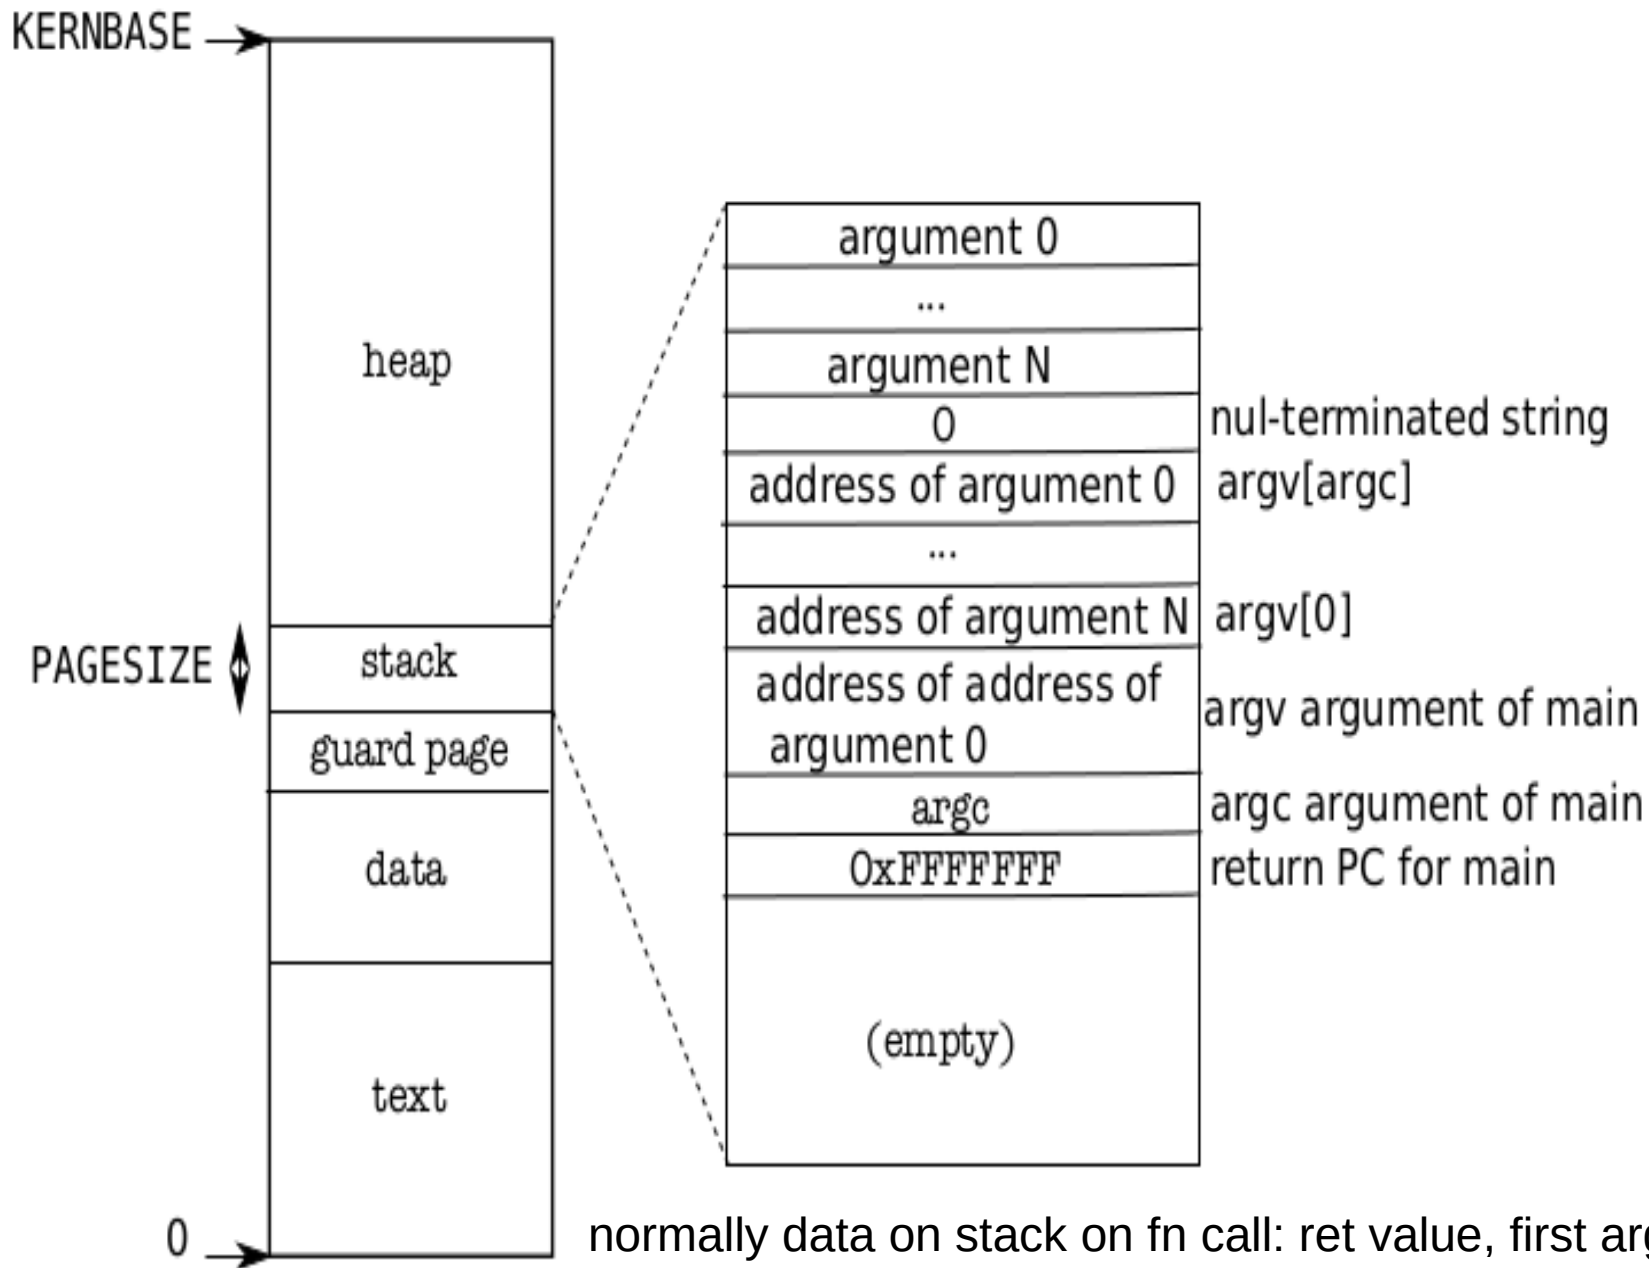
sys_exec()

- the local array argv[]
(allocated on kernel stack,
obviously) set to 0
- fetch every next argument
from array of arguments
 - Sets the address of
argument in argv[1]
- call exec
 - beware: mistake to assume
that this exec() is the exec()
called from user code! NO!

What should `exec()` do?

- **Remember, it came from `fork()`**
 - so proc & within it tf, context, kstack, pgdir-tables-pages, all exist.
 - Code, stack pages exist, and mappings exist through `proc->pgdir`
- **Hence**
 - read the ELF executable file (`argv[0]`)
 - create a new page dir – create mappings for kernel and user code+data; copy data from ELF to these pages (later discard old pagedir)
 - Copy the `argv` onto the user stack – so that when new process starts it has its `main(argc, argv[])` built
 - set values of other fields in proc to start program correctly

User
stack
after
call
to
exec()
is over



normally data on stack on fn call: ret value, first arg, second arg, ...
main(int argc, char *argv[])
argv[] is address of array of string; string itself is an adress. Hence
2 levels of indirection on stack

exec()

```
int
exec(char *path, char **argv)
{
...
    uint argc, sz, sp,
    ustack[3+MAXARG+1];
...
    if((ip = namei(path)) == 0){
        end_op();
        cprintf("exec: fail\n");
        return -1;
    }
```

- **ustack**
 - used to build the arguments to be pushed on user-stack
- **namei**
 - get the inode of the executable file

exec()

// Check ELF header

if(readi(ip, (char*)&elf, 0,
sizeof(elf)) != sizeof(elf))

goto bad;

if(elf.magic != ELF_MAGIC)

goto bad;

if((pgdir = setupkvm()) == 0)

goto bad;

- **readi**

- read ELF header

- **setupkvm()**

- creating a *new* page directory and mapping kernel pages

```

sz = 0;
for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
    if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
        goto bad;
    if(ph.type != ELF_PROG_LOAD)
        continue;
    if(ph.memsz < ph.filesz)
        goto bad;
    if(ph.vaddr + ph.memsz < ph.vaddr)
        goto bad;
    if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
        goto bad;
    if(ph.vaddr % PGSIZE != 0)
        goto bad;
    if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) <
0)
        goto bad;
}

```

exec()

- Read ELF program headers from ELF file
- Map the code/data into pagedir-pagetable-pages
- Copy data from ELF file into the pages allocated

exec()

```
sz = PGROUNDUP(sz);  
if((sz = allocvm(pgdir, sz, sz +  
2*PGSIZE)) == 0)  
    goto bad;  
clearpteu(pgdir, (char*)(sz -  
2*PGSIZE));  
sp = sz;
```

- Allocate 2 pages on top of proc->sz
- One page for stack
- one page for guard page
- Clear the valid flag on guard page

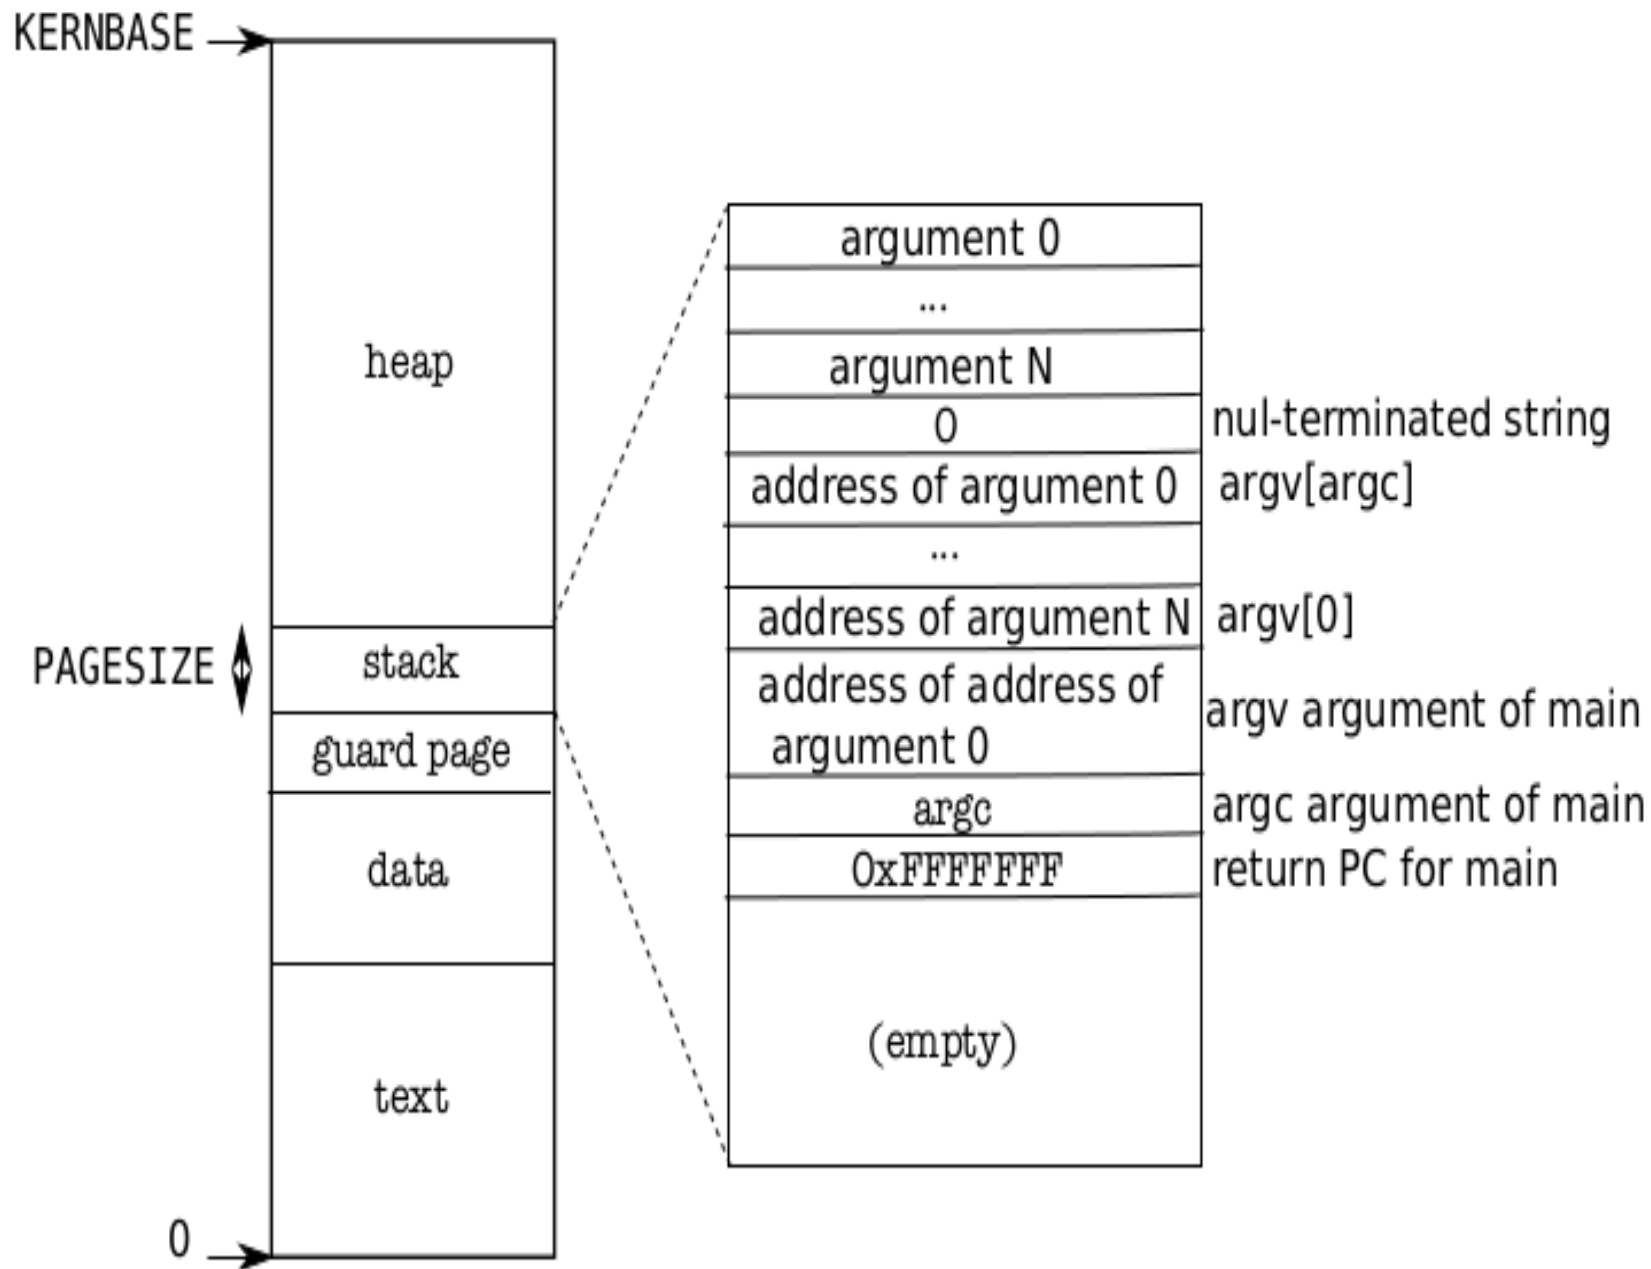
// Push argument strings, prepare rest of stack
in ustack.

```
for(argc = 0; argv[argc]; argc++) {  
    if(argc >= MAXARG)  
        goto bad;  
    sp = (sp - (strlen(argv[argc]) + 1)) & ~3;  
    if(copyout(pgdir, sp, argv[argc],  
strlen(argv[argc]) + 1) < 0)  
        goto bad;  
    ustack[3+argc] = sp;  
}  
ustack[3+argc] = 0;  
ustack[0] = 0xffffffff; // fake return PC  
ustack[1] = argc;  
ustack[2] = sp - (argc+1)*4; // argv pointer  
sp -= (3+argc+1) * 4;  
if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)  
    goto bad;
```

exec()

- For each entry in argv[]
 - copy it on user-stack
 - remember it's location on user stack in ustack
- add extra entries (to be copied to user stack) to ustack
- copy argc, argv pointer
- take sp to bottom
- copy ustack to user stack

**This is
what the
code on
earlier
slide did**



// Save program name for debugging.

```
for(last=s=path; *s; s++)
```

```
    if(*s == '/')
```

```
        last = s+1;
```

```
    safestrcpy(curproc->name, last,  
sizeof(curproc->name));
```

// Commit to the user image.

```
oldpgdir = curproc->pgdir;
```

```
curproc->pgdir = pgdir;
```

```
curproc->sz = sz;
```

```
curproc->tf->eip = elf.entry; // main
```

```
curproc->tf->esp = sp;
```

```
switchvm(curproc);
```

```
freevm(oldpgdir);
```

```
return 0;
```

exec()

- copy name of new process in `proc->name`
- change to new page directory
- change new size
- `tf->eip` will be used when we return from `exec()` to jump to user code. Set to first instruction of code, given by `elf.entry`
- Set user stack pointer to “sp” (bottom of stack of arguments)
- Update TSS, change CR3 to `newpagedir`
- free old page dir

return 0 from exec()?

- We know exec() does not return !
- This was exec() function !
 - Returns to sys_exec()
- sys_exec() also returns , where?
 - Remember we are still in kernel code, running on kernel stack. p->kstack has the trapframe setup
 - There is context struct on stack. Why?
 - sys_exec() returns to trapret(), the trap frame will be popped !
 - with “iret” jump into new program !
 - New program is not old program , which could have accessed return value of sys_exec()